

Image segmentation – part 2



Lihi Zelnik-Manor, Computer Vision

Today

- ▶ **Graph theoretic segmentation**
 - ▶ Normalized cuts
- ▶ **Segmentation as energy minimization**
 - ▶ Markov random fields


Images as graphs

- ▶ Node for every pixel
- ▶ Edge between every pair of pixels (or every pair of “sufficiently close” pixels)
- ▶ Each edge is weighted by the *affinity* or similarity of the two nodes



Segmentation by graph partitioning

- ▶ Break Graph into Segments
 - ▶ Delete links that cross between segments
 - ▶ Easiest to break links that have low affinity
 - ▶ similar pixels should be in the same segments
 - ▶ dissimilar pixels should be in different segments






Measuring affinity

- **Distance** $aff(x, y) = \exp \left\{ -\frac{1}{2\sigma_d^2} \|x - y\|^2 \right\}$
- **Intensity** $aff(x, y) = \exp \left\{ -\frac{1}{2\sigma_d^2} \|I(x) - I(y)\|^2 \right\}$
- **Color** $aff(x, y) = \exp \left\{ -\frac{1}{2\sigma_d^2} \underbrace{dist(c(x), c(y))}_{\text{(some suitable color space distance)}}^2 \right\}$
- **Texture** $aff(x, y) = \exp \left\{ -\frac{1}{2\sigma_d^2} \underbrace{\|f(x) - f(y)\|}_{\text{(vectors of filter outputs)}}^2 \right\}$


Scale affects affinity

- ▶ Small σ : group only nearby points
- ▶ Large σ : group far-away points



Graph cut



- ▶ Set of edges whose removal makes a graph disconnected
- ▶ Cost of a cut: sum of weights of cut edges
- ▶ A graph cut gives us a segmentation
 - ▶ What is a “good” graph cut and how do we find one?



Minimum cut

- We can segment by finding the *minimum cut* in a graph
 - ▶ Efficient algorithms exist for doing this




Minimum cut example



Minimum cut

- We can segment by finding the *minimum cut* in a graph
 - ▶ Efficient algorithms exist for doing this



Minimum cut example



- The cut is defined by the block diagonal structure of the affinity matrix.
- Can this be generalized?

Minimum cut drawback

- ▶ Minimum cut tends to cut off very small, isolated components



Normalized cut (Ncut)

- ▶ The drawback of mincut can be fixed by normalizing the cost by the weight of all the edges incident to the segment
- ▶ The *normalized cut* cost is:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

$cut(A, B)$ = sum of weights of all edges between A and B

$assoc(A, V)$ = sum of weights of all edges in V that touch A

$$Ncut(A, B) = cut(A, B) \frac{1}{\sum_{p \in A} w_{p,q}} + \frac{1}{\sum_{q \in B} w_{p,q}}$$

J. Shi and J. Malik. [Normalized cuts and image segmentation](#). PAMI 2000

Ncut as a generalized eigenvector problem

- Let W be the adjacency matrix of the graph
- Let D be the diagonal matrix with diagonal entries $D(i, i) = \sum_j W(i, j)$
- Then the normalized cut cost can be written as

$$Ncut(A, B) = \frac{y^T (D - W) y}{y^T D y}$$

$$y(p) = \begin{cases} 1 & p \in A \\ \text{negative} & \text{otherwise} \end{cases}$$

J. Shi and J. Malik. [Normalized cuts and image segmentation](#). PAMI 2000

Ncut as a generalized eigenvector problem

- ▶ Problem:

Finding the exact minimum of the normalized cut cost is NP-complete (because y is discrete)

- ▶ Solution:

- ▶ Relax y to take on arbitrary values, then solved by a *generalized eigenvalue problem*

$$N(D - W)y = \lambda Dy$$

- ▶ The solution y is given by the eigenvector corresponding to the second smallest eigenvalue
- ▶ Continuous results need to be converted into discrete

Normalized cut algorithm

1. Construct a weighted graph $G = (V, E)$ from an image
2. Connect each pair of pixels, and assign weights
 $w(i, j) = \text{prob}(i, j \text{ belong to same region})$
3. Compute diagonal matrix $D(i, i) = \sum_j W(i, j)$
4. Solve $(D - W)y = \lambda Dy$ for the eigenvector with the second smallest eigenvalue
5. Threshold eigenvector to get a discrete cut
6. Recursively partition the segmented parts, if necessary


Example results



Example results



Results: Berkeley Segmentation Engine




<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: Pros and cons

- **Pros**
 - ▶ Generic framework, can be used with many different features and affinity formulations
 - ▶ No model or data distribution
- **Cons**
 - ▶ High storage requirement and time complexity
 - ▶ Bias towards partitioning into equal segments






Eigenvectors carry contour information




Avoiding Ncut drawback

- ▶ Do **not** try to find regions from the eigenvectors



Avoiding Ncut drawback

- ▶ Do **not** try to find regions from the eigenvectors
- ▶ Key idea:
 - ▶ Reshape eigenvectors into images
 - ▶ Compute edge probability P_b on eigenvector images
 - ▶ Final edge probability is the sum of all responses



Example results

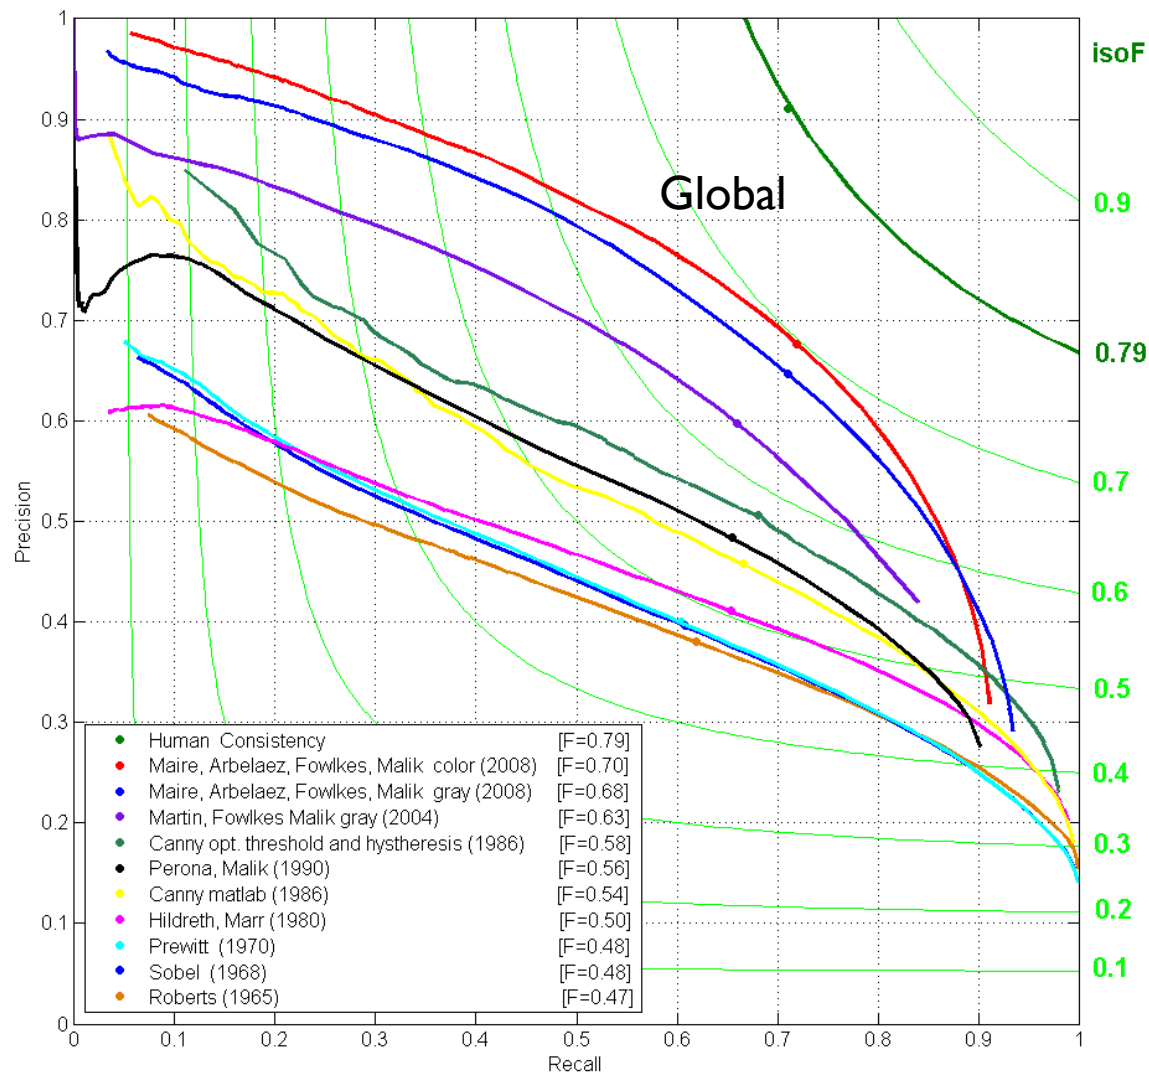
Final Pb

Pb

After threshold continuous




Contour detection ~2008 (color)



Today

- ▶ Graph theoretic segmentation
 - ▶ Normalized cuts
- ▶ Segmentation as energy minimization
 - ▶ Markov random fields

Segmentation as energy minimization



$P(\text{foreground} \mid \text{image})$

1 0 , Y

Normalizing constant called “partition function”

$$P(\mathbf{y}; \theta, \text{data}) = \frac{1}{Z} \prod_{i=1..N} f_1(y_i; \theta, \text{data}) \prod_{i,j \in \text{edges}} f_2(y_i, y_j; \theta, \text{data})$$

Labels to be predicted


Individual predictions

Pairwise predictions

_Y

Writing Likelihood as an “Energy”

$$P(\mathbf{y}; \theta, data) = \frac{1}{Z} \prod_{i=1..N} p_1(y_i; \theta, data) \prod_{i,j \in edges} p_2(y_i, y_j; \theta, data)$$



$$Energy(\mathbf{y}; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j; \theta, data)$$

Cost of assignment y_i

Cost of pairwise assignment y_i, y_j




Notes on energy-based formulation

$$\text{Energy}(\mathbf{y}; \theta, \text{data}) = \sum_i \psi_1(y_i; \theta, \text{data}) + \sum_{i,j \in \text{edges}} \psi_2(y_i, y_j; \theta, \text{data})$$

- ▶ Primarily used when you only care about the most likely solution (not the confidences)
- ▶ Can think of it as a general cost function
- ▶ Can have larger “cliques” than 2
 - ▶ Clique is the set of variables that go into a potential function



Markov Random Fields




Cost to assign a label to
each pixel

Cost to assign a pair of labels to
connected pixels

$$Energy(\mathbf{y}; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j; \theta, data)$$

Markov Random Fields

► Example: “label smoothing” grid



Unary potential

0: $-\log P(y_i = 0 \mid \text{data})$

1: $-\log P(y_i = 1 \mid \text{data})$

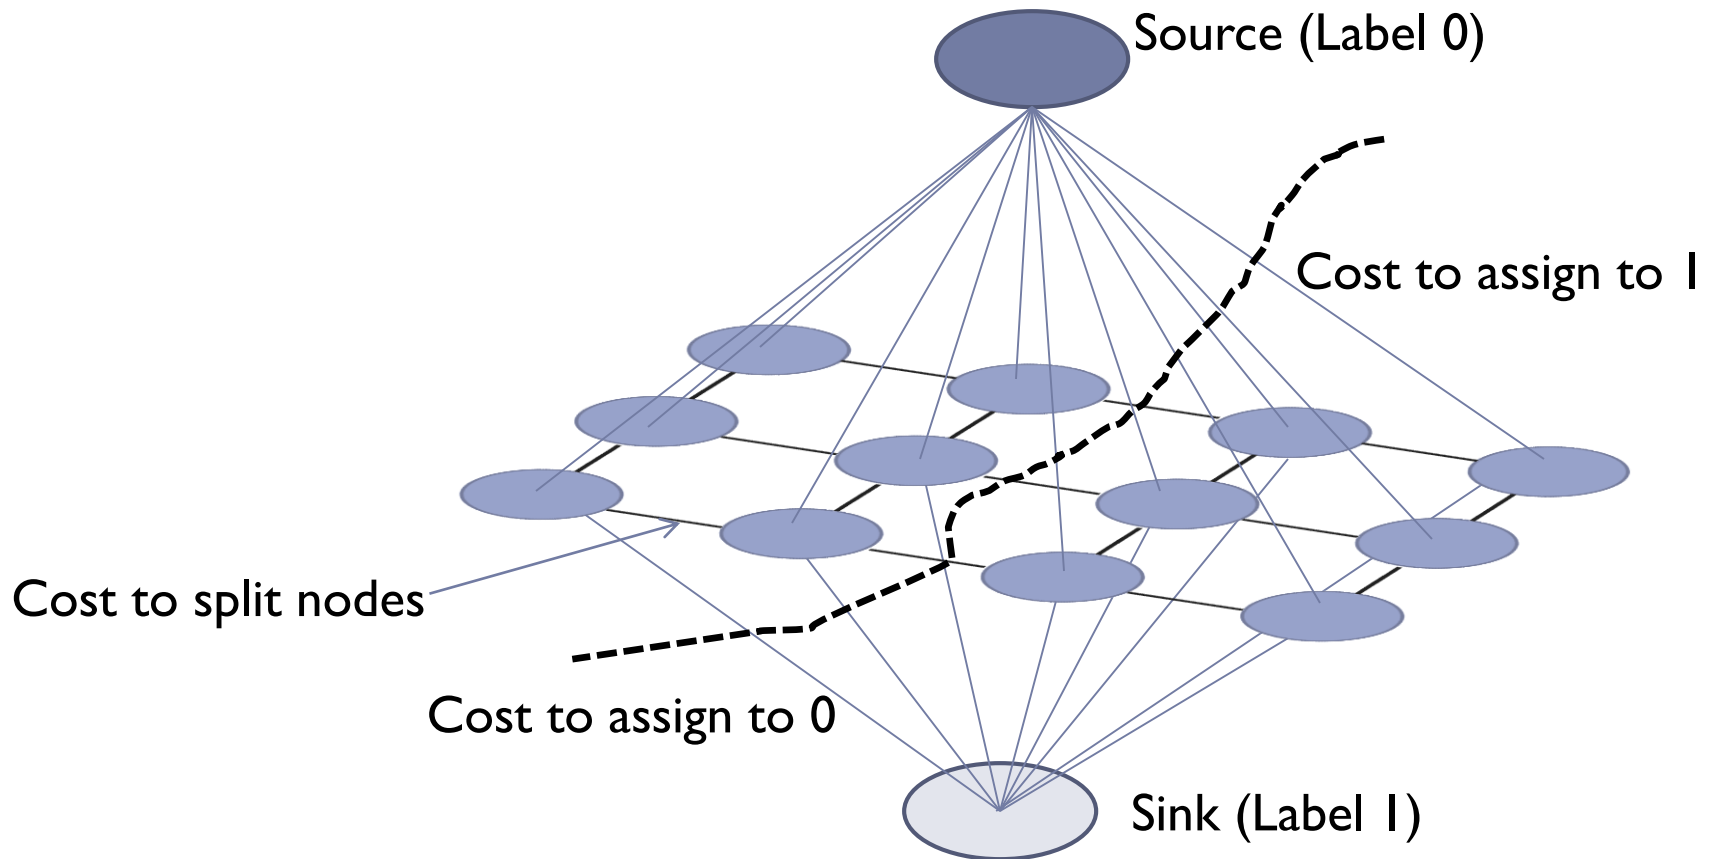
Pairwise Potential

	0	1
0	0	K
1	K	0

$K > 0$

$$\text{Energy}(\mathbf{y}; \theta, \text{data}) = \sum_i \psi_1(y_i; \theta, \text{data}) + \sum_{i,j \in \text{edges}} \psi_2(y_i, y_j; \theta, \text{data})$$

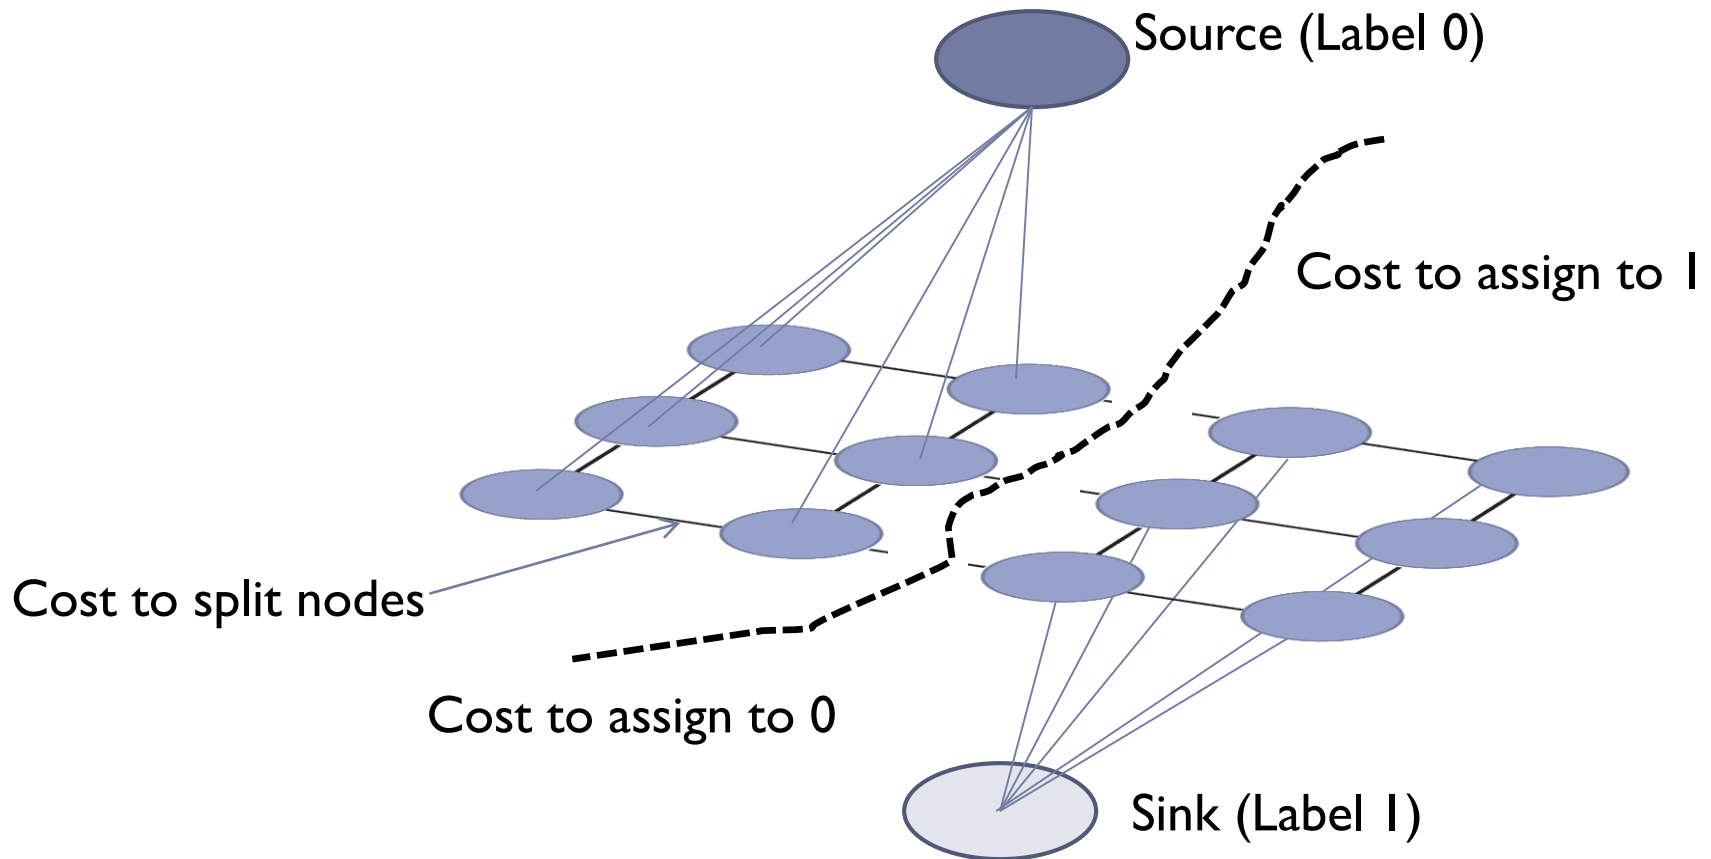
Solving MRFs with graph cuts



$$Energy(\mathbf{y}; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j; \theta, data)$$





Solving MRFs with graph cuts



$$Energy(\mathbf{y}; \theta, data) = \sum_i \psi_1(y_i; \theta, data) + \sum_{i,j \in edges} \psi_2(y_i, y_j; \theta, data)$$

GrabCut segmentation



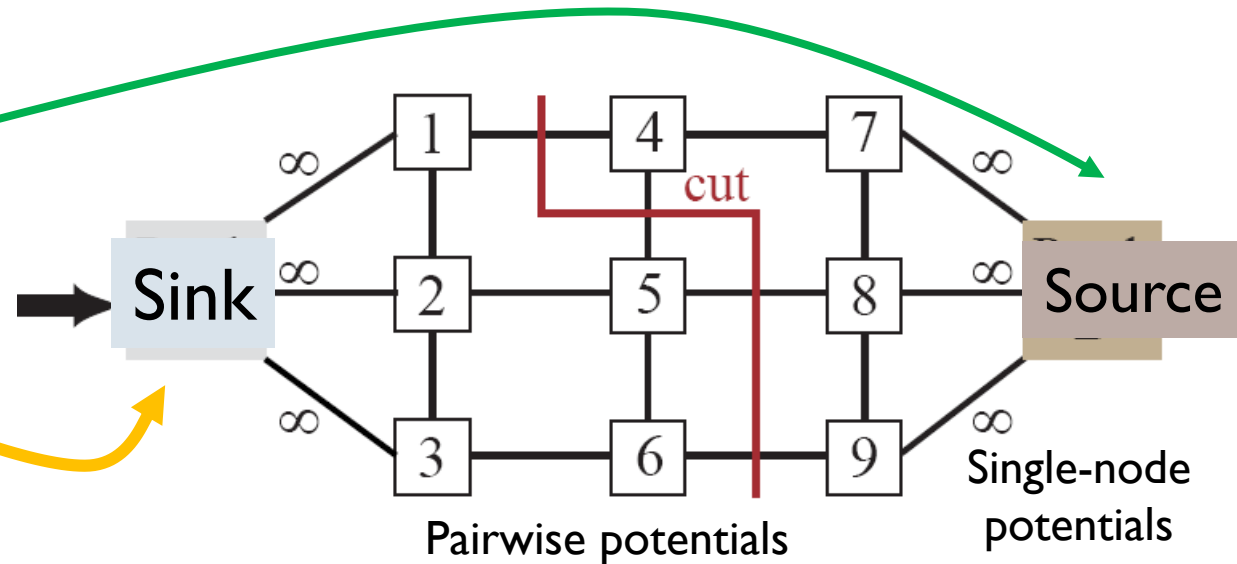
User provides rough indication of foreground region.

Goal: Automatically provide a pixel-level segmentation.



GrabCut

- Convert MRF into source-sink graph





- Minimum cost can be computed in polynomial time

[Kwatra et al., SIGGRAPH 2003]

Graph cuts

Boykov and Jolly (2001)

Image



Cut: separating source and sink; Energy: collection of edges

Min Cut: Global minimal energy in polynomial time

Binary segmentation as energy minimization

- ▶ Define a labeling L as an assignment of each pixel with a 0-1 label (background or foreground)
- ▶ Problem statement: find the labeling L that minimizes

האנרגיה הכוללת

$$E(L) = E_d(L) + \lambda E_s(L)$$

העלות למימוש

match cost

smoothness cost

(how similar is each pixel to the foreground / background?)



העלות למימוש
העלות למימוש
העלות למימוש
העלות למימוש
העלות למימוש

$$E(L) = E_d(L) + \lambda E_s(L)$$



$$E_d(L) = \sum_{(x,y)} C(x, y, L(x, y))$$


$$C(x, y, L(x, y)) = \begin{cases} \infty & \text{if } L(x, y) \neq \tilde{L}(x, y) \\ C'(x, y, L(x, y)) & \text{otherwise} \end{cases}$$

$C'(x, y, 0)$: “distance” from pixel to background pixels
 $C'(x, y, 1)$: “distance” from pixel to foreground pixels


} usually computed by creating a color model from user-labeled pixels

Colour Model

Handwritten notes in Arabic script, possibly indicating a process or model related to the title.



Iterated
graph cut



Gaussian Mixture Model (typically 5-8 components)

$$E(L) = E_d(L) + \lambda E_s(L)$$

- ▶ Neighboring pixels should generally have the same labels
 - ▶ Unless the pixels have very different intensities

Handwritten notes in Hebrew: "אם הפיקסלים הם באותו צבע אז הם צריכים להיות באותו קטגוריה" (If the pixels are the same color, they should be in the same category).



$$E_s(L) = \sum_{\text{neighbors } (p,q)} w_{pq} |L(p) - L(q)|$$

w_{pq} : similarity in intensity of p and q

Hand-drawn boxes containing the weights $w_{pq} = 0.1$ and $w_{pq} = 10.0$, with blue arrows pointing to the highlighted pixels in the image above.

Handwritten notes in Hebrew: "אם הפיקסלים הם באותו צבע אז הם צריכים להיות באותו קטגוריה" (If the pixels are the same color, they should be in the same category).


Binary segmentation as energy minimization

$$E(L) = E_d(L) + \lambda E_s(L)$$

- ▶ For this problem, we can easily find the global minimum!
- ▶ Use max flow / min cut algorithm




Graph min cut problem



- ▶ Given a weighted graph G with source and sink nodes (s and t), partition the nodes into two sets, S and T such that the sum of edge weights spanning the partition is minimized
 - ▶ and $s \in S$ and $t \in T$


Segmentation by min cut



Graph

- ▶ node for each pixel, link between adjacent pixels
 - ▶ specify a few pixels as foreground and background
 - ▶ create an infinite cost link from each bg pixel to the t node
 - ▶ create an infinite cost link from each fg pixel to the s node
 - ▶ create finite cost links from s and t to each other node
 - ▶ compute min cut that separates s from t
-
- ▶ The min-cut max-flow theorem [Ford and Fulkerson 1956]

Segmentation by min cut



- ▶ The partitions S and T formed by the min cut give the optimal foreground and background segmentation
- ▶ I.e., the resulting labels minimize

$$E(d) = E_d(d) + \lambda E_s(d)$$

GrabCut segmentation

1. Define graph

- ▶ usually 4-connected or 8-connected
 - ▶ Divide diagonal potentials by $\sqrt{2}$

2. Define unary potentials

- ▶ Color histogram or mixture of Gaussians for background and foreground

$$\text{unary_potential}(x) = -\log \left(\frac{P(c(x); \theta_{\text{foreground}})}{P(c(x); \theta_{\text{background}})} \right)$$

3. Define pairwise potentials

$$\text{edge_potential}(x, y) = k_1 + k_2 \exp \left\{ \frac{-\|c(x) - c(y)\|^2}{2\sigma^2} \right\}$$

4. Apply graph cuts

5. Return to 2, using current labels to compute foreground, background models



GrabCut

אובייקט של המבוקש
מסירה על ידי
user

שינוי לאלמנט
המבוקש

Grabcut [[Rother et al., SIGGRAPH 2004](#)]




GrabCut

- Implemented in MS office Let's try it



Reported results

Easier examples



More difficult Examples

Camouflage &



Fine structure



Harder Case



Graph cuts with multiple labels

▶ Alpha expansion

Repeat until no change

For $\alpha = 1..M$

Assign each pixel to current label or α (2-class graphcut)

▶ Achieves “strong” local minimum

▶ Alpha-beta swap

Repeat until no change

For $\alpha = 1..M$, $\beta = 1..M$ (except α)

Re-assign all pixels currently labeled as α or β to one of those two labels while keeping all other pixels fixed



Other application: synthesis

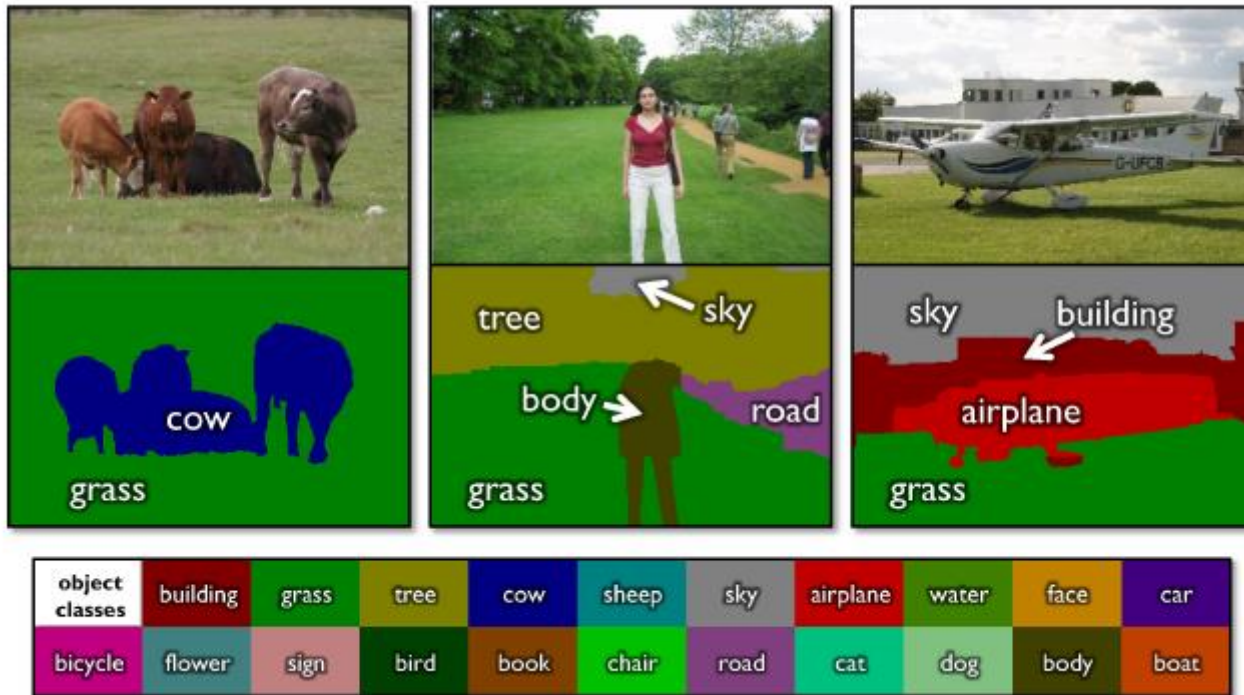


Handwritten notes:

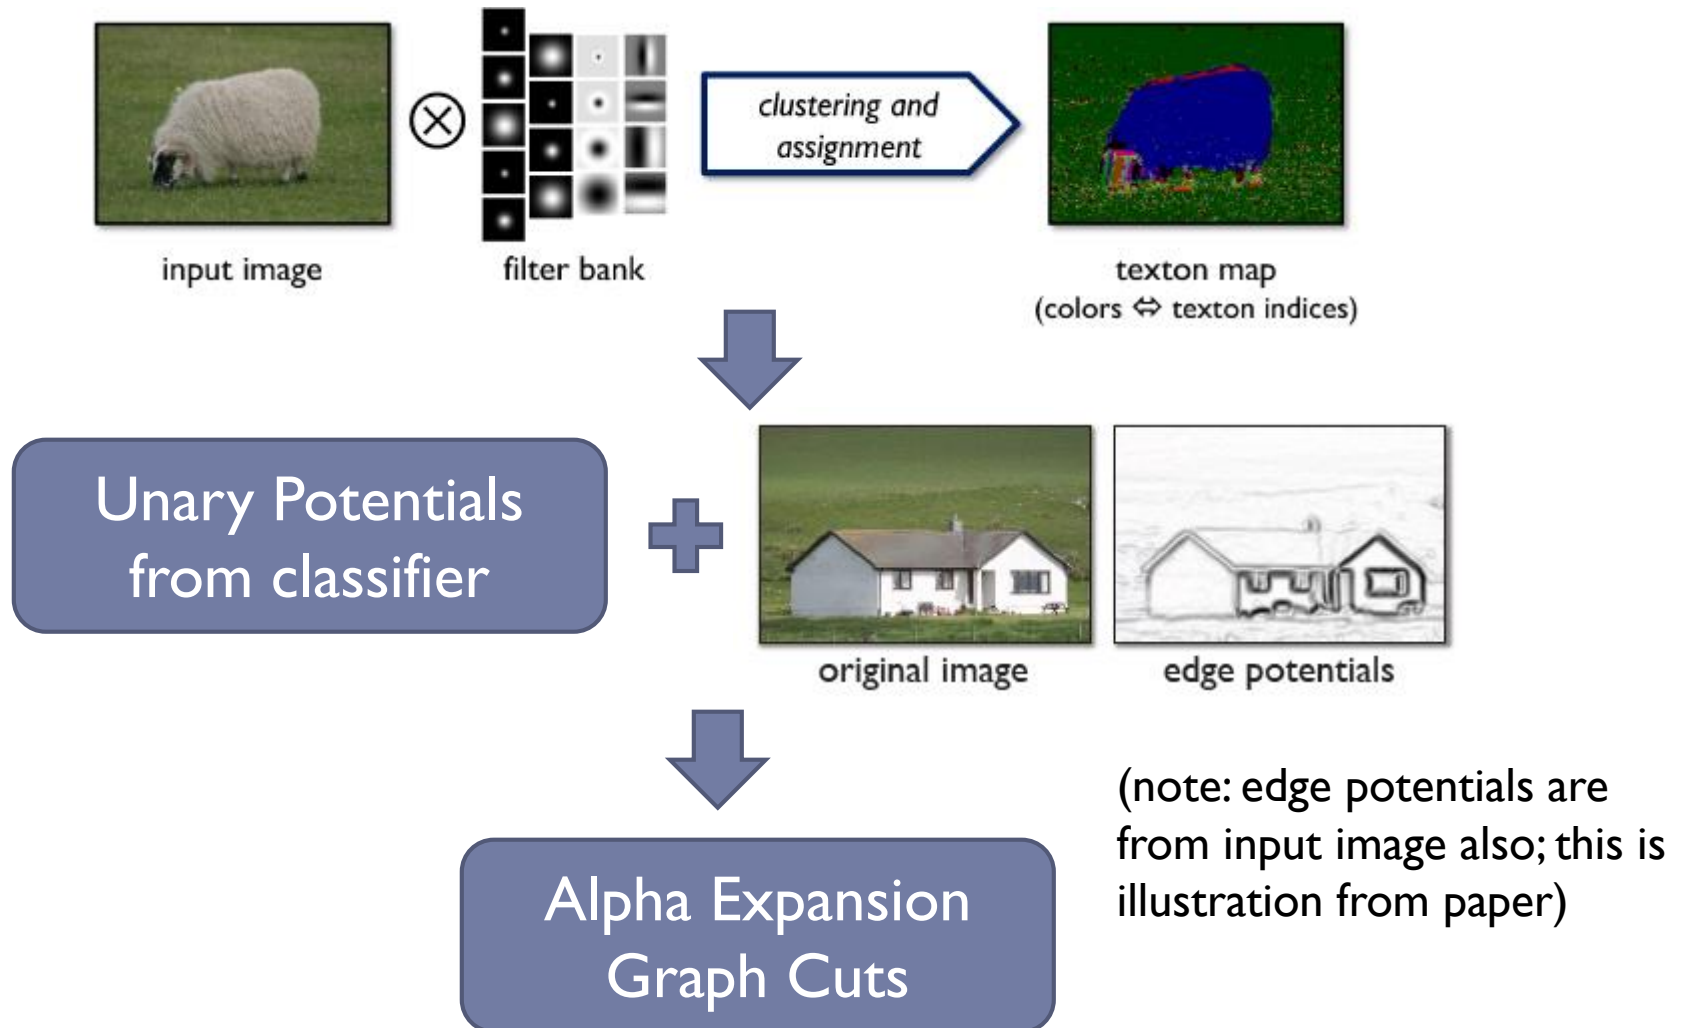
- 215W
- 712
- 715W
- Graph
- 5/5



Using graph cuts for recognition



Using graph cuts for recognition



Summary: MRF and graph-cuts

▶ Pros:

- ▶ Powerful, based on probabilistic model (MRF)
- ▶ Applicable to a wide range of problems
- ▶ Very efficient algorithms available for many problems
- ▶ Becoming a standard for segmentation

▶ Cons

- ▶ Graph-cuts can only solve a limited class of problems:
 - ▶ Sub-modular energy functions
 - ▶ Can only capture part of the power of MRF
- ▶ Only approximate solutions available for multi-label case

Graph cuts: Pros and Cons

► Pros

- Very fast inference
- Can incorporate data likelihoods and priors
- Applies to a wide range of problems (stereo, image labeling, recognition)

► Cons

- Not always applicable (associative only)
- Need unary terms (not used for bottom-up segmentation, for example)

► Use whenever applicable

use whenever applicable

Further reading and resources

▶ Graph cuts

- ▶ <http://www.cs.cornell.edu/~rdz/graphcuts.html>
- ▶ Classic paper: [What Energy Functions can be Minimized via Graph Cuts?](#) (Kolmogorov and Zabih, ECCV '02/PAMI '04)

▶ Belief propagation

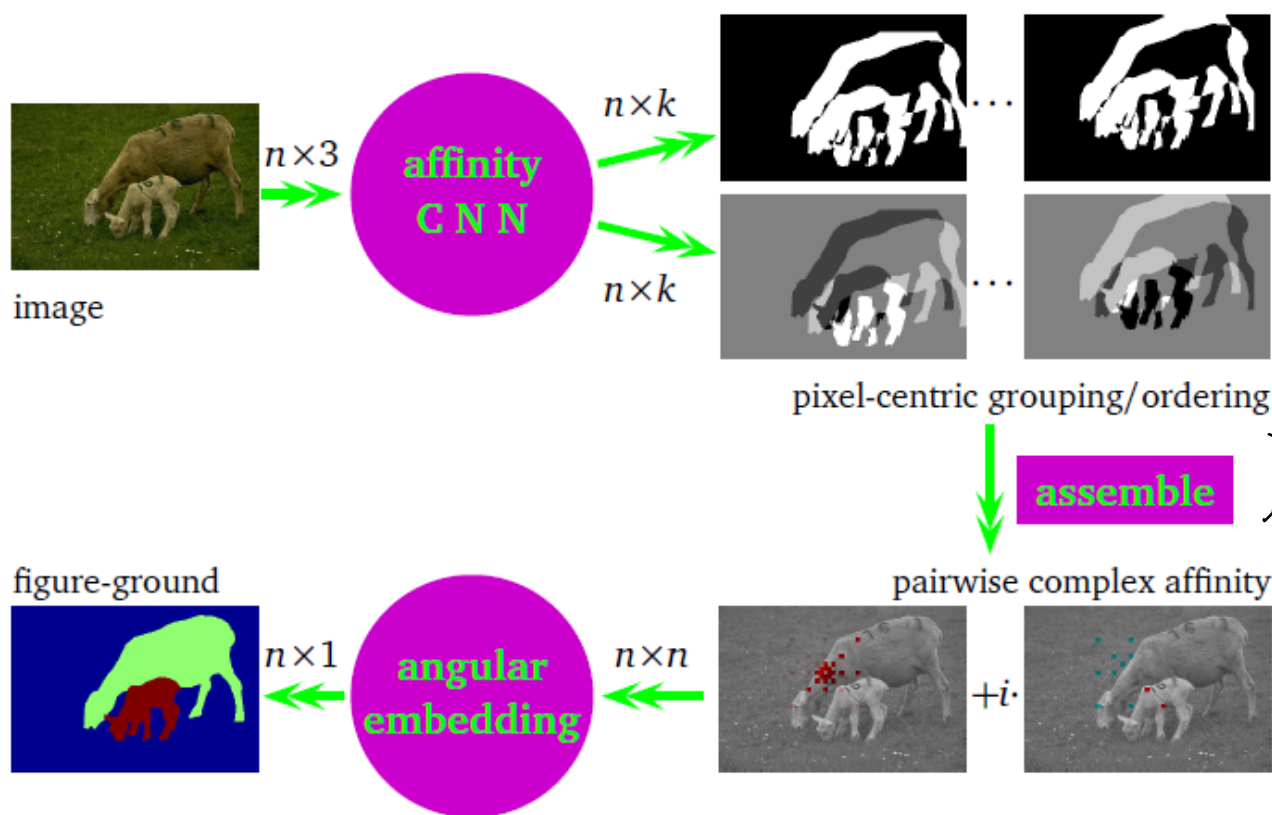
Yedidia, J.S.; Freeman, W.T.; Weiss, Y., "Understanding Belief Propagation and Its Generalizations", Technical Report, 2001:

<http://www.merl.com/publications/TR2001-022/>



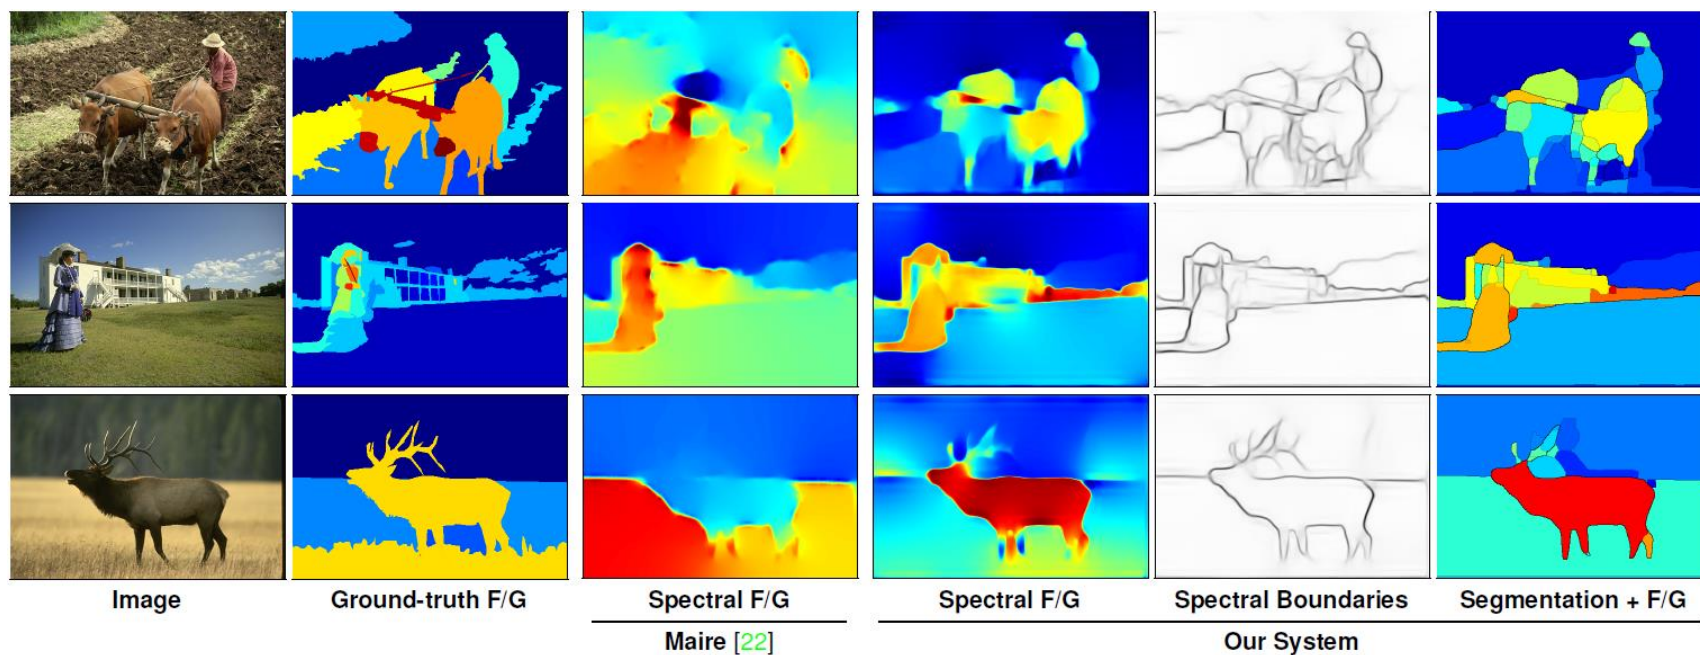
CNN for affinity learning

► Maire and Yu, 2015



CNN for affinity learning

► Maire and Yu, 2015



End – image segmentation part 2

Now you know how it works