

From points to regions

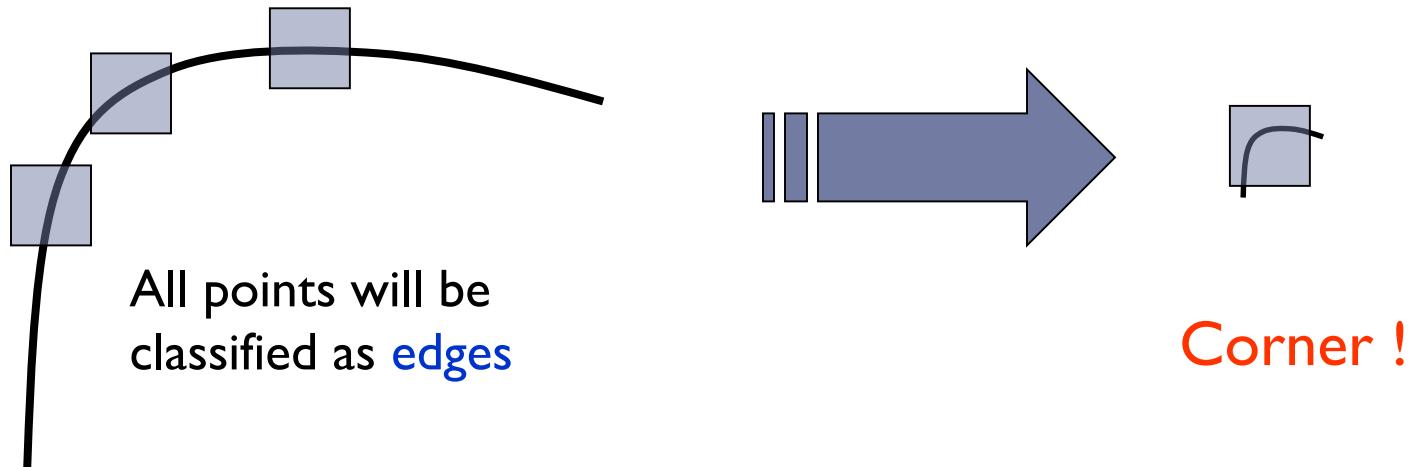
Lihai Zelnik-Manor, Computer Vision

Today

- ▶ Local invariant features
 - ▶ Motivation
 - ▶ Requirements, invariances
- ▶ Keypoint localization
 - ▶ Harris corner detector
 - ▶ Hessian detector
- ▶ Scale invariant region selection
 - ▶ Automatic scale selection
 - ▶ Laplacian-of-Gaussian detector
 - ▶ Difference-of-Gaussian detector

Harris Detector: Properties

- ▶ Is it invariant to image scale?



No: Not invariant to image scale!

From points to regions

- ▶ The Harris and Hessian operators define interest points
 - ▶ Precise localization
 - ▶ High repeatability

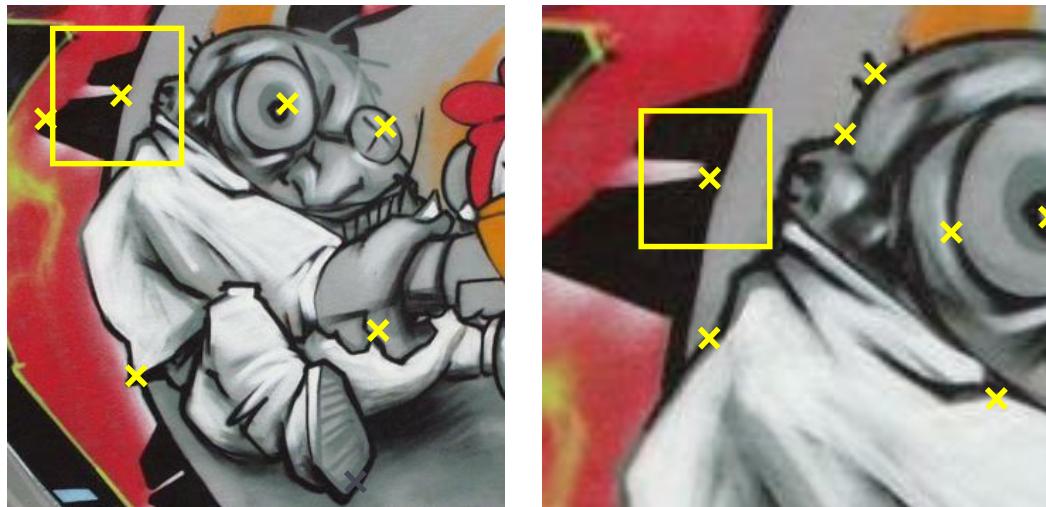
מעוניינים בഗלי סקאלה, שיגיד לנו מה הסקלה שכדי להסתכל עליה



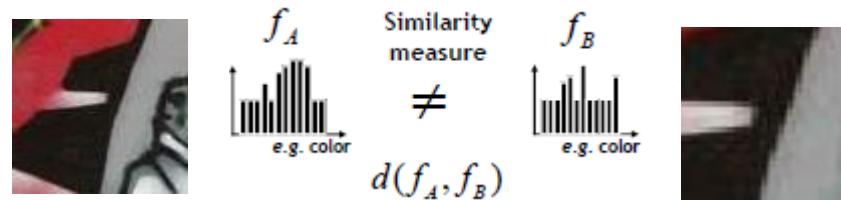
- ▶ In order to compare (and match) those points, we need to compute a descriptor over their local region
 - ▶ How can we define such a region in a scale invariant manner?
 - ▶ **How can we detect scale invariant interest region?**

Naïve approach: Exhaustive search

- ▶ Multi-scale procedure
 - ▶ Compare descriptors while varying patch size

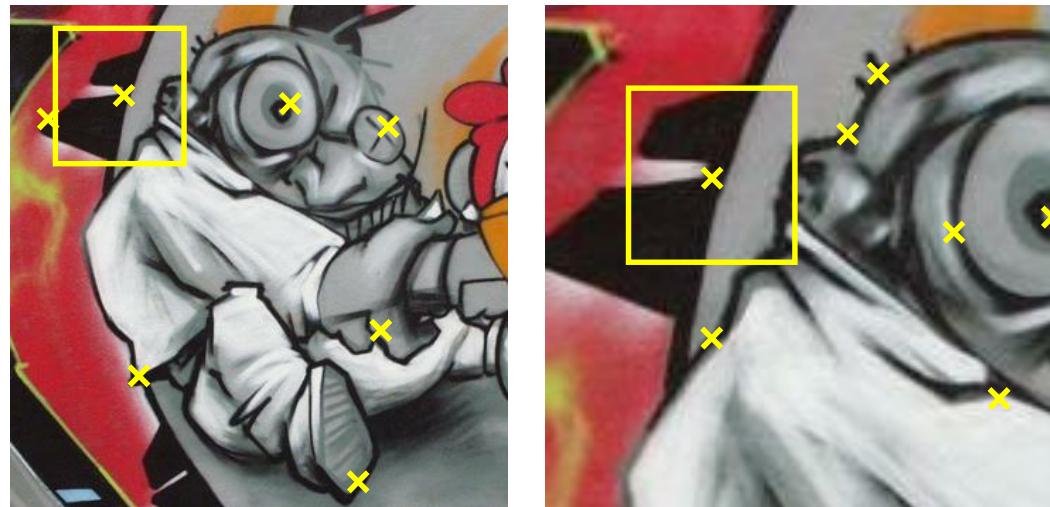


הריובעים באותו גודל אך התמונה לא באותו זום

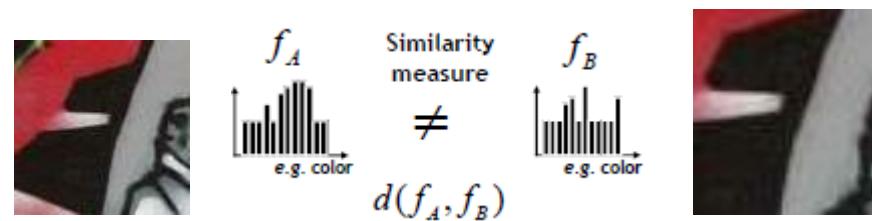


Naïve approach: Exhaustive search

- ▶ Multi-scale procedure
 - ▶ Compare descriptors while varying patch size

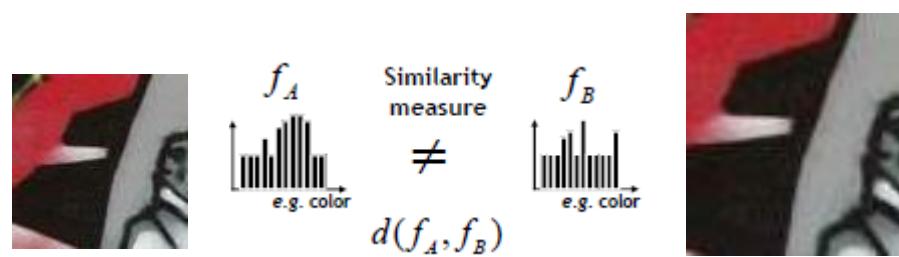
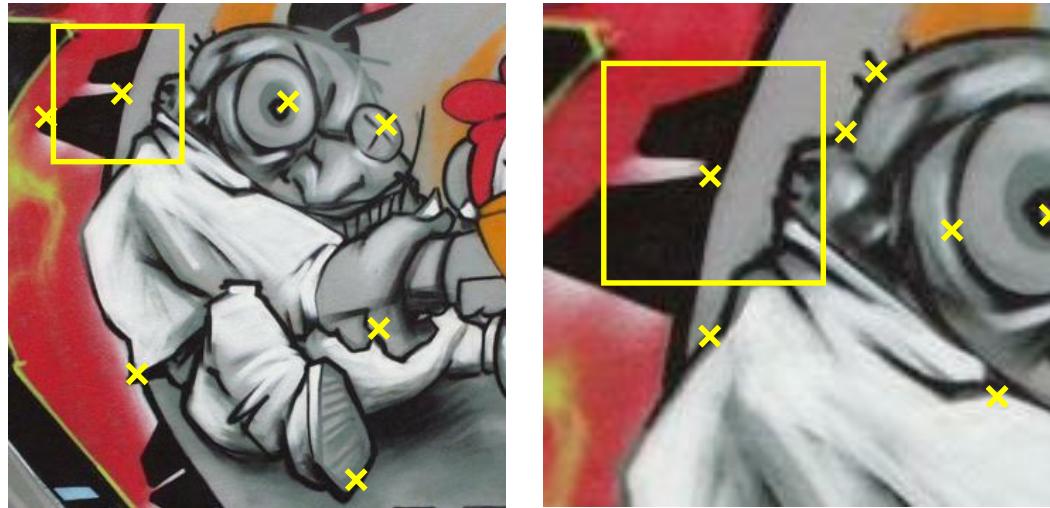


לכן נתחילה להגדיל את הריבוע בצד
ימין ובאיזשהו שלב יהיה לנו שווין



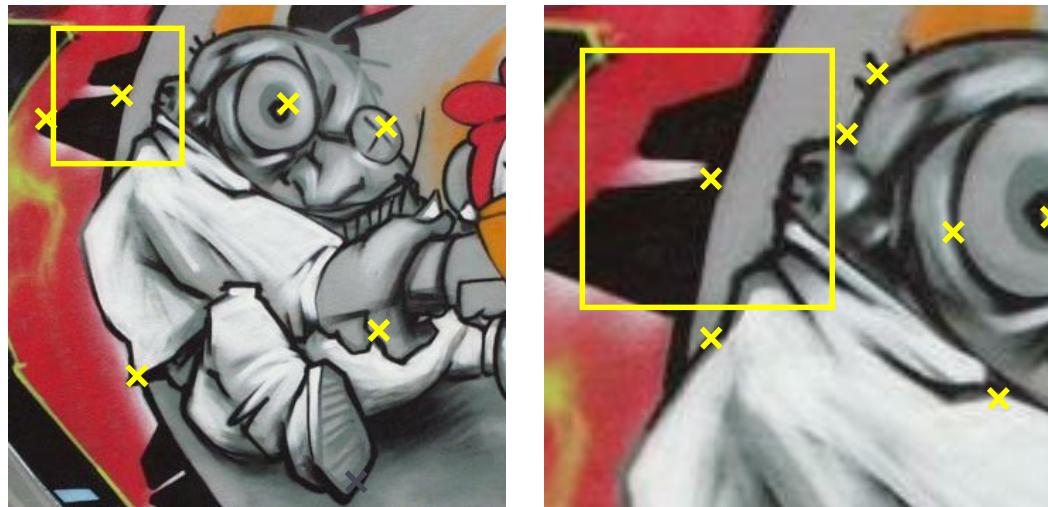
Naïve approach: Exhaustive search

- ▶ Multi-scale procedure
 - ▶ Compare descriptors while varying patch size

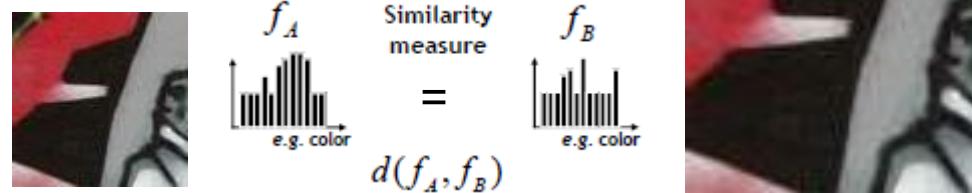


Naïve approach: Exhaustive search

- ▶ Multi-scale procedure
 - ▶ Compare descriptors while varying patch size



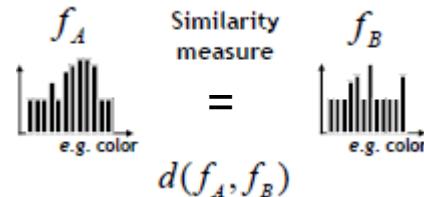
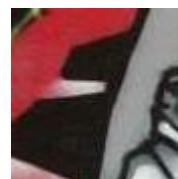
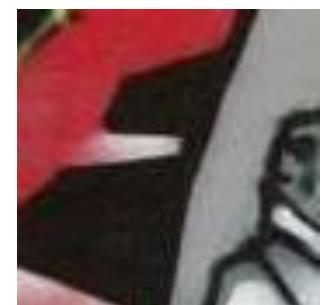
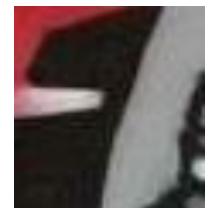
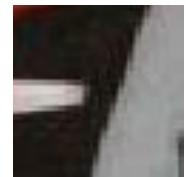
זו השיטה הנאייבית
לקחים אוסף של סקלולות
ורצים עליהם ,



Naïve approach: Exhaustive search

- ▶ Compare descriptors while varying patch size
 - ▶ Computationally inefficient
 - ▶ Prohibitive for retrieval in large databases
 - ▶ Prohibitive for recognition

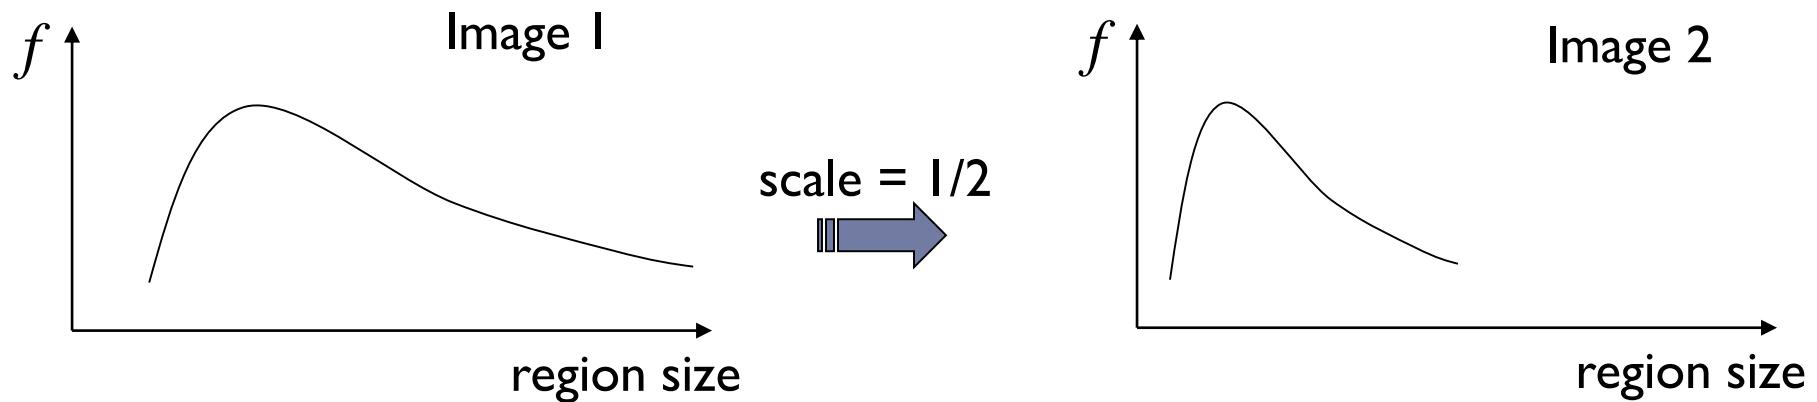
בשיטת הנאייבית נוכל להתאים דברים שלא בהכרח מתאימים



Automatic scale selection

- ▶ Solution:
 - ▶ Design a function on the region, which is “scale invariant” (*the same for corresponding regions, even if they are at different scales*)
Example: average intensity. For corresponding regions (even of different sizes) it will be the same.
 - ▶ For a point in one image, we can consider it as a function of region size (patch width)

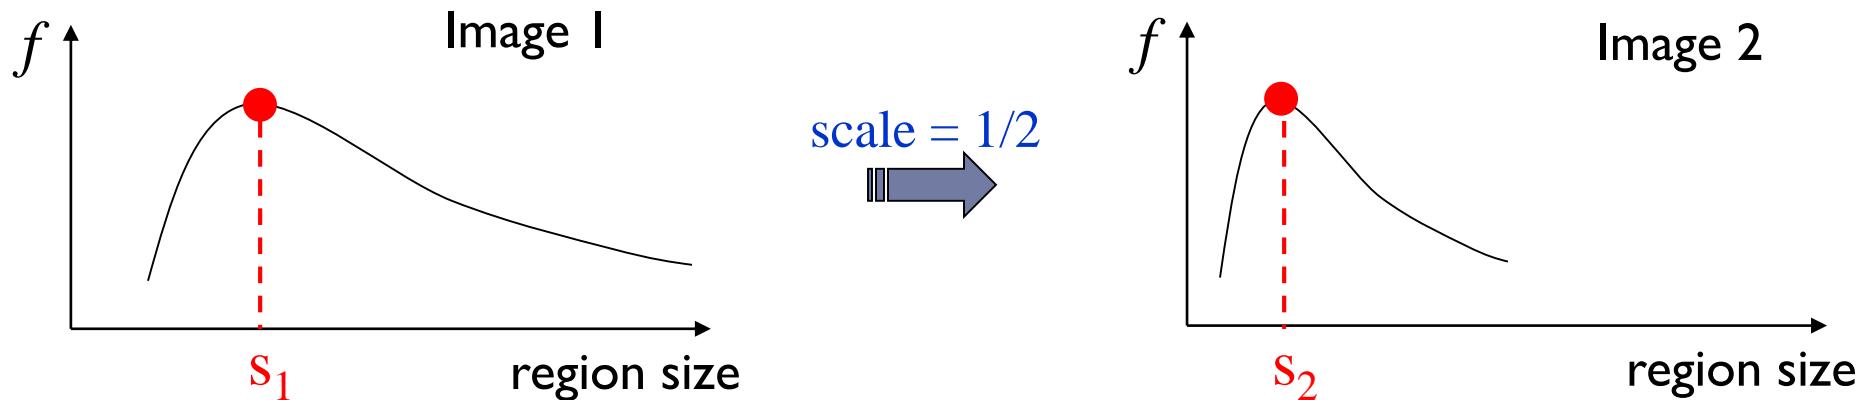
לפונקציה הזאת יש אותה צורה מותאמת
אבל עם אותה מקסימום



Automatic scale selection

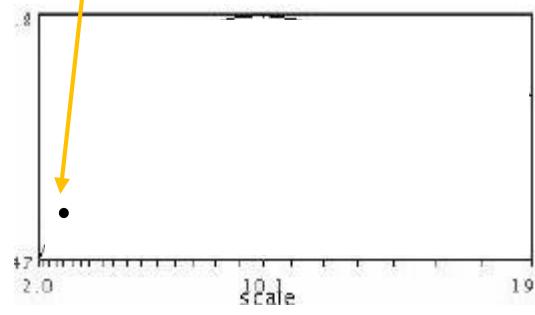
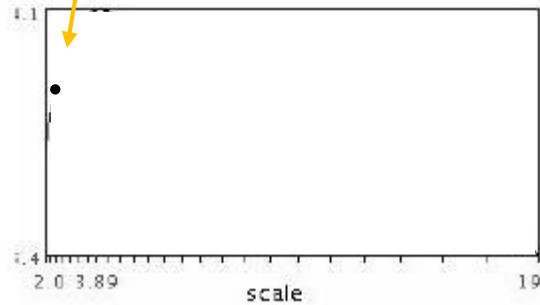
- ▶ Common approach:
 - ▶ Take a local maximum of this function
 - ▶ Observation: region size, for which the maximum is achieved, should be *invariant* to image scale.

Important: this scale invariant region size is found in each image independently!



Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)

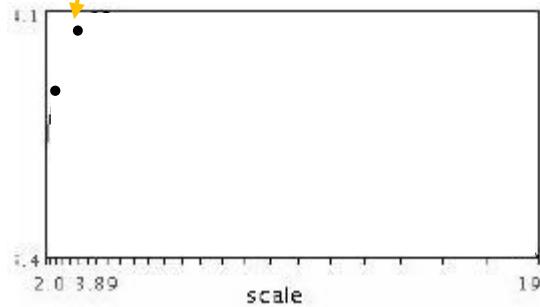


$$f(I_{i_1 \dots i_m}(x, \sigma))$$

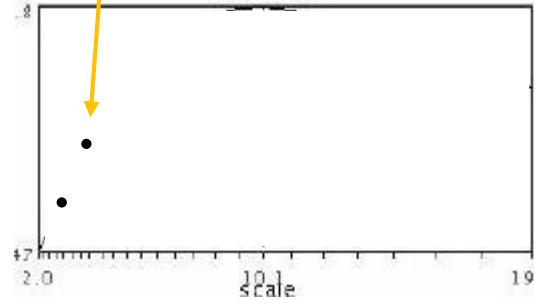
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)



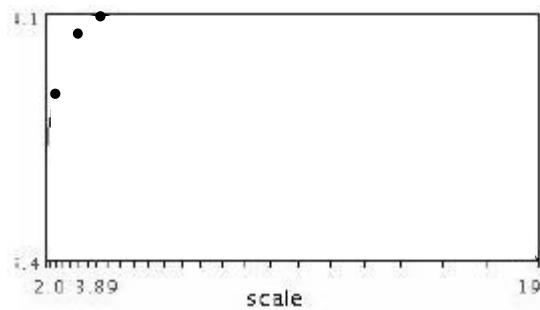
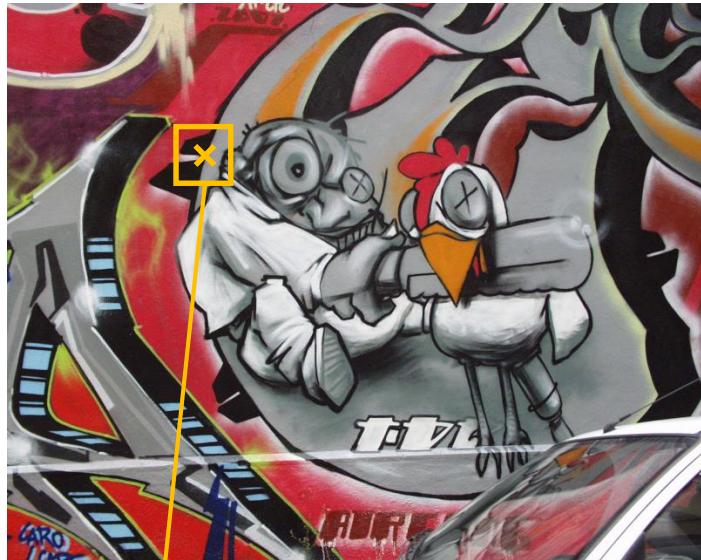
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



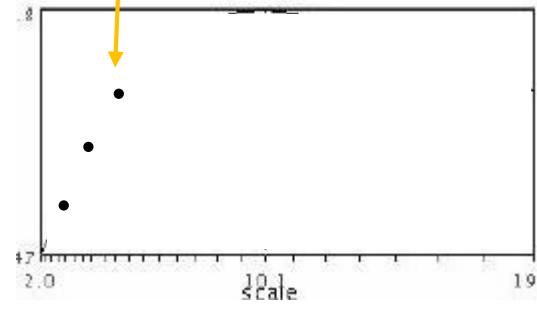
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)



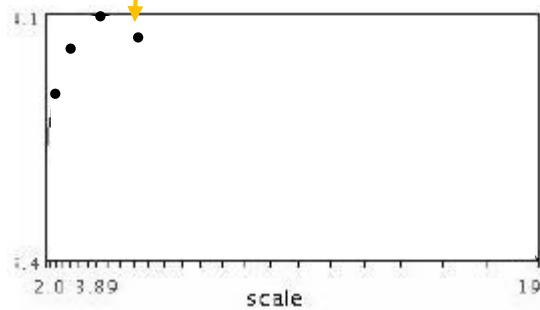
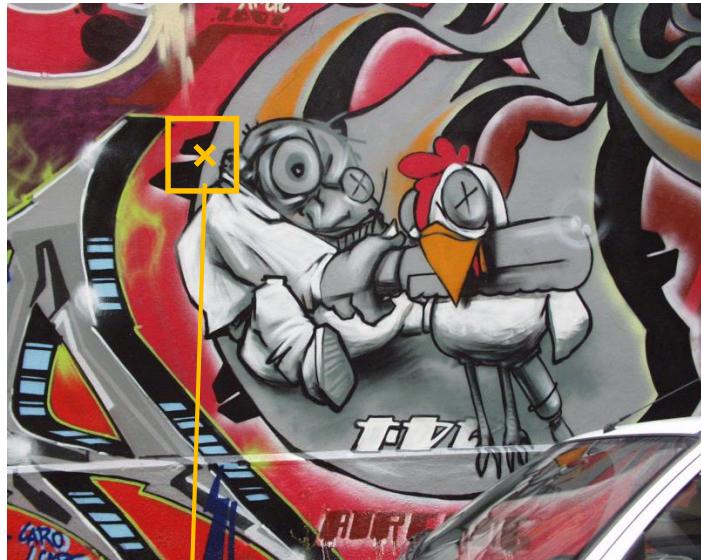
$$f(I_{i_1\dots i_m}(x, \sigma))$$



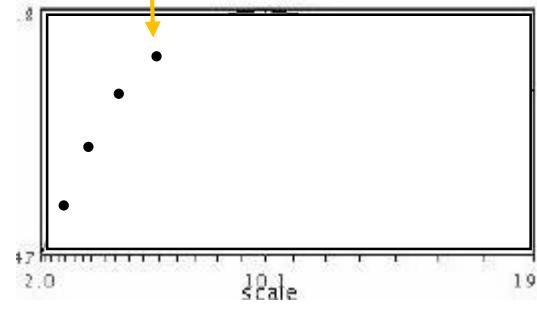
$$f(I_{i_1\dots i_m}(x', \sigma))$$

Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)



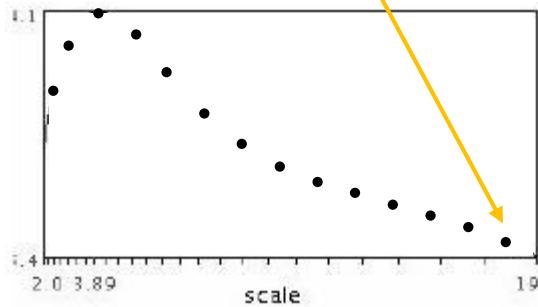
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



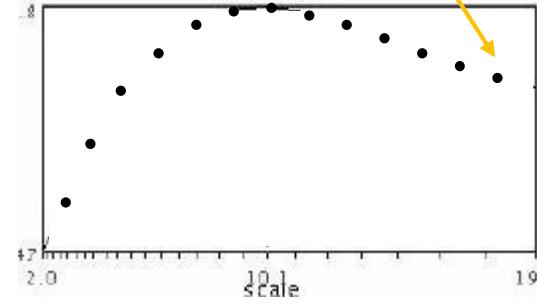
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)



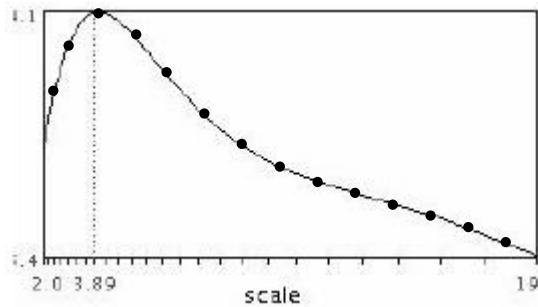
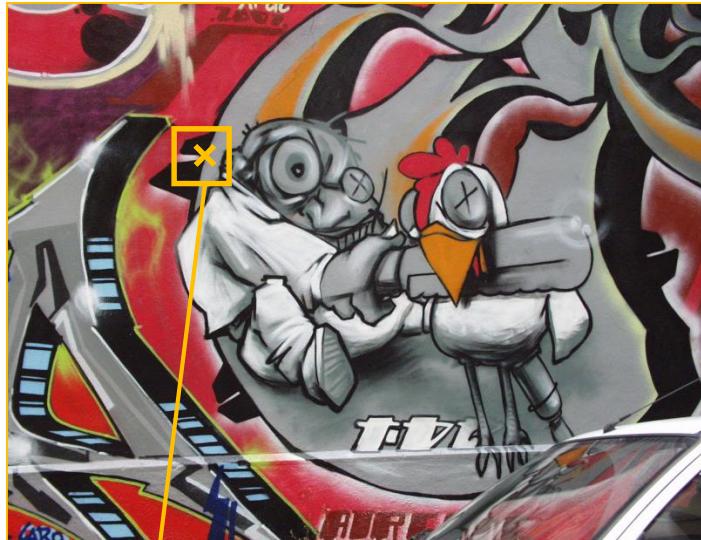
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



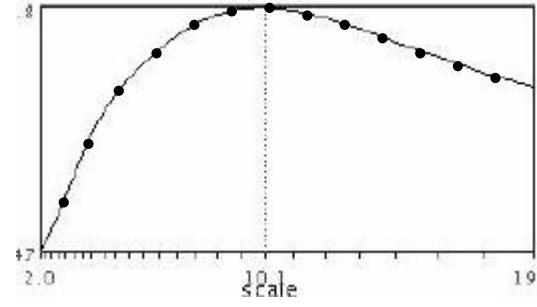
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- ▶ Function responses for increasing scale (scale signature)



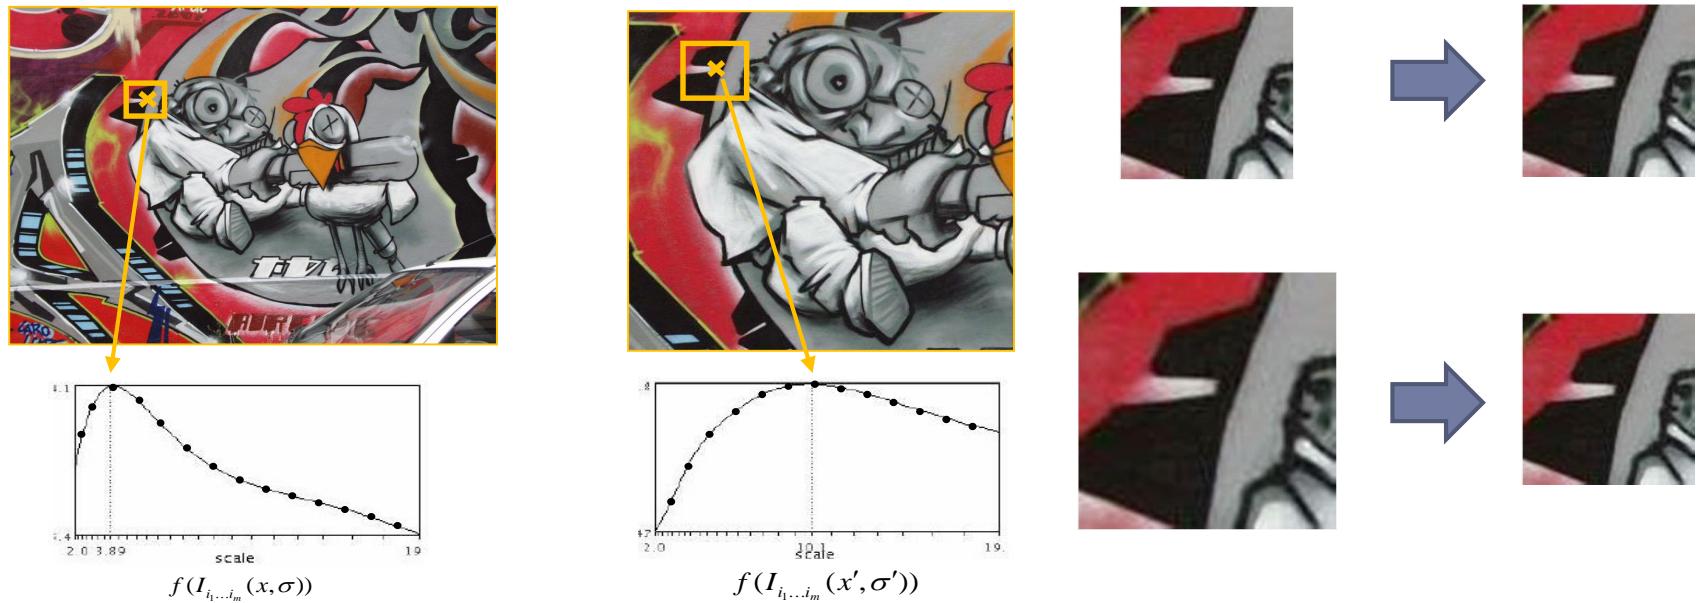
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma'))$$

Automatic Scale Selection

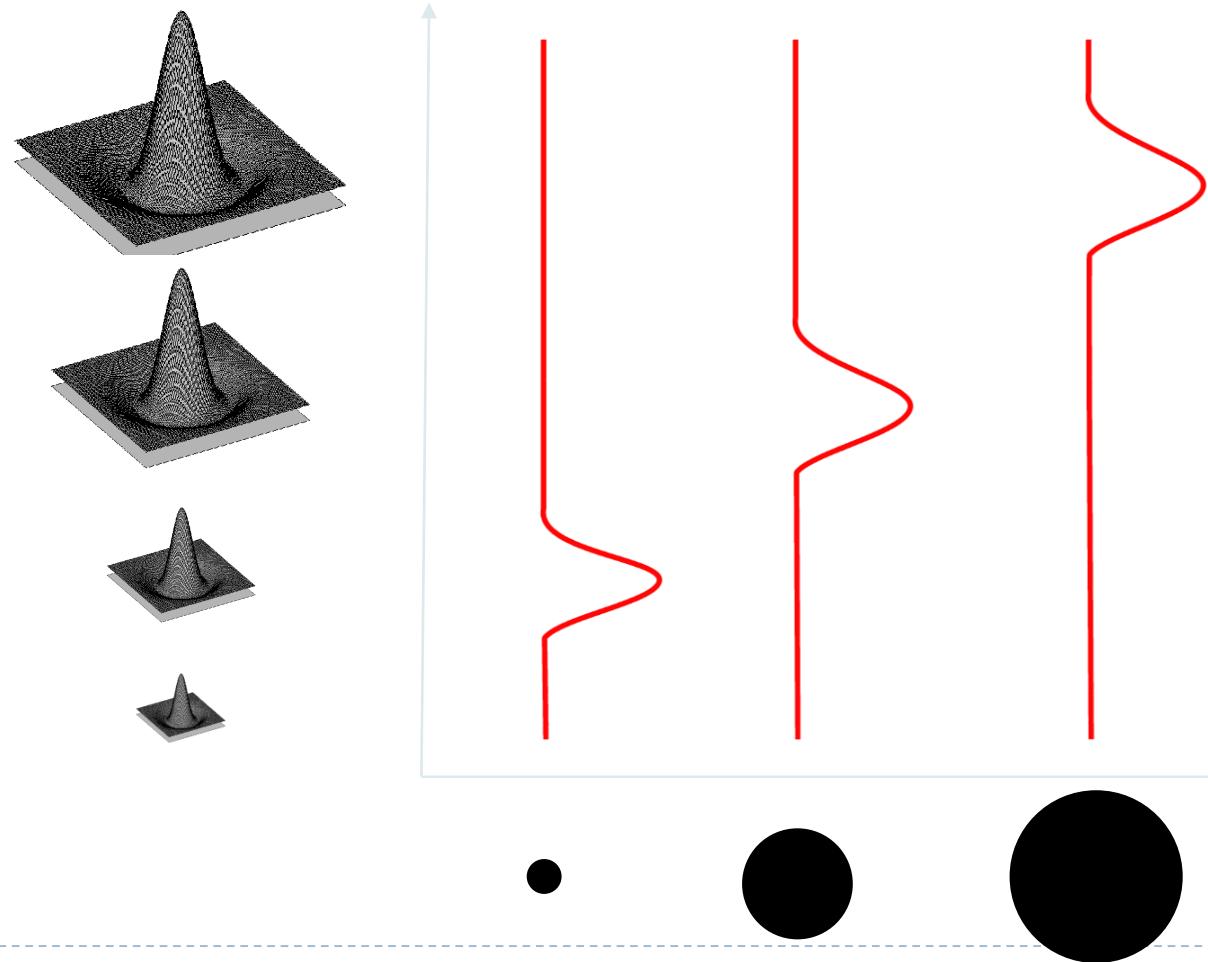
- ▶ Normalize: rescale selected regions to a fixed size



התוצאה היא הסקלוט השומן לריבועית שאינו רוצה להשווות אותו

A useful scale “signature” function

- ▶ Laplacian-of-Gaussian = “blob” detector 



Harris – Laplace [Mikolajczyk'01]

1. Keypoint detection:

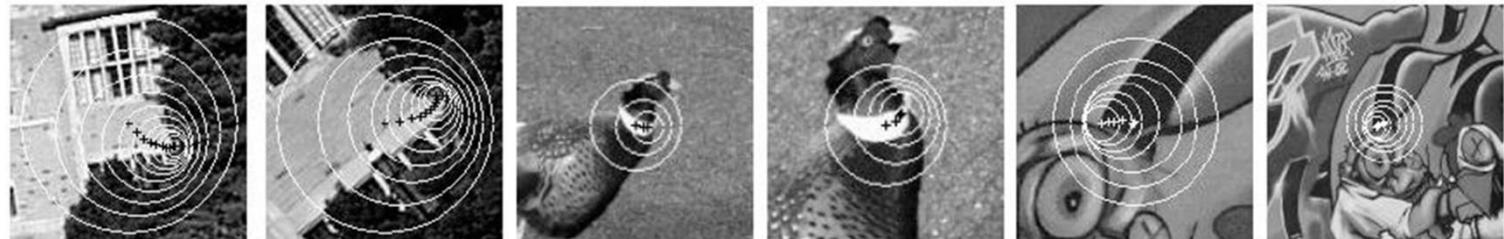
- ▶ Multi-scale Harris corner detection

2. Scale selection

- ▶ Scale selection based on Laplacian signature

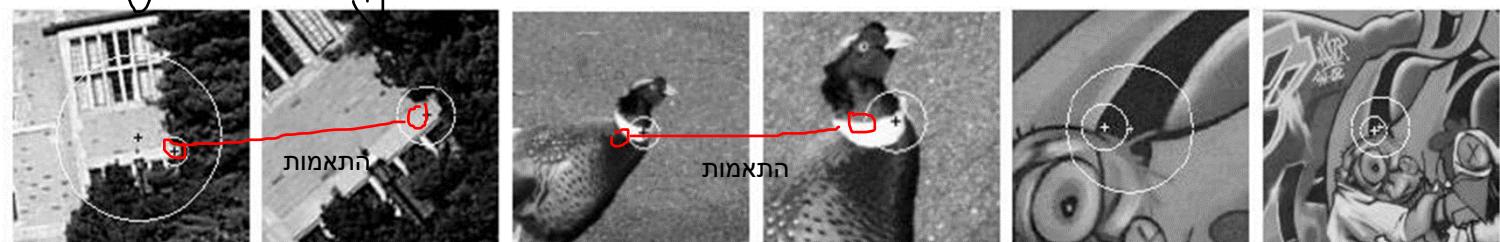
מסתכל על התמונה שלי בסביבות בגודל שונה
ובכל סביבה אני מוצא נקודות הריס

Harris
points



אקסים שונים
שהתקבלו מהאריס
העיגולים מיצגים את הסקלות
השונות בתמונה

Harris-Laplace
points



אחרי הוספה של הלפלסיאן אני
נשאר עם שתי נקודות של הריס

מה שאני מקבל לאחר הפעלת הלפלסיאן הספציפי
הנקודות שנשארו זה

התרומה של הלפלסיאן זה בעצם להגיד לי מה הסקלות הנכונות
להסתכל על הסביבה שבחורת

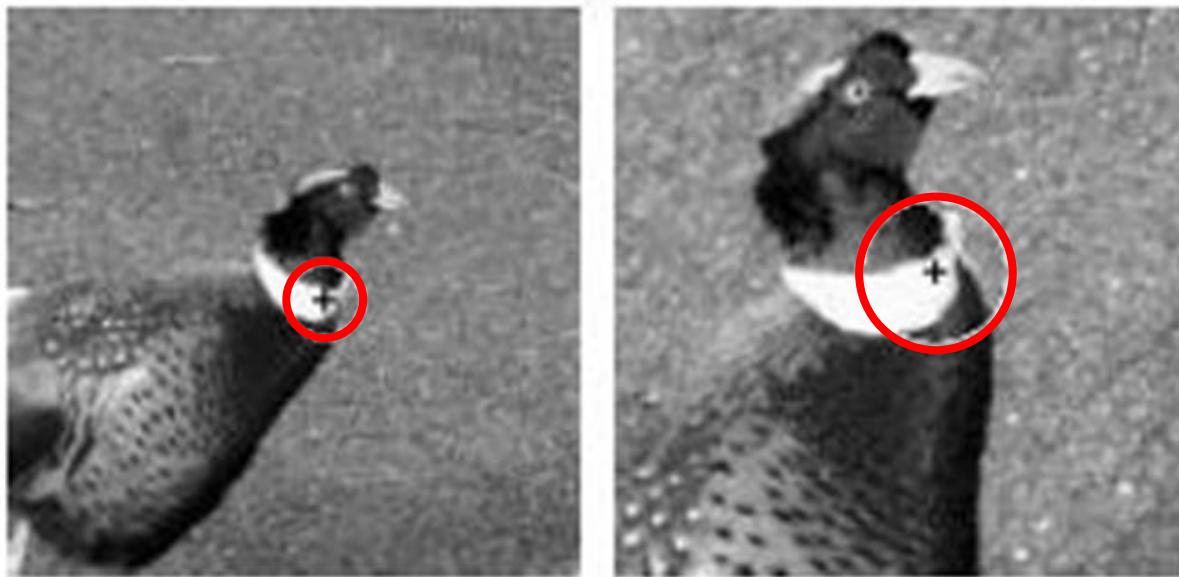
Harris – Laplace [Mikolajczyk'01]

1. Keypoint detection:

- ▶ Multi-scale Harris corner detection

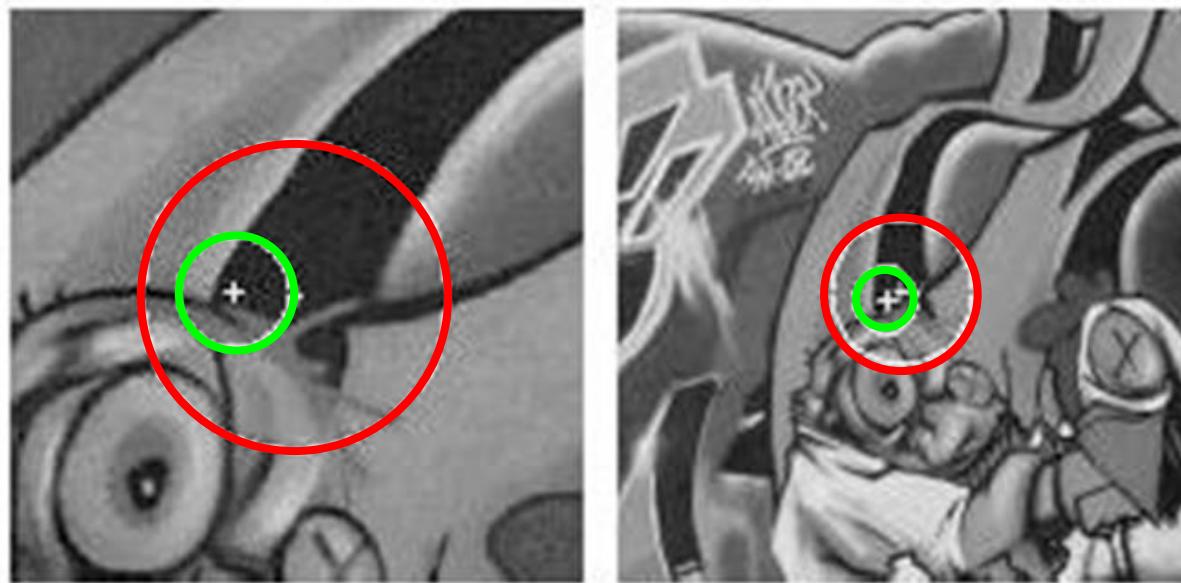
2. Scale selection

- ▶ Scale selection based on Laplacian signature



Harris – Laplace [Mikolajczyk'01]

1. Keypoint detection:
 - ▶ Multi-scale Harris corner detection
2. Scale selection
 - ▶ Scale selection based on Laplacian signature



Today

- ▶ Local invariant features
 - ▶ Motivation
 - ▶ Requirements, invariances
- ▶ Keypoint localization
 - ▶ Harris corner detector
 - ▶ Hessian detector
- ▶ Scale invariant region selection
 - ▶ Automatic scale selection
 - ▶ Laplacian-of-Gaussian detector
 - ▶ Difference-of-Gaussian detector

blob deduction

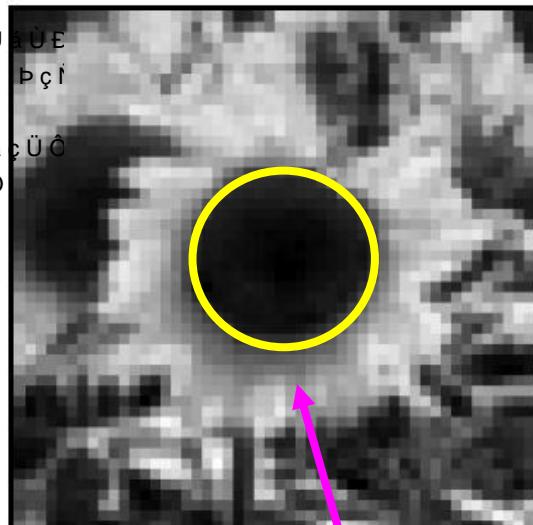
גודל של הבלוב בתמונה

יגדר לי את הסקיל שבו אני רוצה להסתכל על האיזור הספציפי

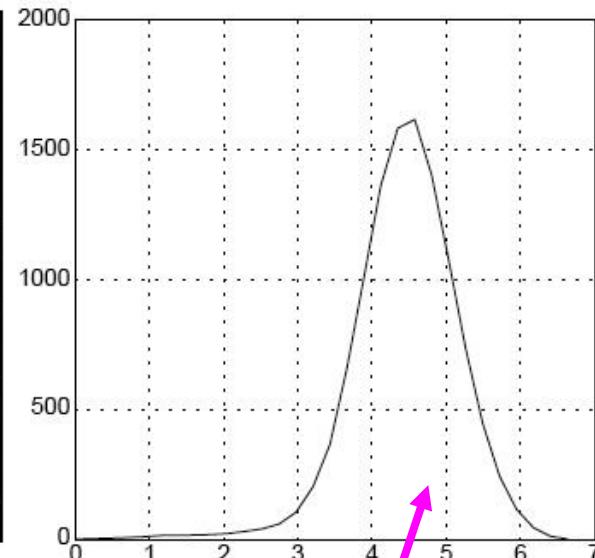
Characteristic scale

- ▶ We define the *characteristic scale* as the scale that produces peak of Laplacian response

ĐàÙ àÔÑè áÝ ÚþÔ ÜäÜáUE
ÕĐàÙ èÕĐÔ ĐÙäÔ ĐàÙ þçí
ôÒÓÖÜÔ ÑÙÖèè
ÕéþÔ ĐàÙ çÕÑâ Đê ÔáçÜÔ
éÑÔ ĐàÙ èÕæÔ ÜáÑÖÔ



characteristic scale

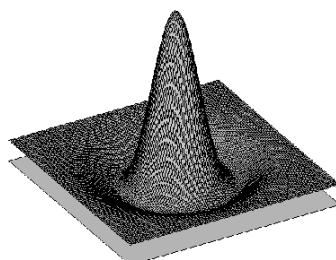


T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)
International Journal of Computer Vision 30 (2): pp 77--116.

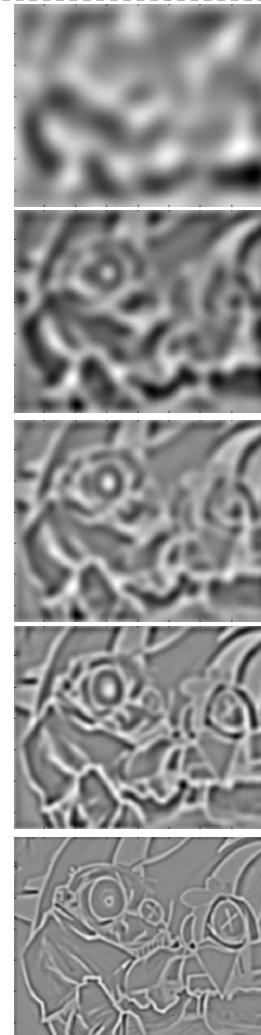
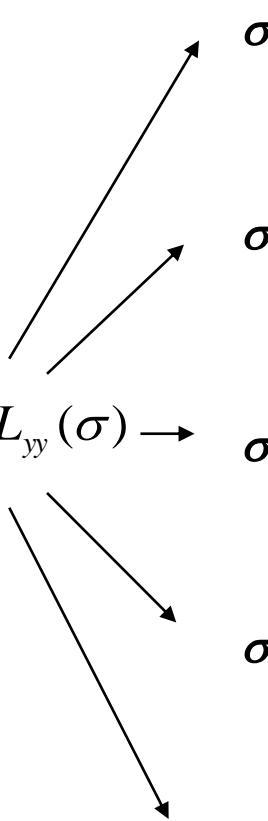
Laplacian-of-Gaussian (LoG)

► Interest points:

Local maxima in scale
space of Laplacian-of-
Gaussian



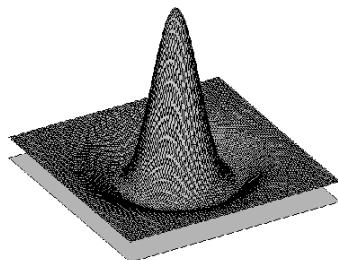
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow$$



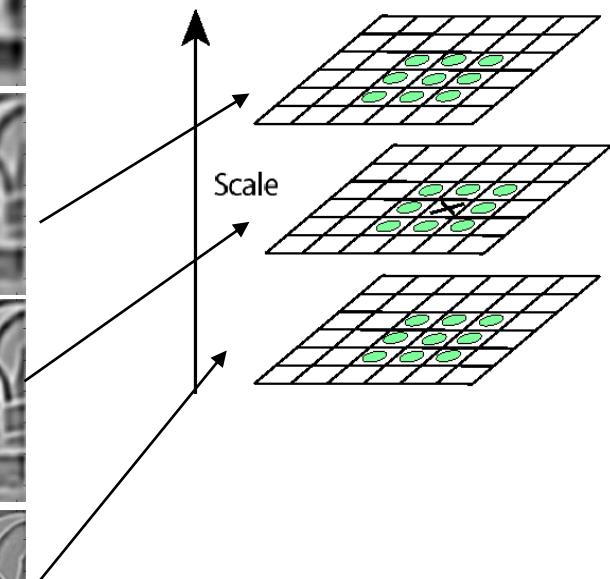
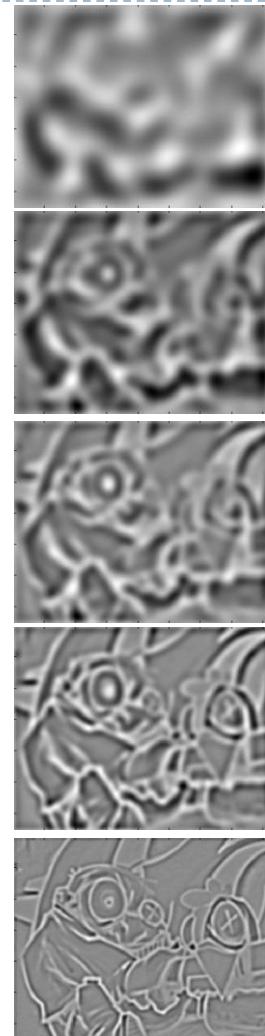
Laplacian-of-Gaussian (LoG)

► Interest points:

Local maxima in scale
space of Laplacian-of-
Gaussian



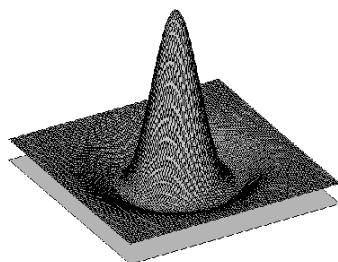
$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow$$
 σ^5
 σ^4
 σ^3
 σ^2
 σ



Laplacian-of-Gaussian (LoG)

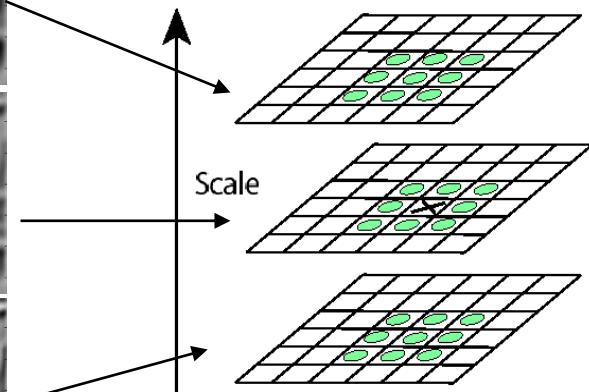
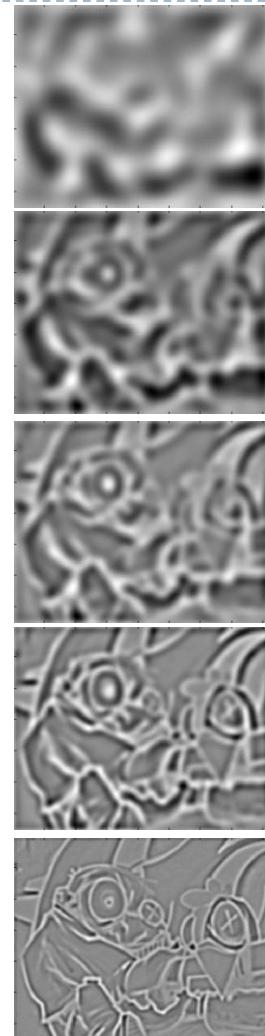
► Interest points:

Local maxima in scale
space of Laplacian-of-
Gaussian



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow$$

A diagram showing a 3D surface plot of a bell-shaped peak, representing the response of a LoG filter at a specific scale and position. Four arrows point from this diagram to five horizontal images of the same cartoon character, each at a different scale. The scales are labeled σ^5 , σ^4 , σ^3 , σ^2 , and σ from top to bottom.



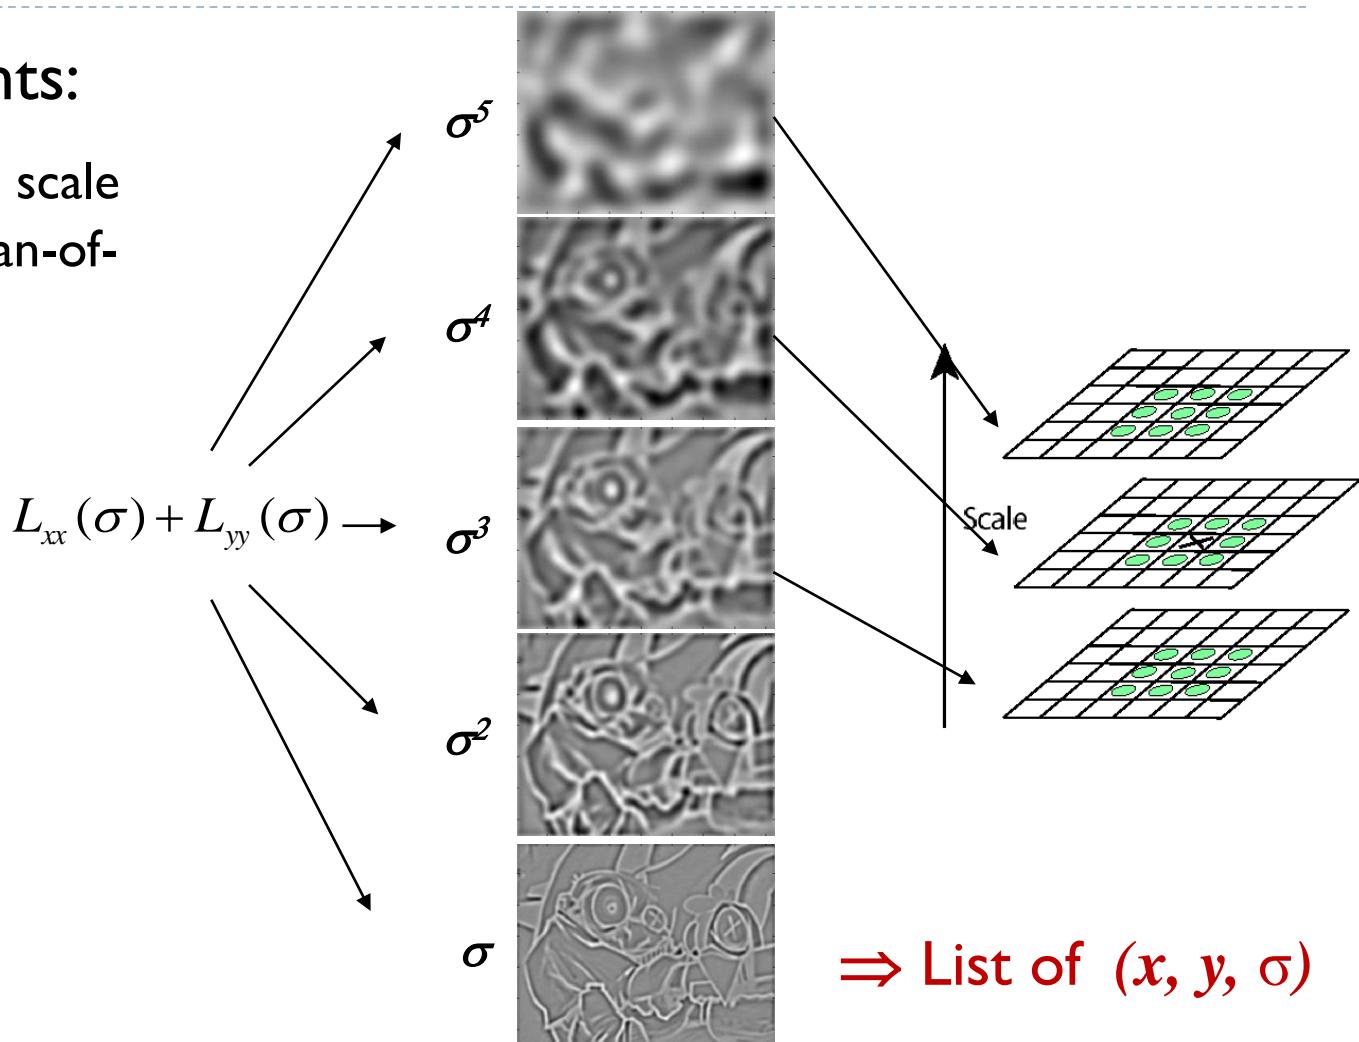
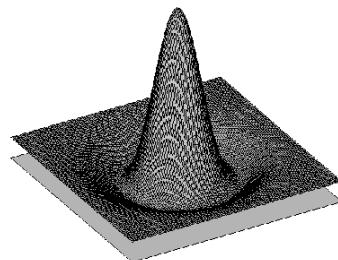
Laplacian-of-Gaussian (LoG)

► Interest points:

Local maxima in scale
space of Laplacian-of-
Gaussian



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow$$



LoG detector: workflow



אנו לוקחים הרבה הרבה סיגמות

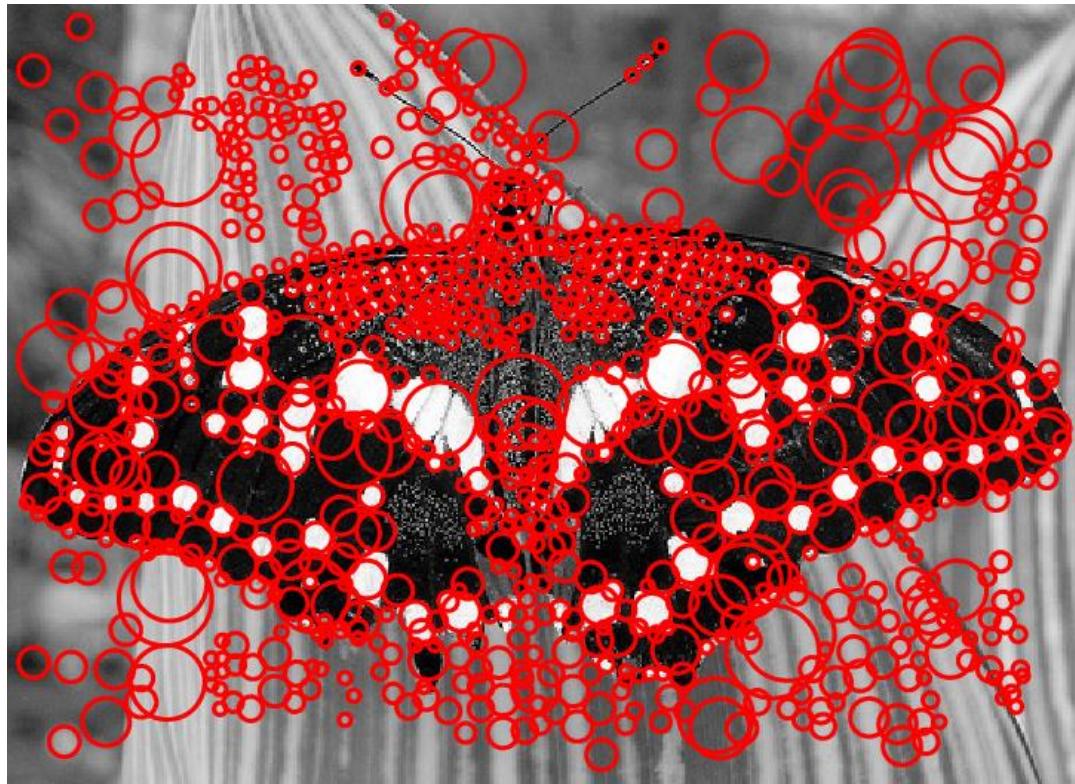
ולוקחים לוג דיטיקטור עם הרבה סיגמות שונות

LoG detector: workflow



$\sigma = 11.9912$

LoG detector: example result



כasher ani mastcal ul
kol hblowim batmuna achat
hblowim matkbelim bahtam le-sigmot

blow katan sigmat ktna
blow gadol sigma gdola

LoG technical details

- ▶ We can efficiently approximate the Laplacian with a difference of Gaussians:
- ▶ **Laplacian:**

$$L = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

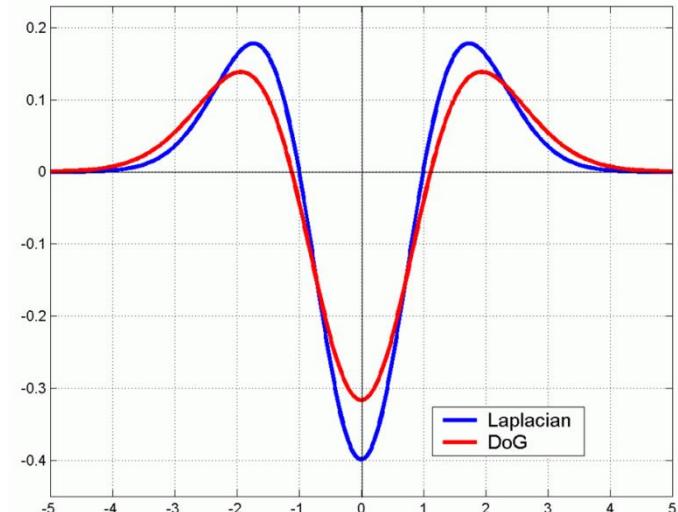
- ▶ **DoG = Difference of Gaussians:**

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma) \approx L$$

ענין של חישוב זריז
דימוי של לבסיאן

גauss גודל

גauss קטן

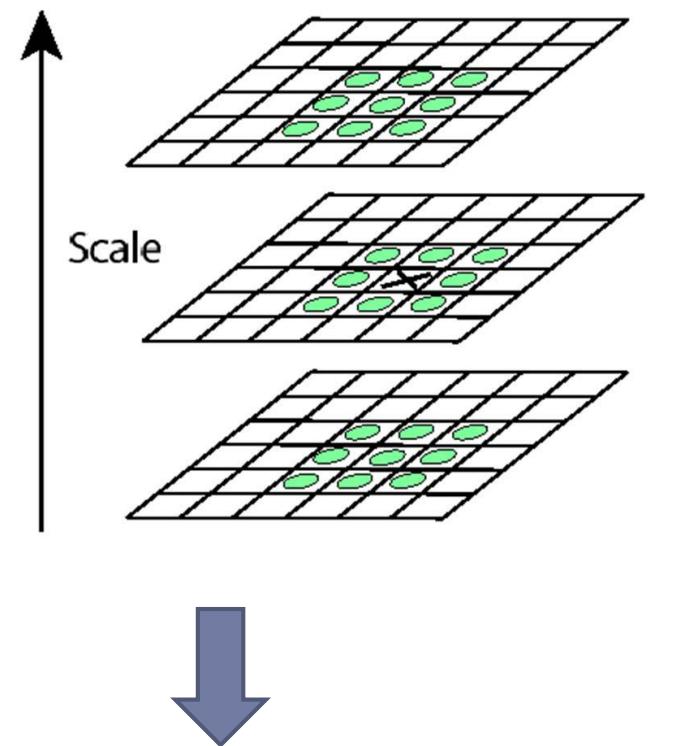


Today

- ▶ Local invariant features
 - ▶ Motivation
 - ▶ Requirements, invariances
- ▶ Keypoint localization
 - ▶ Harris corner detector
 - ▶ Hessian detector
- ▶ Scale invariant region selection
 - ▶ Automatic scale selection
 - ▶ Laplacian-of-Gaussian detector
 - ▶ **Difference-of-Gaussian detector**

Key point localization with DoG

- Detect maxima of difference-of-Gaussian (DoG) in scale space
- Then reject points with low contrast (threshold)
- Eliminate edge responses



Candidate keypoints: list of (x, y, σ)



DoG – efficient computation

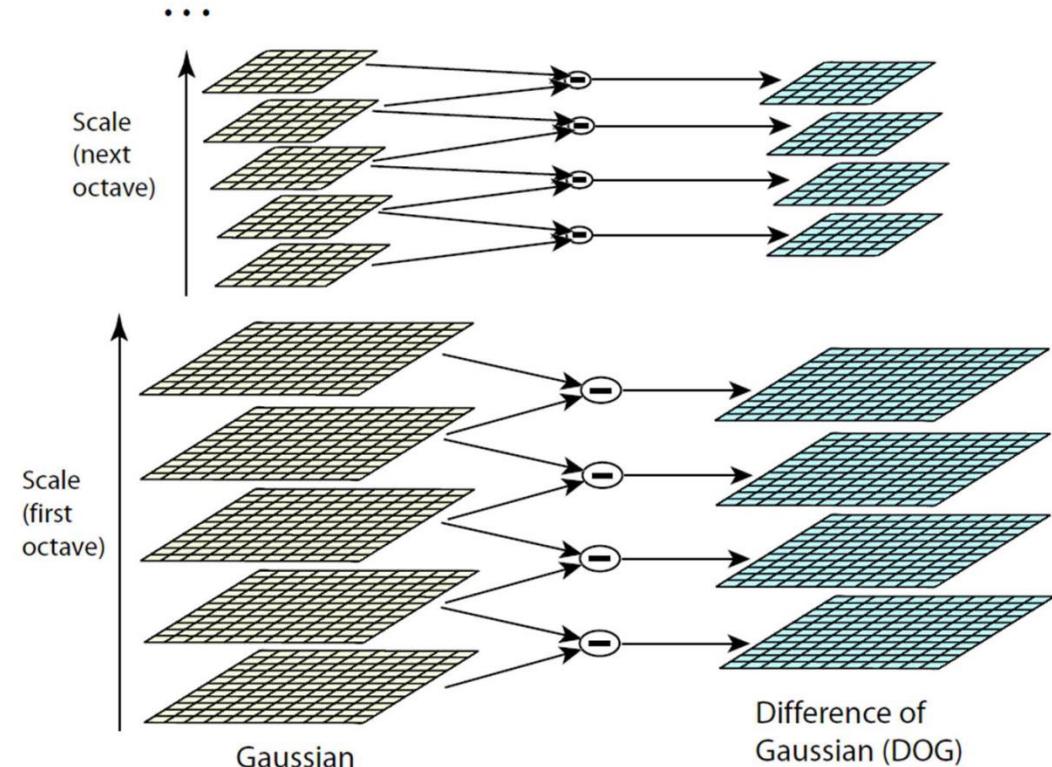
▶ Computation in Gaussian scale pyramid



Sampling with step $\sigma^4 = 2$



Original image

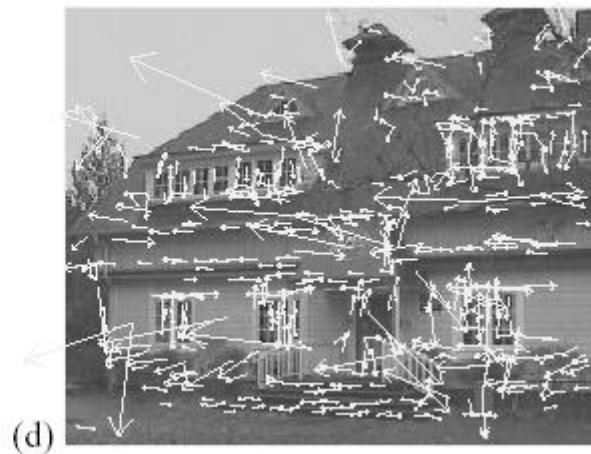
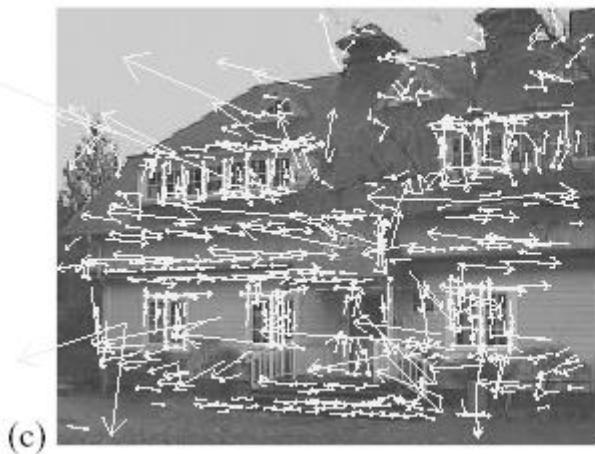


[Lowe 2004]

DoG – example result



DoG – example result



(a) 233x189 image

(b) 832 DoG extrema

(c) 729 left after peak value threshold

(d) 536 left after testing ratio of principle curvatures (removing edge responses)

Scale Invariant Detection: Summary

- ▶ **Given:** two images of the same scene with a large *scale difference* between them
- ▶ **Goal:** find *the same* interest points *independently* in each image
- ▶ **Solution:** search for *maxima* of suitable functions in *scale* and in *space* (over the image)
 1. Laplacian of Gaussian (LoG)
 2. Difference of Gaussians (DoG) – fast approximation of LoG
 3. Harris + Laplace
 4. Hessian + Laplace

Invariant regions

- ▶ So far our regions are:
 - ▶ Translation invariant (since they are centered at keypoints)
 - ▶ Scale-invariant
- ▶ What about:
 - ▶ Orientation?



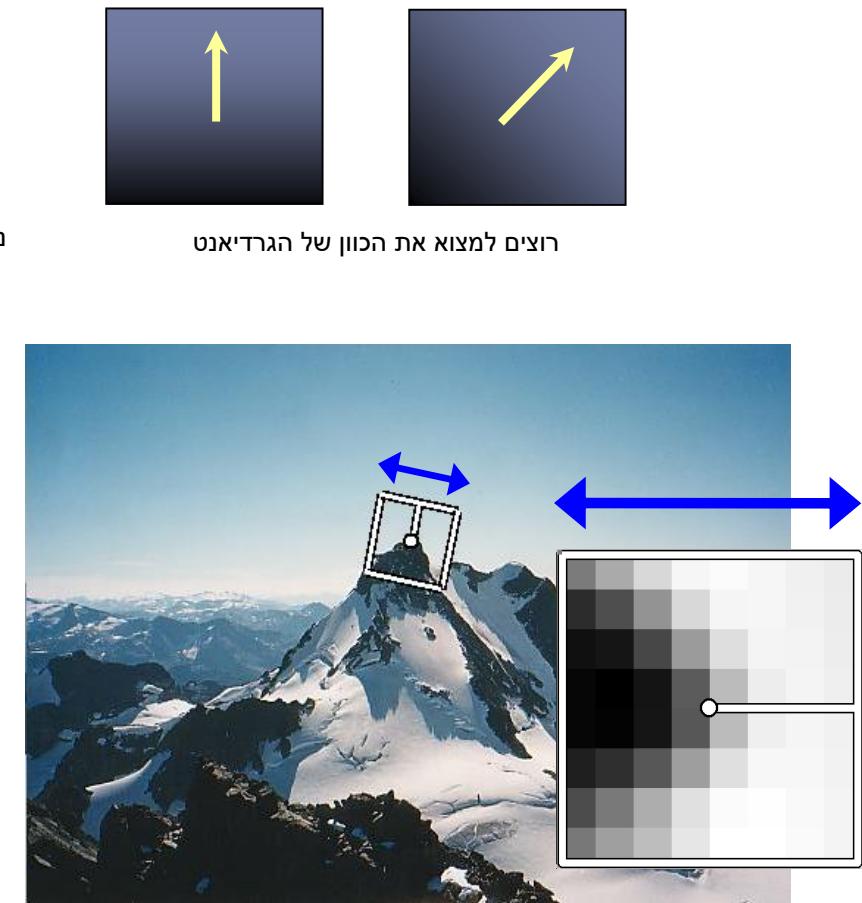
Rotation invariant regions

- ▶ Find local orientation
 - ▶ Compute a weighted histogram of gradient directions

בננה הוסטרגמת כווני הגרדיינט על ה patch

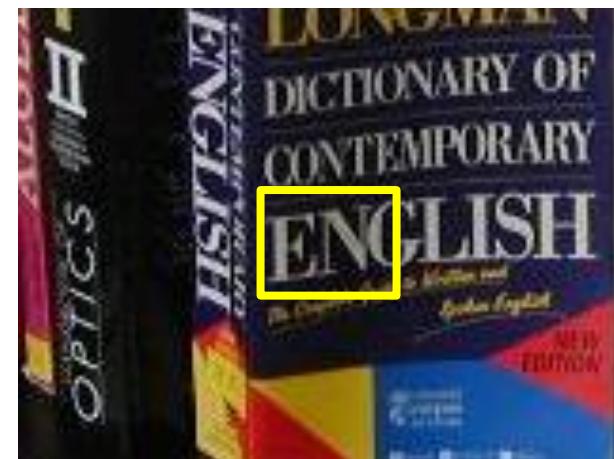
רוצים למצוא את הכוון של הגרדיינט
 - ▶ Find the dominant direction

רחלציית הזרויות, אני קובע אותה לפי הפינים של הפסטוגרמה
אפשר לא להכליל גרדינטים שהם לא חזקים
כלומר לא עוביים ערך סף
- ▶ Normalize orientation
 - ▶ Rotate patch according to this angle
 - ▶ This puts the patches into a canonical orientation.



Invariant regions

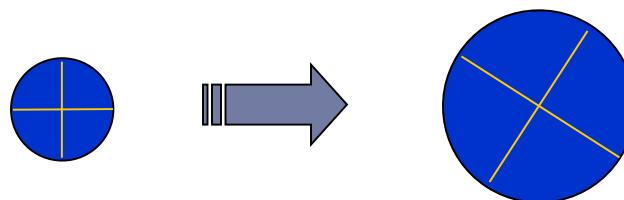
- ▶ So far our regions are:
 - ▶ Translation invariant (since they are centered at keypoints)
 - ▶ Scale-invariant
 - ▶ Rotation-invariant
- ▶ What about:
 - ▶ Other transformations?



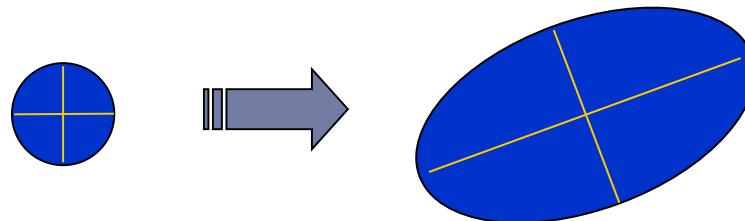
יש כאן טרנספורמציה שיר, שזה טרנספורמציה אפינית אחרת

Affine invariant regions

- ▶ Above we considered:
Similarity transform (rotation + uniform scale)



- ▶ Now we go on to:
Affine transform (rotation + non-uniform scale)



Affine adaption

- ▶ Problem
 - ▶ Determine the characteristic shape of the region
- ▶ Assumption
 - ▶ Shape can be described by a “local affine frame”
- ▶ Solution
 - ▶ Use a circular window to compute the second-moment matrix
 - ▶ Compute eigenvectors to adapt the circle to an ellipse
 - ▶ Re-compute second-moment matrix using the new window
 - ▶ Iterate...

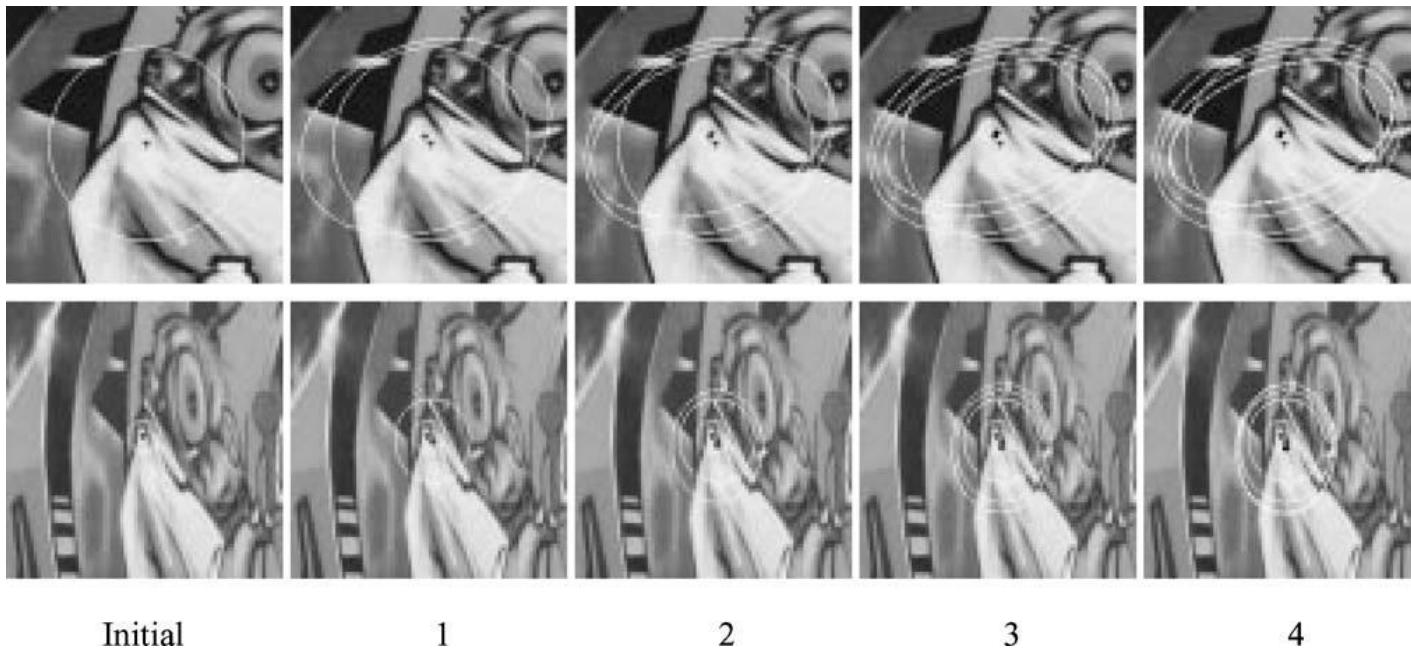
צורה שקל לשנות לה את הצורה
עוטפת את
האובייקט בצורה מדויקת יחסית



כל שמותוויצות יותר טרנספורמציות
התאמת אכוטית תרד

Iterative affine adaptation

1. Detect keypoint (e.g., multi-scale Harris)
2. Automatically select the scales (e.g., Laplace signature)
3. Adapt affine shape based on second order moment matrix
4. Refine point location



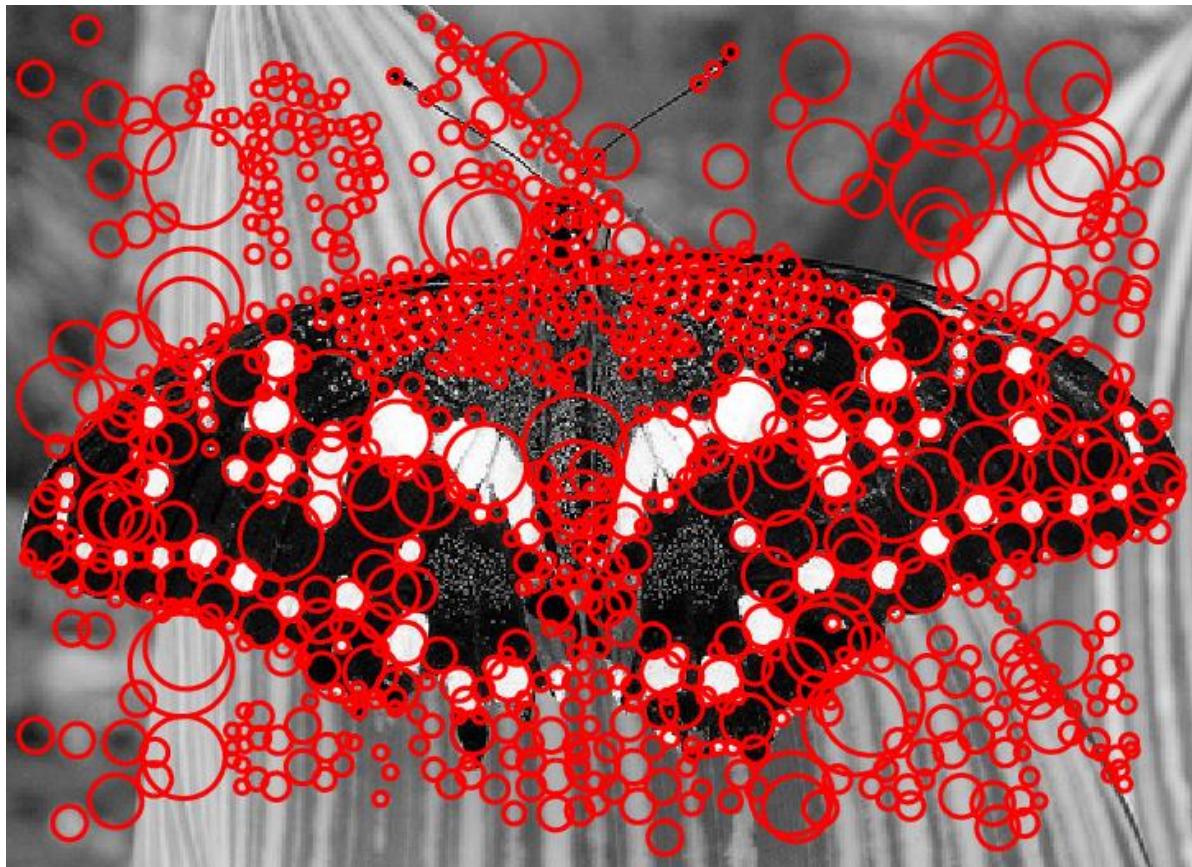
Affine normalization (deskewing)

► Steps:

- Rotate the ellipse's main axis to horizontal
- Scale the x axis, such that it forms a circle



Affine adaptation example



Scale-invariant regions

Affine adaptation example



Affine-adapted regions

Affine invariant regions example

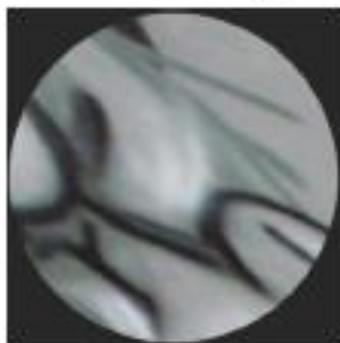


Affine invariant regions - summary

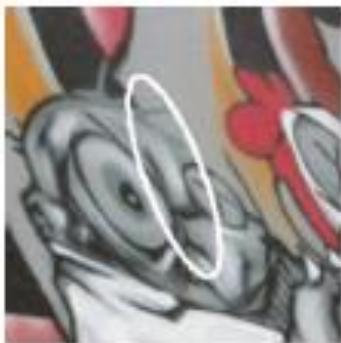
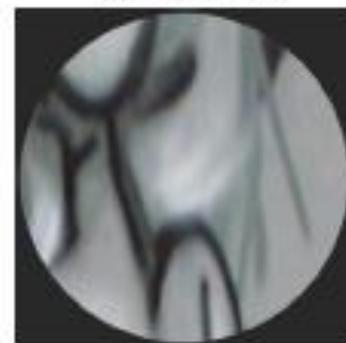
Extract affine regions



Normalize regions



Eliminate rotational ambiguity



End – From points to regions

Now you know how it works