



Alignment

Lihi Zelnik-Manor, Computer Vision

Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ Computing parametric transformations
- ▶ Panoramas


Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ Computing parametric transformations
- ▶ Panoramas

Feature matching

We know how to detect **and describe** good points



Next question: **How to match them?**



Feature matching

Given a feature in *Image1*, how to find the best match in *Image2*



1. Define distance function that compares two descriptors
2. Compute distances between all pairs features and find the ones with min distance



Feature matching - better


Given a feature in *Image1*, how to find the best match in *Image2*

1. Define distance function that compares two descriptors
2. Compute distances between all pairs features and find matches that $\text{dist}(f_1, \text{best match}) / \text{dist}(f_1, \text{second-best match}) > \text{threshold}$



Typical feature matching results

- ▶ Some matches are correct
- ▶ Some matches are incorrect



- ▶ Solution: search for a set of geometrically consistent matches

What we need to solve

- ▶ Given source and target images, how do we compute the transformation between them?



Today


- ▶ Matching local features
- ▶ **Parametric transformations**
- ▶ Computing parametric transformations
- ▶ Panoramas

Image alignment

- Turns out that in many cases there's a global transformation relating points in two images:

$$p' = Hp$$

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = H \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$





Parametric (global) warping

► Examples of parametric transformations:



original



translation



In-plane rotation



Aspect ratio




affine



perspective


Scaling

- ▶ *Scaling* a coordinate means multiplying each of its components by a scalar
- ▶ *Uniform scaling* means this scalar is the same for all components:



Scaling

- ▶ *Non-uniform scaling*: different scalars per component:



Scaling

- ▶ Scaling operation:

$$x' = ax$$

$$y' = by$$


- ▶ Or, in matrix form:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


scaling matrix S




2-D Rotation



$$\begin{aligned}x' &= x \cos(\theta) - y \sin(\theta) \\ y' &= x \sin(\theta) + y \cos(\theta)\end{aligned}$$

2-D Rotation



Polar coordinates...

$$x = r \cos(\phi)$$

$$y = r \sin(\phi)$$

$$x' = r \cos(\phi + \theta)$$

$$y' = r \sin(\phi + \theta)$$

Trig Identity...

$$x' = r \cos(\phi) \cos(\theta) - r \sin(\phi) \sin(\theta)$$

$$y' = r \sin(\phi) \cos(\theta) + r \cos(\phi) \sin(\theta)$$

Substitute...

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = x \sin(\theta) + y \cos(\theta)$$



2-D Rotation

This is easy to capture in matrix form:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Even though $\sin(\theta)$ and $\cos(\theta)$ are nonlinear functions of θ ,

- ▶ ***x' is a linear combination of x and y***
- ▶ ***y' is a linear combination of x and y***

What is the inverse transformation?

- ▶ Rotation by $-\theta$
- ▶ For rotation matrices $R^{-1} = R^T$



Basic 2D→2D Transformations

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Rotate in-plane

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale/Aspect ratio

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Affine

Combination of translation, scale,
rotation, shear

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & a & 0 \\ b & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

Affine Transformations

- ▶ Affine transformations are combinations of ...

- ▶ Linear transformations, and
- ▶ Translations

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- ▶ Properties of affine transformations:

- ▶ Origin does not necessarily map to origin
- ▶ Lines map to lines
- ▶ Parallel lines remain parallel
- ▶ Ratios are preserved
- ▶ Closed under composition



Projective Transformations

- ▶ Projective transformations ...

- ▶ Affine transformations, and
- ▶ Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- ▶ Properties of projective transformations:

- ▶ Lines map to lines
- ▶ Parallel lines do not necessarily remain parallel
- ▶ Ratios are not preserved
- ▶ Closed under composition
- ▶ Projective matrix is defined up to a scale (8 DOF)



Basic 2D→2D Transformations

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Rotate in-plane

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale/Aspect ratio


$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Affine

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Projective
(Homography)

2D image transformations (reference table)



Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

When do we get affine or homography?

- ▶ Camera does not translate (only rotation and scale)

Transformation between coordinates in 3D:


$$P' = RP + T$$

Without translation:

$$P' = RP$$

And the projections:

$$p' = Hp$$



When do we get affine or homography?

- ▶ Camera does not translate (only rotation and scale)
- ▶ Object is planar
- ▶ Works fine for small viewpoint changes


Homographies == Planar Perspective Maps

$$\mathbf{H} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix}$$

Called a *homography*
(or *planar perspective map*)



Fitting an affine transformation




Affine model approximates perspective projection of planar objects.

Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ **Computing parametric transformations**
 - ▶ **Least-squares**
 - ▶ RANSAC
- ▶ Panoramas

Fitting an affine transformation


- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \Rightarrow \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Fitting an affine transformation

- ▶ Solve with Least-squares
$$\begin{bmatrix} & & & & & & \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$
- ▶ How many matches (correspondence pairs) do we need to solve for the transformation parameters?
- ▶ Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for any pixel (x_{new}, y_{new}) ?



Fitting a projective transformation

- Recall working with homogenous coordinates

$$\begin{bmatrix} x'_i \\ y'_i \\ w'_i \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \begin{aligned} x'_i &\rightarrow \frac{x'_i}{w'_i}, \\ y'_i &\rightarrow \frac{y'_i}{w'_i} \end{aligned}$$

- The equations we get are


$$\begin{aligned} x'_i &= \frac{ax_i + by_i + c}{gx_i + hy_i + a} \\ y'_i &= \frac{dx_i + ey_i + f}{gx_i + hy_i + a} \end{aligned}$$



$$\begin{bmatrix} \dots \\ \dots \\ \dots \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \\ g \\ h \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ \vdots \end{bmatrix}$$

- Solve with SVD


Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ **Computing parametric transformations**
 - ▶ Least-squares
 - ▶ **RANSAC** 
- ▶ Panoramas


RANSAC

1. Randomly select a *seed group* of matches
2. Compute transformation (using Least-squares) from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers
5. Keep the transformation with the largest number of inliers


RANSAC example: Translation




RANSAC example: Translation




RANSAC example: Translation




RANSAC example: Translation




RANSAC example: Translation



RANSAC example: Translation




RANSAC example: Translation




RANSAC example: homography

All matches




RANSAC example: homography

After RANSAC




RANSAC example: homography


Applying the homography



Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ Computing parametric transformations
 - ▶ Least-squares
 - ▶ RANSAC
- ▶ **Panoramas**
 - ▶ **Multi-frame estimation**
 - ▶ Warping 
 - ▶ Blending

Panoramas



Obtain a wider angle view by combining multiple images.

Panoramas



Panoramas




Panoramas



Problem: Drift


- ▶ **Error accumulation**

- ▶ small (vertical) errors accumulate over time
- ▶ apply correction so that sum = 0 (for 360° pan.)



Aligning multiple images

► Bundle adjustment



$$E = \sum_{i=1}^{\text{all pairs}} \sum_{j=1}^{\text{matches in } i} f(r_{ij})$$

$$f(\mathbf{x}) = \begin{cases} |\mathbf{x}|, & \text{if } |\mathbf{x}| < x_{max} \\ x_{max}, & \text{if } |\mathbf{x}| \geq x_{max} \end{cases}$$


f , x_{max} let us bound the effect of outliers

Gradient descent

$$\mathbf{x}_{min} = \underset{\mathbf{x}}{\operatorname{argmin}} f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$x_1 = x_0 + \delta x$$

Given a starting location x_0 , we can look at the $\frac{df}{dx}$ and move in the *downhill* direction to generate a new estimate, $x_1 = x_0 + \delta x$.



Newton's method

Fit a quadratic approximation to $f(x)$ using both gradient and curvature information at x .

- Expand $f(x)$ locally using a Taylor series.

$$f(x + \delta x) = f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) + \text{h.o.t}$$

- Find the δx which minimizes this local quadratic approximation.

$$\delta x = -\frac{f'(x)}{f''(x)}$$

- Update x

$$x_{n+1} = x_n - \frac{f'(x)}{f''(x)}$$

Newton in N dimensions

$$f(\mathbf{x} + \mathbf{h}) \approx f(\mathbf{x}) + \nabla f(\mathbf{x})^\top \mathbf{h} + \frac{1}{2} \mathbf{h}^\top \mathbf{H}(\mathbf{x}) \mathbf{h}$$

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \lambda_n \mathbf{H}(\mathbf{x}_n)^{-1} \nabla f(\mathbf{x}_n)$$

The gradient $\nabla f(\mathbf{x})$ of $f(\mathbf{x})$ is the vector

$$\nabla f(\mathbf{x}) = \left[\frac{\partial f}{\partial x_1} \cdots \frac{\partial f}{\partial x_N} \right]$$

The Hessian $\mathbf{H}(\mathbf{x})$ of $f(\mathbf{x})$ is the symmetric matrix

$$\mathbf{H}(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_N} \\ \vdots & \ddots & \\ \frac{\partial^2 f}{\partial x_1 \partial x_N} & & \frac{\partial^2 f}{\partial x_N^2} \end{bmatrix}$$

Minimizing sum of residuals

$$f(\mathbf{x}) = \sum_{i=1}^M r_i^2$$

$$\mathbf{J}(\mathbf{x}) = \begin{pmatrix} \frac{\partial r_1}{\partial x_1} & \cdots & \frac{\partial r_1}{\partial x_N} \\ \vdots & \ddots & \vdots \\ \frac{\partial r_M}{\partial x_1} & \cdots & \frac{\partial r_M}{\partial x_N} \end{pmatrix}$$

$$\nabla f(\mathbf{x}) = 2\mathbf{J}^\top \mathbf{r}$$

$$\mathbf{H}(\mathbf{x}) = 2\mathbf{J}^\top \mathbf{J} + \sum_{i=1}^M r_i \frac{d^2 r_i}{d\mathbf{x}^2}$$

Levenberg-Marquardt

$$H(\mathbf{x}, \lambda) = 2\mathbf{J}^\top \mathbf{J} + \lambda \mathbf{I}$$

Levenberg-Marquardt

$$H(\mathbf{x}, \lambda) = 2\mathbf{J}^\top \mathbf{J} + \lambda \mathbf{I}$$

All unknowns

$$\Theta = (\mathbf{J}^T \mathbf{J} + \sigma^2 \mathbf{C}_p^{-1})^{-1} \mathbf{J}^T \mathbf{r}$$

Compute analytically

$$J = \frac{\partial r}{\partial \Theta}$$

Jacobian

Prior parameter
Covariance matrix

$$\mathbf{C}_p$$

$$\sigma_\theta = \pi/16 \quad \sigma_f = \bar{f}/10$$

Step size

$$\sigma$$

Multi-frame optimization

Rotation matrix

$$\mathbf{R}_i = e^{[\boldsymbol{\theta}_i]_{\times}}, \quad [\boldsymbol{\theta}_i]_{\times} = \begin{bmatrix} 0 & -\theta_{i3} & \theta_{i2} \\ \theta_{i3} & 0 & -\theta_{i1} \\ -\theta_{i2} & \theta_{i1} & 0 \end{bmatrix}$$

Calibration matrix

$$\mathbf{K}_i = \begin{bmatrix} f_i & 0 & 0 \\ 0 & f_i & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Unknown parameters for image i

$$\boldsymbol{\theta}_i = [\theta_{i1} \quad \theta_{i2} \quad \theta_{i3} \quad f_i]^T$$

Bundle adjustment formulations

All pairs optimization:

$$E_{\text{all-pairs-2D}} = \sum_i \sum_{jk} c_{ij} c_{ik} \|\tilde{x}_{ik}(\hat{x}_{ij}; R_j, f_j, R_k, f_k) - \hat{x}_{ik}\|^2,$$

Confidence / uncertainty of point i in image j

Map 2D point i in image j to 2D point in image k

Full bundle adjustment, using 3-D point positions $\{x_i\}$

$$E_{\text{BA-2D}} = \sum_i \sum_j c_{ij} \|\tilde{x}_{ij}(x_i; R_j, f_j) - \hat{x}_{ij}\|^2,$$

Map 3D point i to 2D point in image i

Bundle adjustment using 3-D ray:

$$E_{\text{BA-3D}} = \sum_i \sum_j c_{ij} \|\tilde{x}_i(\hat{x}_{ij}; R_j, f_j) - x_i\|^2,$$

3-D ray from point i

All-pairs 3-D ray formulation:

$$E_{\text{all-pairs-3D}} = \sum_i \sum_{jk} c_{ij} c_{ik} \|\tilde{x}_i(\hat{x}_{ij}; R_j, f_j) - \tilde{x}_i(\hat{x}_{ik}; R_k, f_k)\|^2.$$

3-D ray from points i and j

Projected point

$\tilde{x}_{ij} \sim K_j R_j x_i \text{ and } x_i \sim R_j^{-1} K_j^{-1} \tilde{x}_{ij},$


3-D ray from point

Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ Computing parametric transformations
 - ▶ Least-squares
 - ▶ RANSAC
- ▶ **Panoramas**
 - ▶ Multi-frame estimation
 - ▶ **Warping**
 - ▶ Blending


What we need to solve

- ▶ Given source and target images, and the transformation between them, how do we align them?
- ▶ Send each pixel \mathbf{x} in *image 1* to its corresponding location \mathbf{x}' in *image 2*



Forward Warping

- ▶ What if pixel lands “between” two pixels?
- ▶ Answer: add “contribution” to several pixels and normalize (*splatting*)




Forward Warping

- ▶ What if pixel lands “between” two pixels?
- ▶ Answer: add “contribution” to several pixels and normalize (*splatting*)
- ▶ Limitation: Holes (some pixels are never visited)




Inverse Warping

- ▶ For each pixel \mathbf{x}' in image 2 find its origin \mathbf{x} in image 1
- ▶ Problem: What if pixel comes from “between” two pixels?




Inverse Warping

- ▶ For each pixel \mathbf{x}' in image 2 find its origin \mathbf{x} in image 1
- ▶ Problem: What if pixel comes from “between” two pixels?
- ▶ Answer: *interpolate* color value from *neighbors*



Bilinear interpolation

Sampling at $f(x,y)$:



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$

Interpolation

- ▶ Possible interpolation filters:
 - ▶ nearest neighbor
 - ▶ Bilinear interpolation
 - ▶ bicubic interpolation
 - ▶ sinc / FIR
- ▶ Needed to prevent “jaggies” and “texture crawl”




Today

- ▶ Matching local features
- ▶ Parametric transformations
- ▶ Computing parametric transformations
 - ▶ Least-squares
 - ▶ RANSAC
- ▶ **Panoramas**
 - ▶ Multi-frame estimation
 - ▶ Warping
 - ▶ **Blending**

Image feathering

- ▶ Weight each image proportional to its distance from the edge (distance map [Danielsson, CVGIP 1980])



- ▶ 1. Generate *weight map* for each image
- ▶ 2. Sum up all of the weights and divide by sum:
weights sum up to 1: $w_i' = w_i / (\sum_i w_i)$




Image feathering






Image feathering




Panoramas – summary

- ▶ Detect features
- ▶ Compute transformations between pairs of frames
- ▶ Refine transformations using bundle-adjustment
- ▶ Warp all images onto a single coordinate system
- ▶ Blend









End – Alignment



Now you know how it works