

Object Category Detection: Statistical Templates


Slides borrowed from Derek Hoiem and
then modified some

Today's class: Object Category Detection

- Overview of object category detection
- Statistical template matching
 - Dalal-Triggs pedestrian detector (basic concept)
 - Viola-Jones detector (cascades, integral images)
 - R-CNN detector (object proposals/CNN)

Object Category Detection

- Focus on object search: “Where is it?”
- Build templates that quickly differentiate object patch from background patch




Challenges in modeling the object class



Illumination



Object pose



Clutter



Occlusions



Intra-class
appearance




Viewpoint

Challenges in modeling the non-object class


True
Detections





Bad
Localization



Confused with
Similar Object



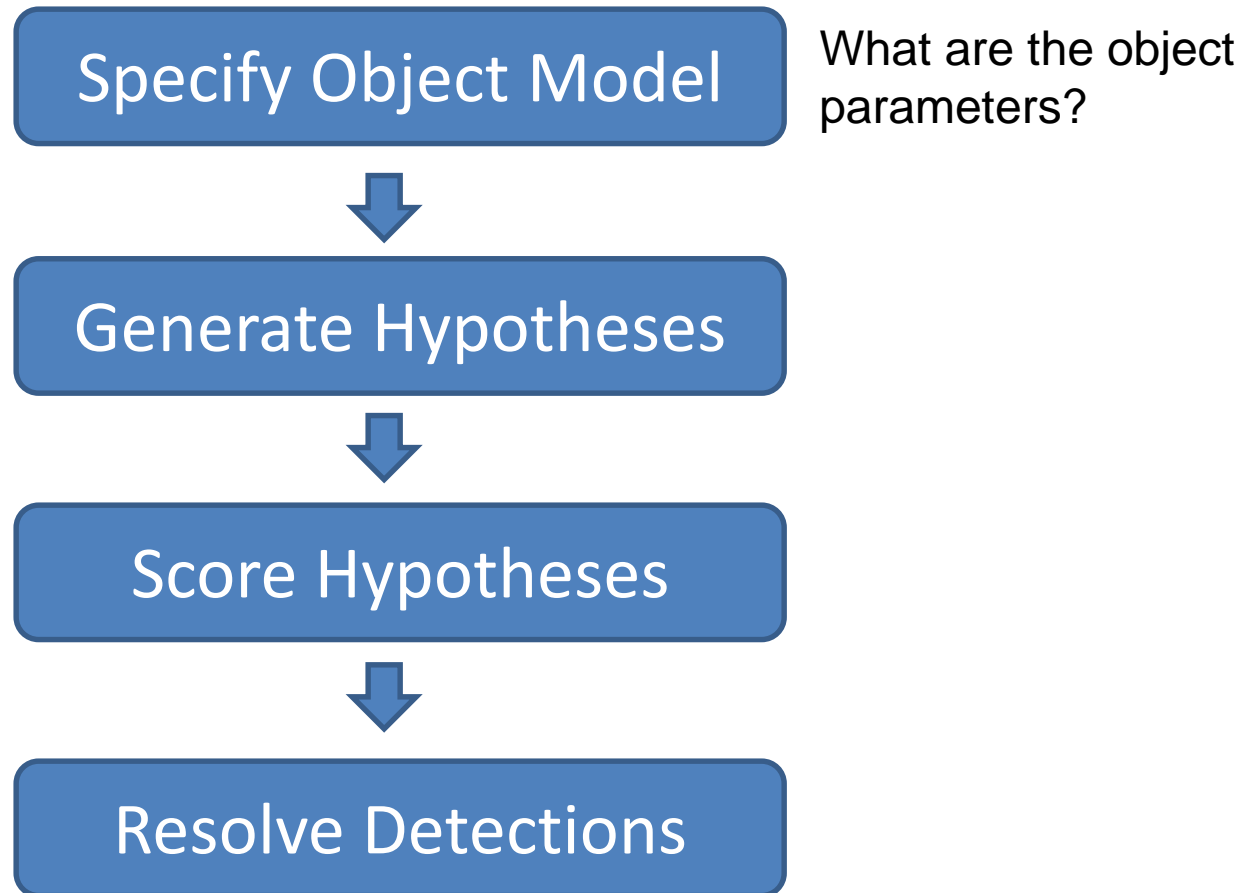
Misc. Background



Confused with
Dissimilar Objects



General Process of Object Recognition




Specifying an object model

1. Statistical Template in Bounding Box

- Object is some (x,y,w,h) in image
- Features defined wrt bounding box coordinates



Image





Template Visualization

Specifying an object model

2. Articulated parts model








- Object is configuration of parts
- Each part is detectable



Specifying an object model

3. Hybrid template/parts model

Detections



Template Visualization

root filters
coarse resolution


part filters
finer resolution

deformation
models

Specifying an object model

4. 3D-ish model

- Object is collection of 3D planar patches under affine transformation



General Process of Object Recognition

Specify Object Model



Generate Hypotheses

Propose an alignment of the model to the image



Score Hypotheses



Resolve Detections

Generating hypotheses

1. Sliding window

- Test patch at each location and scale



Generating hypotheses

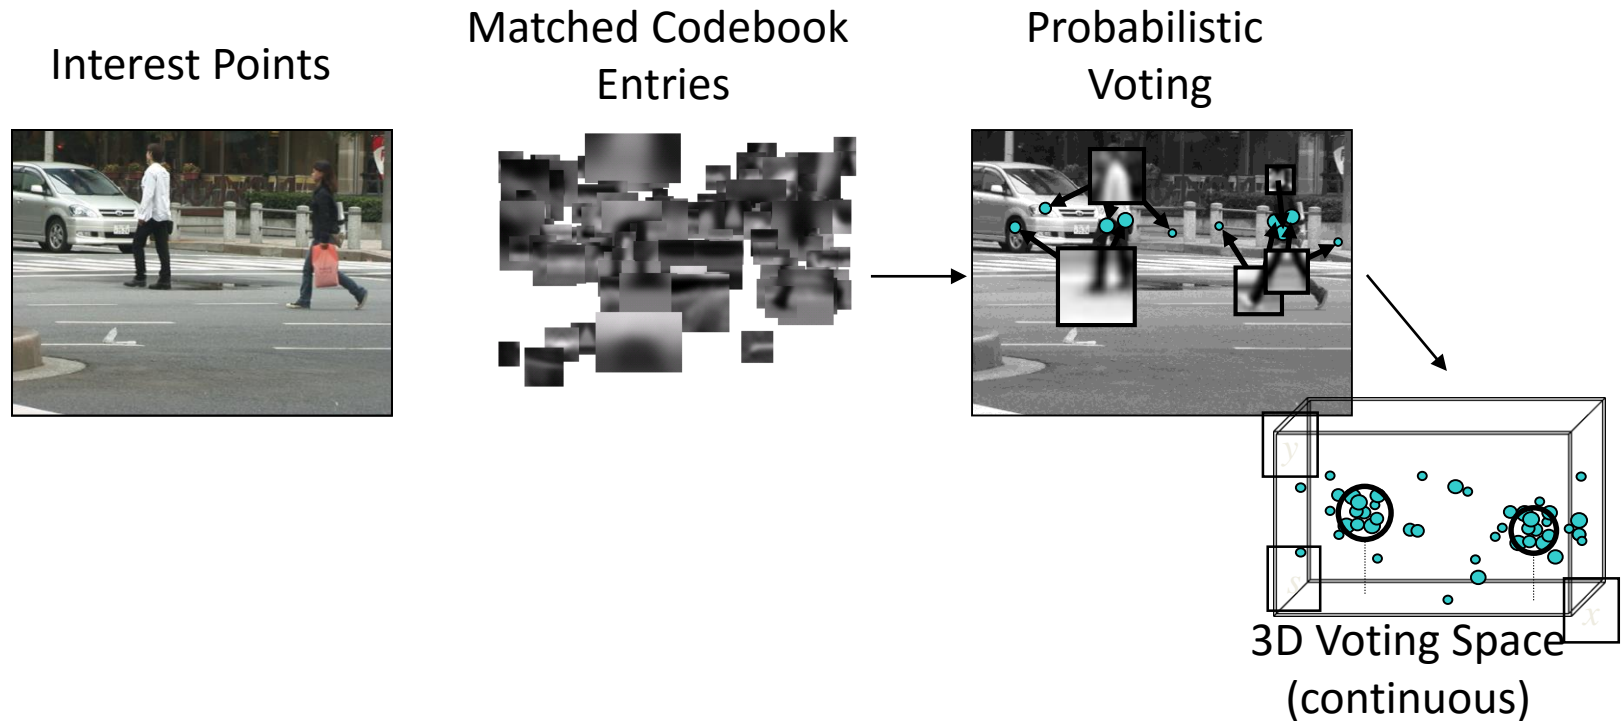
1. Sliding window

- Test patch at each location and scale








Generating hypotheses

2. Voting from patches/keypoints

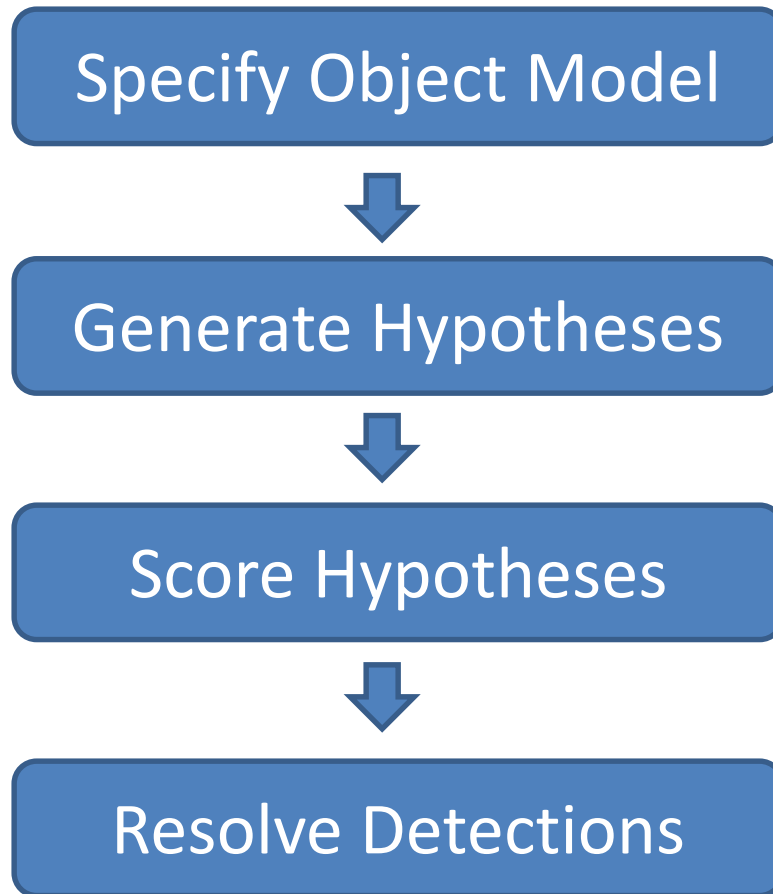


Generating hypotheses

3. Region-based proposal

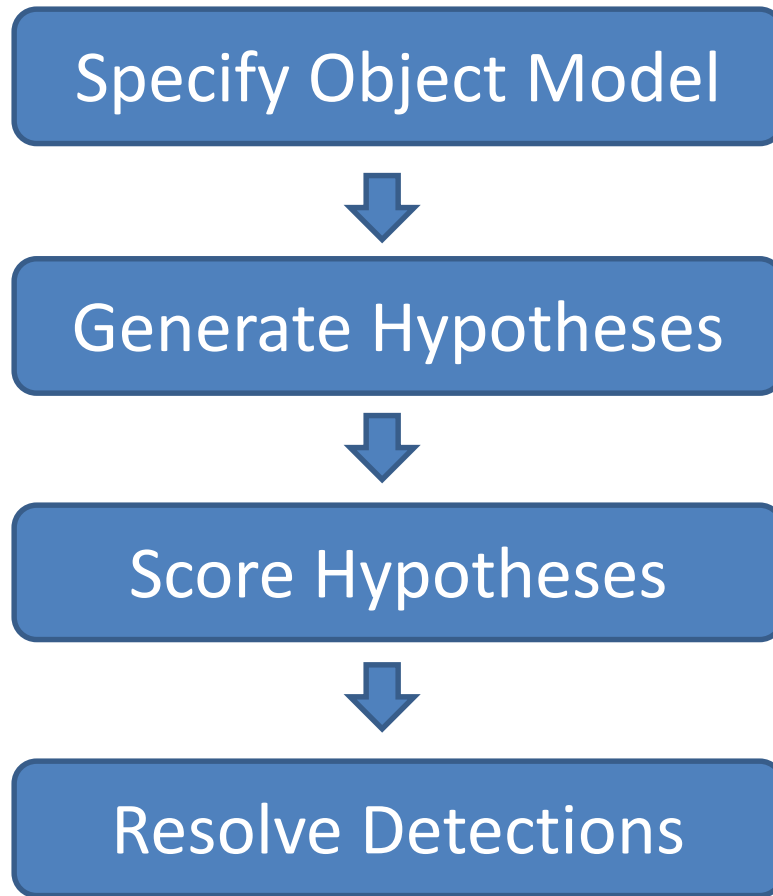


General Process of Object Recognition



Mainly-gradient based or
CNN features, usually based
on summary representation,
many classifiers

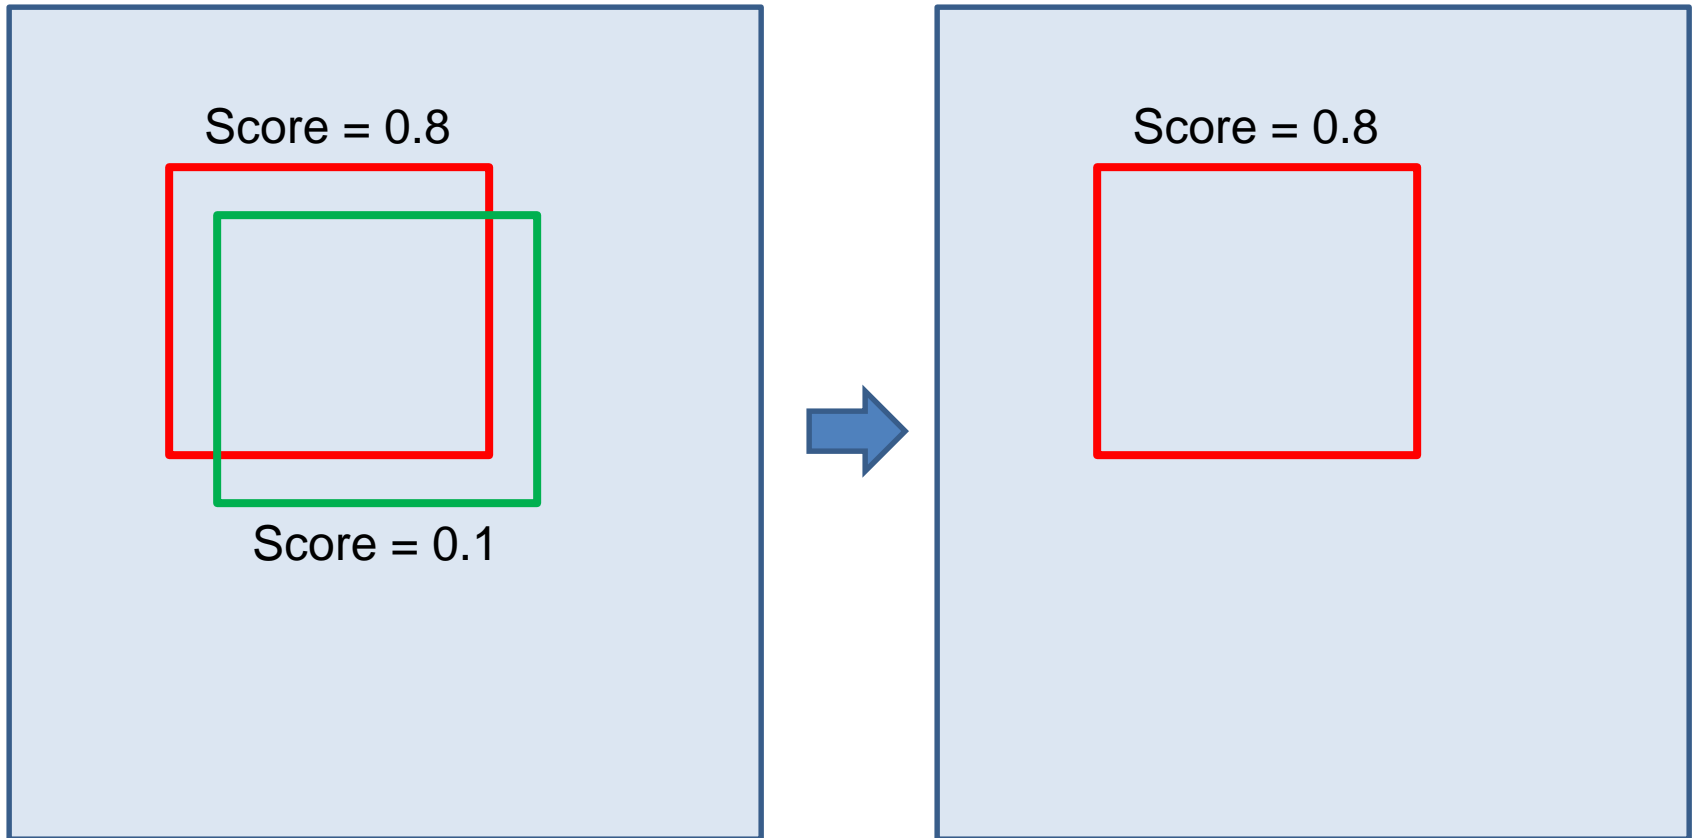
General Process of Object Recognition



Rescore each proposed object based on whole set


Resolving detection scores

1. Non-max suppression



Resolving detection scores


2. Context/reasoning



(g) Car Detections: Local



(h) Ped Detections: Local



Object category detection in computer vision


Goal: detect all pedestrians, cars, monkeys, etc in image



Basic Steps of Category Detection


1. Align

- E.g., choose position, scale orientation
- How to make this tractable?



2. Compare

- Compute similarity to an example object or to a summary representation
- Which differences in appearance are important?



Sliding window: a simple alignment solution



Each window is separately classified




Statistical Template

- Object model = sum of scores of features at fixed positions



$$+3 +2 -2 -1 -2.5 = -0.5 \stackrel{?}{>} 7.5$$

Non-object



$$+4 +1 +0.5 +3 +0.5 = 10.5 \stackrel{?}{>} 7.5$$

Object


Design challenges


- How to efficiently search for likely objects
 - Even simple models require searching hundreds of thousands of positions and scales
- Feature design and scoring
 - How should appearance be modeled? What features correspond to the object?
- How to deal with different viewpoints?
 - Often train different models for a few different viewpoints
- Implementation details
 - Window size
 - Aspect ratio
 - Translation/scale step size
 - Non-maxima suppression

Example: Dalal-Triggs pedestrian detector



1. Extract fixed-sized (64x128 pixel) window at each position and scale
2. Compute HOG (histogram of gradient) features within each window
3. Score the window with a linear SVM classifier
4. Perform non-maxima suppression to remove overlapping detections with lower scores






- Tested with
 - RGB
 - LAB
 - Grayscale

} Slightly better performance vs. grayscale
- Gamma Normalization and Compression
 - Square root
 - Log

} Very slightly better performance vs. no adjustment



Outperforms

-1	0	1
----	---	---


centered

-1	1
----	---

uncentered

1	-8	0	8	-1
---	----	---	---	----

cubic-corrected



0	1
-1	0

diagonal


-1	0	1
-2	0	2
-1	0	1

Sobel




- Histogram of gradient orientations


Orientation: 9 bins
(for unsigned angles)



Histograms in
8x8 pixel cells




- Votes weighted by magnitude
- Bilinear interpolation between cells






R-HOG

Cell




Normalize with respect to surrounding cells

$$L2 - norm : v \longrightarrow v / \sqrt{\|v\|_2^2 + \epsilon^2}$$



X=





orientations

features = 15 x 7 x 9 x 4 = 3780


cells



normalizations by neighboring cells



pos w

neg w





$$0.16 = w^T x - b$$

$$\text{sign}(0.16) = 1$$

\Rightarrow pedestrian

Detection examples




2 minute break

Something to think about...

- Sliding window detectors work
 - *very well* for faces
 - *fairly well* for cars and pedestrians
 - *badly* for cats and dogs
- Why are some classes easier than others?

Face detection

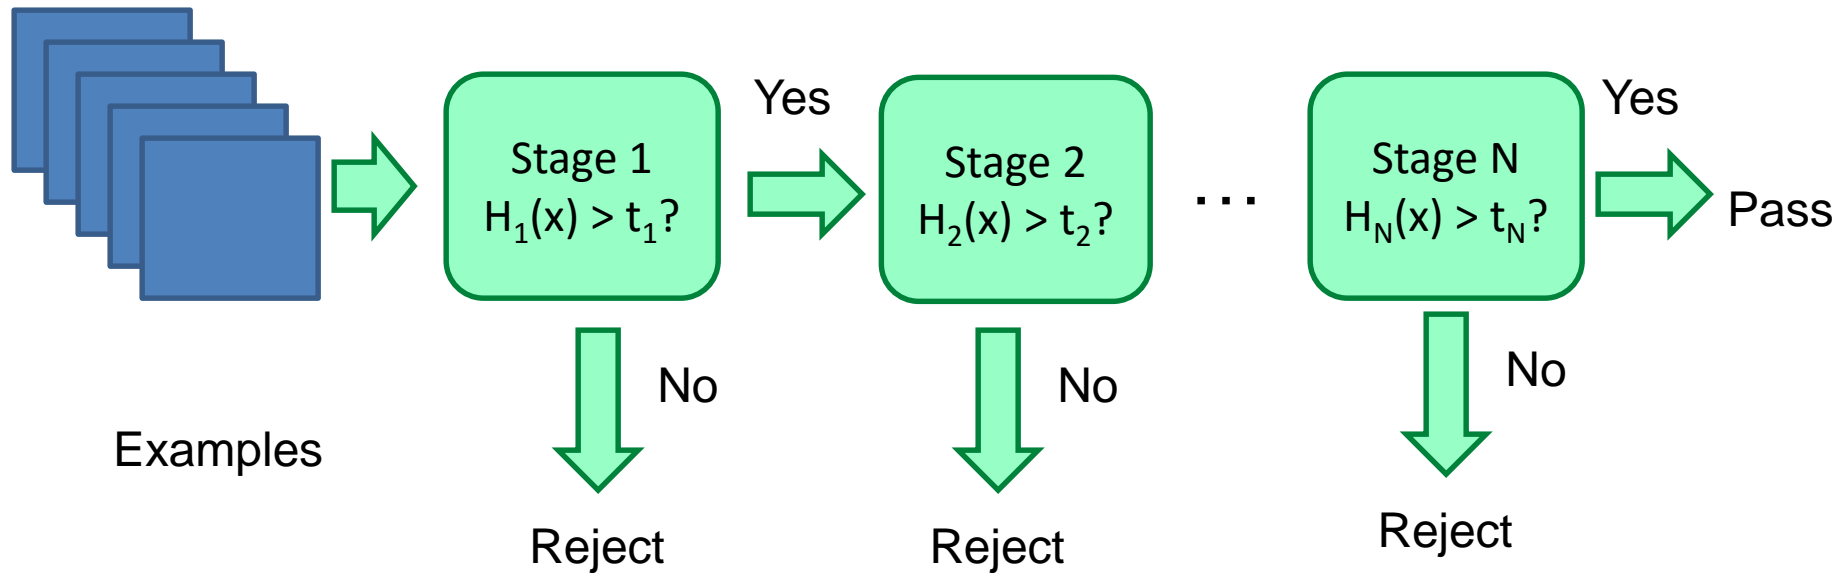


Viola-Jones sliding window detector

Fast detection through two mechanisms

- Quickly eliminate unlikely windows
- Use features that are fast to compute

Cascade for Fast Detection




- Choose threshold for low false negative rate
- Fast classifiers early in cascade
- Slow classifiers later, but most examples don't get there

Features that are fast to compute


- “Haar-like features”
 - Differences of sums of intensity
 - Thousands, computed at various positions and scales within detection window



Two-rectangle features



Three-rectangle features



Etc.

Features that are fast to compute

- “Haar-like features”
 - Differences of sums of intensity
 - Thousands, computed at various positions and scales within detection window




Two-rectangle features




Integral Images

- `ii = cumsum(cumsum(im, 1), 2)`



$ii(x,y)$ = Sum of the values in the grey region



How to compute $B-A$?

How to compute $A+D-B-C$?

Feature selection with Adaboost

- Create a large pool of features (180K)
- Select features that are discriminative and work well together
 - “Weak learner” = feature + threshold + parity

$$h_j(x) = \begin{cases} 1 & \text{if } p_j f_j(x) < p_j \theta_j \\ 0 & \text{otherwise} \end{cases}$$

- Choose weak learner that minimizes error on the weighted training set
- Reweight

Adaboost

- Given example images $(x_1, y_1), \dots, (x_n, y_n)$ where $y_i = 0, 1$ for negative and positive examples respectively.
- Initialize weights $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$ for $y_i = 0, 1$ respectively, where m and l are the number of negatives and positives respectively.
- For $t = 1, \dots, T$:

1. Normalize the weights,

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

so that w_t is a probability distribution.

2. For each feature, j , train a classifier h_j which is restricted to using a single feature. The error is evaluated with respect to w_t , $\epsilon_j = \sum_i w_i |h_j(x_i) - y_i|$.
3. Choose the classifier, h_t , with the lowest error ϵ_t .
4. Update the weights:

$$w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$$


where $e_i = 0$ if example x_i is classified correctly, $e_i = 1$ otherwise, and $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$.

- The final strong classifier is:

$$h(x) = \begin{cases} 1 & \sum_{t=1}^T \alpha_t h_t(x) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

where $\alpha_t = \log \frac{1}{\beta_t}$

Top 2 selected features




Viola-Jones details

- 38 stages with 1, 10, 25, 50 ... features
 - 6061 total used out of 180K candidates
 - 10 features evaluated on average
- Training Examples
 - 4916 positive examples
 - 10000 negative examples collected after each stage
- Scanning
 - Scale detector rather than image
 - Scale steps = 1.25 (factor between two consecutive scales)
 - Translation $1 \times \text{scale}$ (# pixels between two consecutive windows)
- Non-max suppression: average coordinates of overlapping boxes
- Train 3 classifiers and take vote

Viola Jones Results

Speed = 15 FPS (in 2001)




Detector	False detections						
	10	31	50	65	78	95	167
Viola-Jones	76.1%	88.4%	91.4%	92.0%	92.1%	92.9%	93.9%
Viola-Jones (voting)	81.1%	89.7%	92.1%	93.1%	93.1%	93.2 %	93.7%
Rowley-Baluja-Kanade	83.2%	86.0%	-	-	-	89.2%	90.1%
Schneiderman-Kanade	-	-	-	94.4%	-	-	-
Roth-Yang-Ahuja	-	-	-	-	(94.8%)	-	-

MIT + CMU face dataset

Viola Jones Results



Viola Jones Results



Fails in commercial face detection

Who's in These Photos?

The photos you uploaded were grouped automatically so you can quickly label and notify friends of these pictures. (Friends can always untag themselves.)



Who is this?



Who is this?



Unnamed people

4 Group(s), 67 Face(s)

Select someone you know and add a name, or click the "x" to ignore that person.



Add a name

Add a name

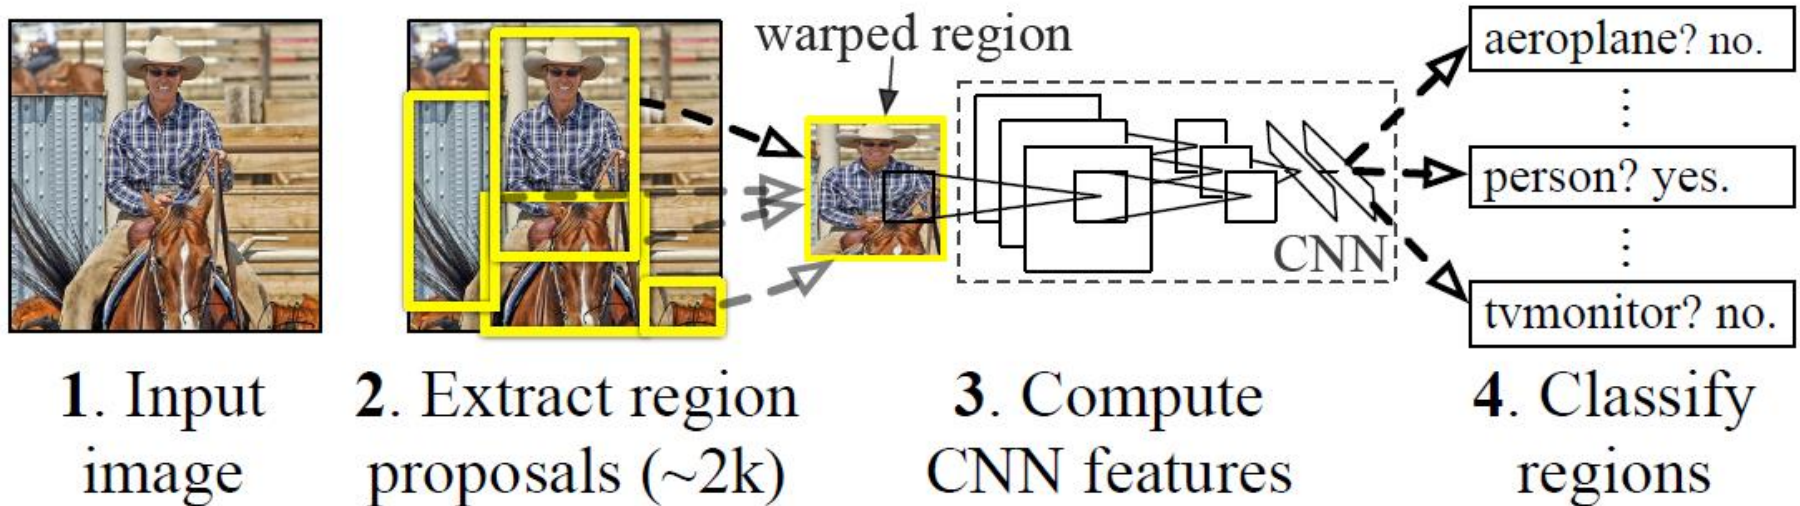
Add a name



http://www.oddee.com/item_98248.aspx


Detection with CNNs

R-CNN (Girshick et al. CVPR 2014)



- Replace sliding windows with “selective search” region proposals (Uijilings et al. IJCV 2013)
- Extract rectangles around regions and resize to 227x227
- Extract features with fine-tuned CNN (that was initialized with network trained on ImageNet before training)
- Classify last layer of network features with SVM

Region proposals



- Hierarchical image segmentation
- Random sampling of regions as object proposals

Sliding window vs. region proposals


Sliding window

- Comprehensive search over position, scale (sometimes aspect, though expensive)
- Typically 100K candidates
- Simple
- Speed boost through convolution often possible
- Repeatable
- Even with many candidates, may not be a good fit to object

Region proposals


- Search over regions guided by image contours/patterns with varying aspect/size
- Typically 2-10K candidates
- Random (not repeatable)
- Requires a preprocess (currently 1-5s)
- Often requires resizing patch to fit fixed size
- More likely to provide candidates with very good object fit

Improvements in Object Detection



Statistical Template
Matching

Improvements in Object Detection



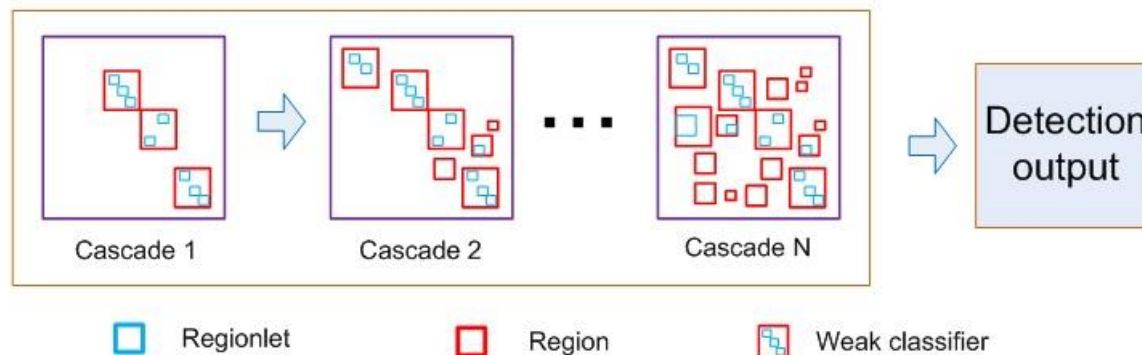
Regionlets, Wang et al

- Feature histograms are built in variable regions (vs fixed size cells, 8x8 HOG for example)
- Feature extraction regions are normalized to detection windows.
- Deformation handling is learned from data.
- Regionlets model is not limited by a fixed scale or aspect ratio.


Generate candidate detection bounding boxes




Boosting classifier cascades



Improvements in Object Detection






Improvements in Object Detection



Key Advance: Learn effective features from massive amounts of labeled data *and* adapt to new tasks with less data

Better Features

False positive type distribution



Loc = localization

Oth = other / dissimilar classes

Sim = similar classes

BG = background

Analysis software: D. Hoiem, Y. Chodpathumwan, and Q. Dai.
Diagnosing Error in Object Detectors. ECCV, 2012.

Mistakes are often reasonable

Bicycle: AP = 72.8%




bicycle (loc): ov=0.41 1-r=0.64



bicycle (loc): ov=0.35 1-r=0.61



bicycle (loc): ov=0.15 1-r=0.59



bicycle (loc): ov=0.44 1-r=0.57



bicycle (sim): ov=0.00 1-r=0.56



bicycle (bg): ov=0.00 1-r=0.52




bicycle (loc): ov=0.55 1-r=0.47



bicycle (bg): ov=0.00 1-r=0.47




icycle (loc): ov=0.46 1-r=0.45



bicycle (loc): ov=0.10 1-r=0.45



bicycle (loc): ov=0.42 1-r=0.45



bicycle (bg): ov=0.00 1-r=0.44

Strengths and Weaknesses of Statistical Template Approach

Strengths

- Works very well for non-deformable objects: faces, cars, upright pedestrians
- Fast detection

Weaknesses

- Sliding window has difficulty with deformable objects (proposals works with flexible features works better)
- Not robust to occlusion
- Requires lots of training data


Tricks of the trade

- Details in feature computation really matter
 - E.g., normalization in Dalal-Triggs improves detection rate by 27% at fixed false positive rate
- Template size
 - Typical choice for sliding window is size of smallest detectable object
 - For CNNs, typically based on what pretrained features are available
- “Jittering” to create synthetic positive examples
 - Create slightly rotated, translated, scaled, mirrored versions as extra positive examples
- Bootstrapping to get hard negative examples
 1. Randomly sample negative examples
 2. Train detector
 3. Sample negative examples that score > -1
 4. Repeat until all high-scoring negative examples fit in memory


Influential Works in Detection

- Sung-Poggio (1994, 1998) : ~2100 citations
 - Basic idea of statistical template detection (I think), bootstrapping to get “face-like” negative examples, multiple whole-face prototypes (in 1994)
- Rowley-Baluja-Kanade (1996-1998) : ~4200
 - “Parts” at fixed position, non-maxima suppression, simple cascade, rotation, pretty good accuracy, fast
- Schneiderman-Kanade (1998-2000,2004) : ~2250
 - Careful feature/classifier engineering, excellent results, cascade
- Viola-Jones (2001, 2004) : ~20,000
 - Haar-like features, Adaboost as feature selection, hyper-cascade, very fast, easy to implement
- Dalal-Triggs (2005) : ~11000
 - Careful feature engineering, excellent results, HOG feature, online code
- Felzenszwalb-Huttenlocher (2000): ~1600
 - Efficient way to solve part-based detectors
- Felzenszwalb-McAllester-Ramanan (2008,2010)? ~4000
 - Excellent template/parts-based blend
- Girshick-Donahue-Darrell-Malik (2014) ~300
 - Region proposals + fine-tuned CNN features (marks significant advance in accuracy over hog-based methods)


Summary: statistical templates




Sliding window: scan image pyramid




Region proposals: edge/region-based, resize to fixed window



HOG



Fast randomized features



CNN features

SVM

Boosted stabs

Neural network

Non-max suppression

Segment or refine localization