Image Features and Categorization

Slides borrowed from Derek Hoiem

Today: Image features and categorization

- General concepts of categorization
  - Why? What? How?
- Image features
  - Color, texture, gradient, shape, interest points
  - Histograms, feature encoding, and pooling
  - CNN as feature
- Image and region categorization

What do you see in this image?

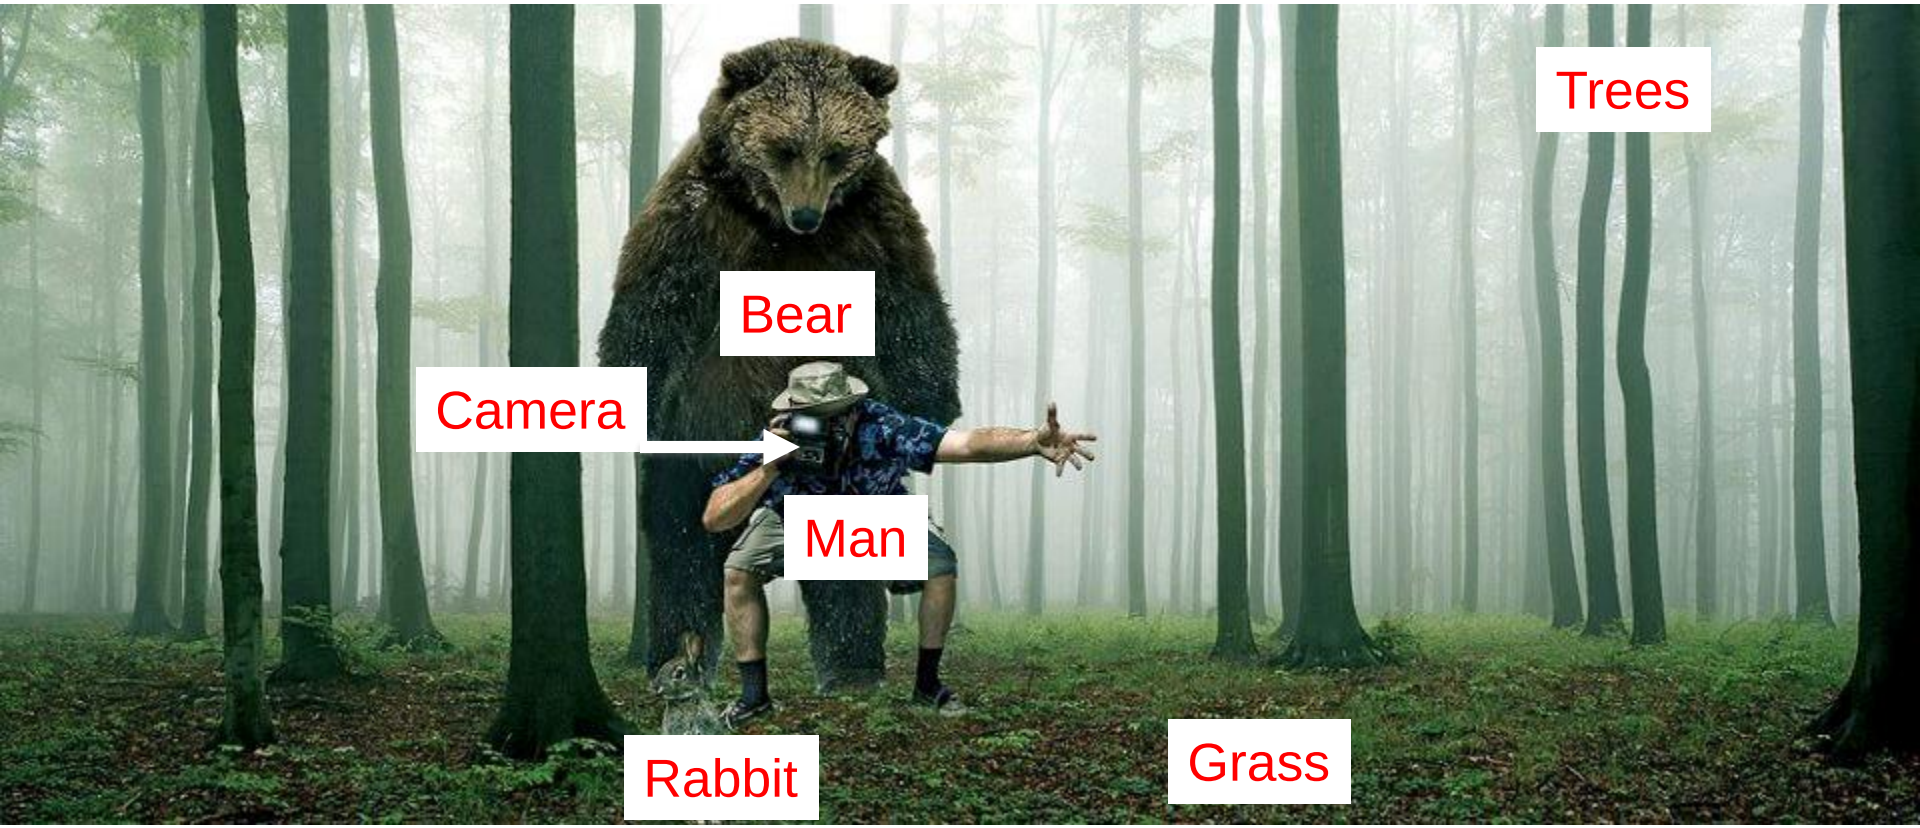

Forest

Describe, predict, or interact with the object based on visual cues

Is it **dangerous**?

Is it **alive**?

How **fast** does it run?

Is it **soft**?

Does it have a **tail**?

Can I **poke with it**?

Why do we care about categories?

- From an object's category, we can make predictions about its behavior in the future, beyond of what is immediately perceived.
- Pointers to knowledge
  - Help to understand individual cases not previously encountered
- Communication

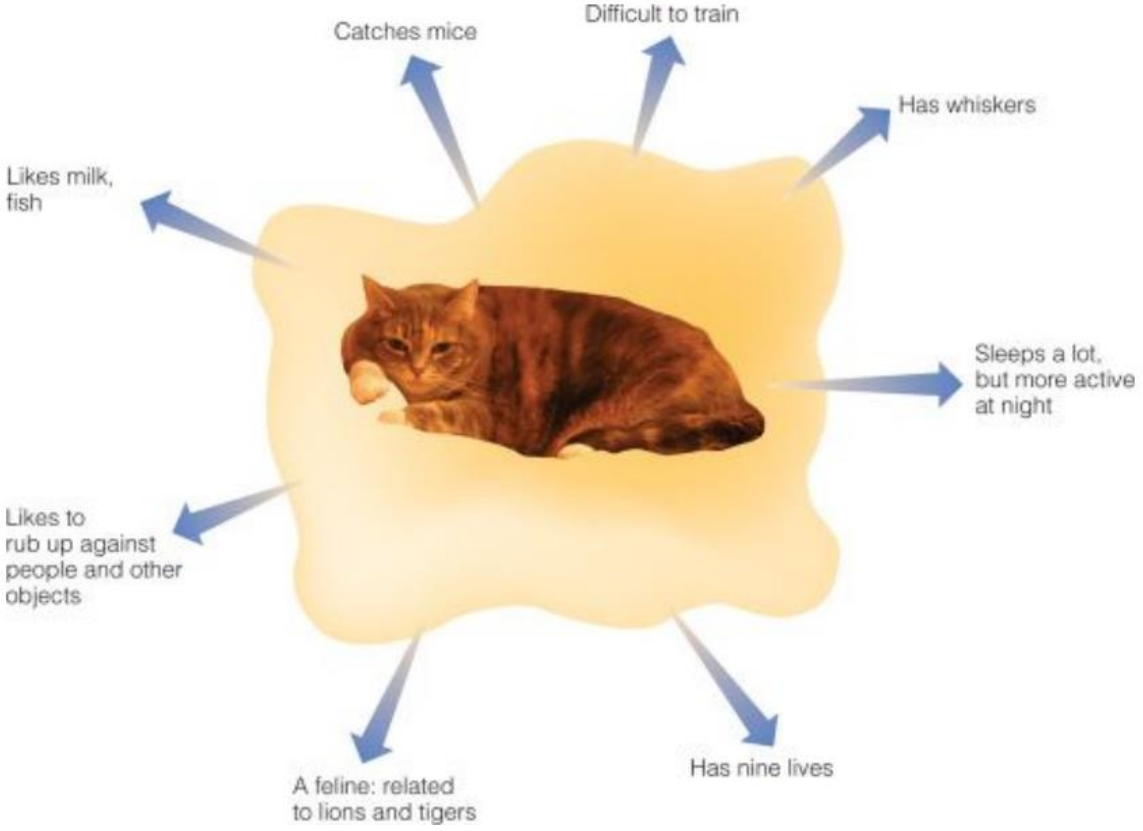

Theory of categorization

How do we determine if something is a member of a particular category?

- Definitional approach
- Prototype approach
- Exemplar approach

Definitional approach: classical view of categories

- Plato & Aristotle
  - Categories are defined by a list of properties shared by all elements in a category
  - Category membership is binary
  - Every member in the category is equal

Aristotle by Francesco Hayez

The Categories (Aristotle) :

https://en.wikipedia.org/wiki/Categories_(Aristotle)

Slide Credit: A. A. Efros

Prototype or sum of exemplars ?

Prototype Model

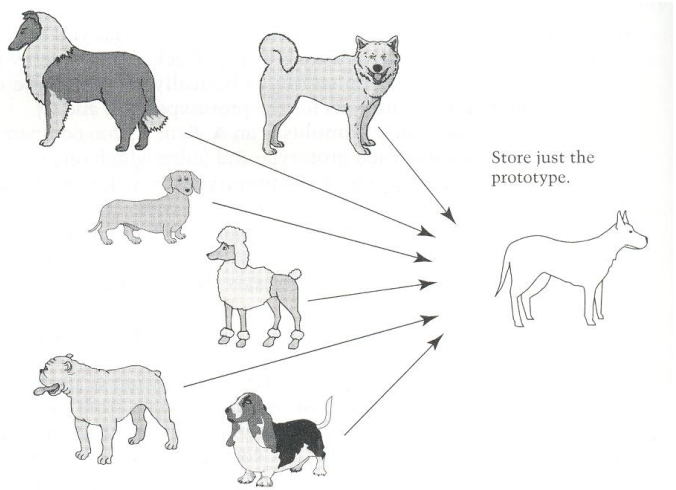

Figure 7.3. *Schematic of the prototype model. Although many exemplars are seen, only the prototype is stored. The prototype is updated continually to incorporate more experience with new exemplars.*

Category judgments are made by comparing a new exemplar to the prototype.

Exemplars Model

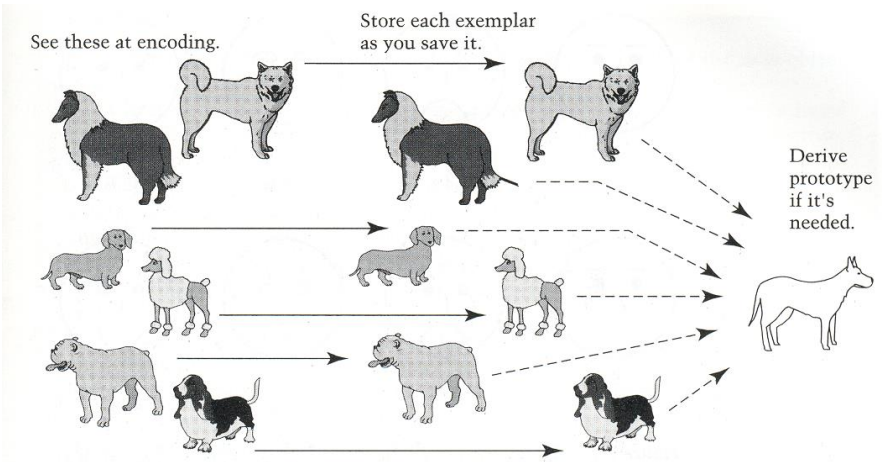

Figure 7.4. *Schematic of the exemplar model. As each exemplar is seen, it is encoded into memory. A prototype is abstracted only when it is needed, for example, when a new exemplar must be categorized.*

Category judgments are made by comparing a new exemplar to all the old exemplars of a category or to the exemplar that is the most appropriate

Slide Credit: Torralba

Levels of categorization [Rosch 70s]

Definition of Basic Level:

- **Similar shape**: Basic level categories are the highest-level category for which their members have similar shapes.
- **Similar motor interactions**: … for which people interact with its members using similar motor sequences.
- **Common attributes**: … there are a significant number of attributes in common between pairs of members.

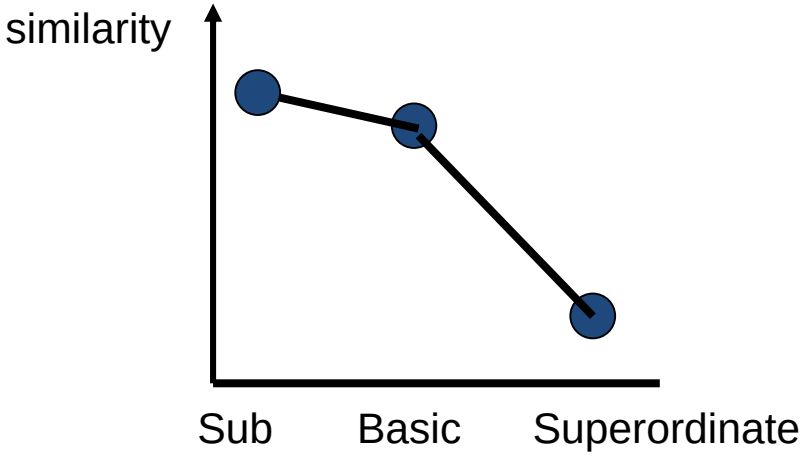

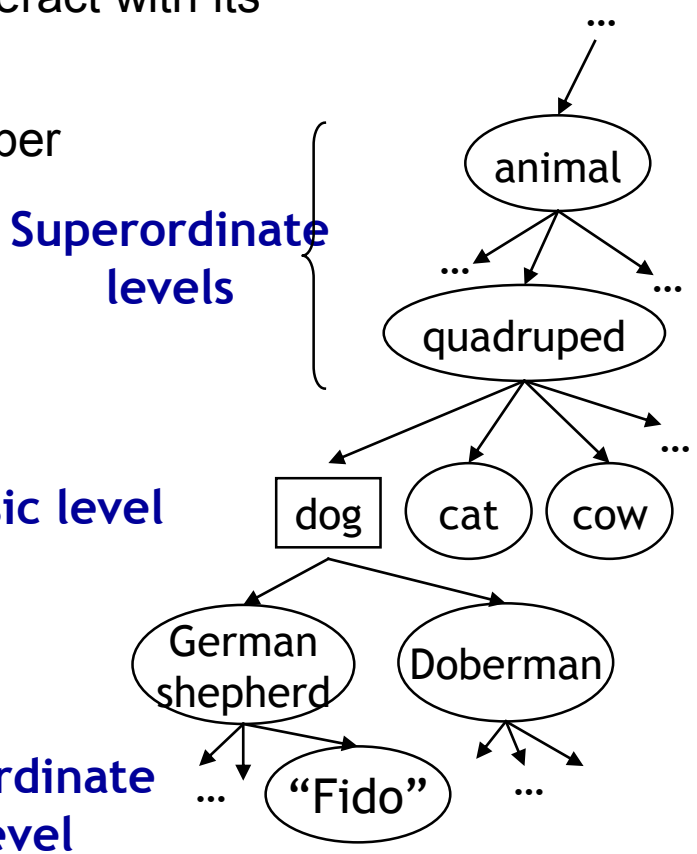

Rosch et a. Principle of categorization, 1978

Image categorization

- Cat vs Dog

Image categorization

- Object recognition

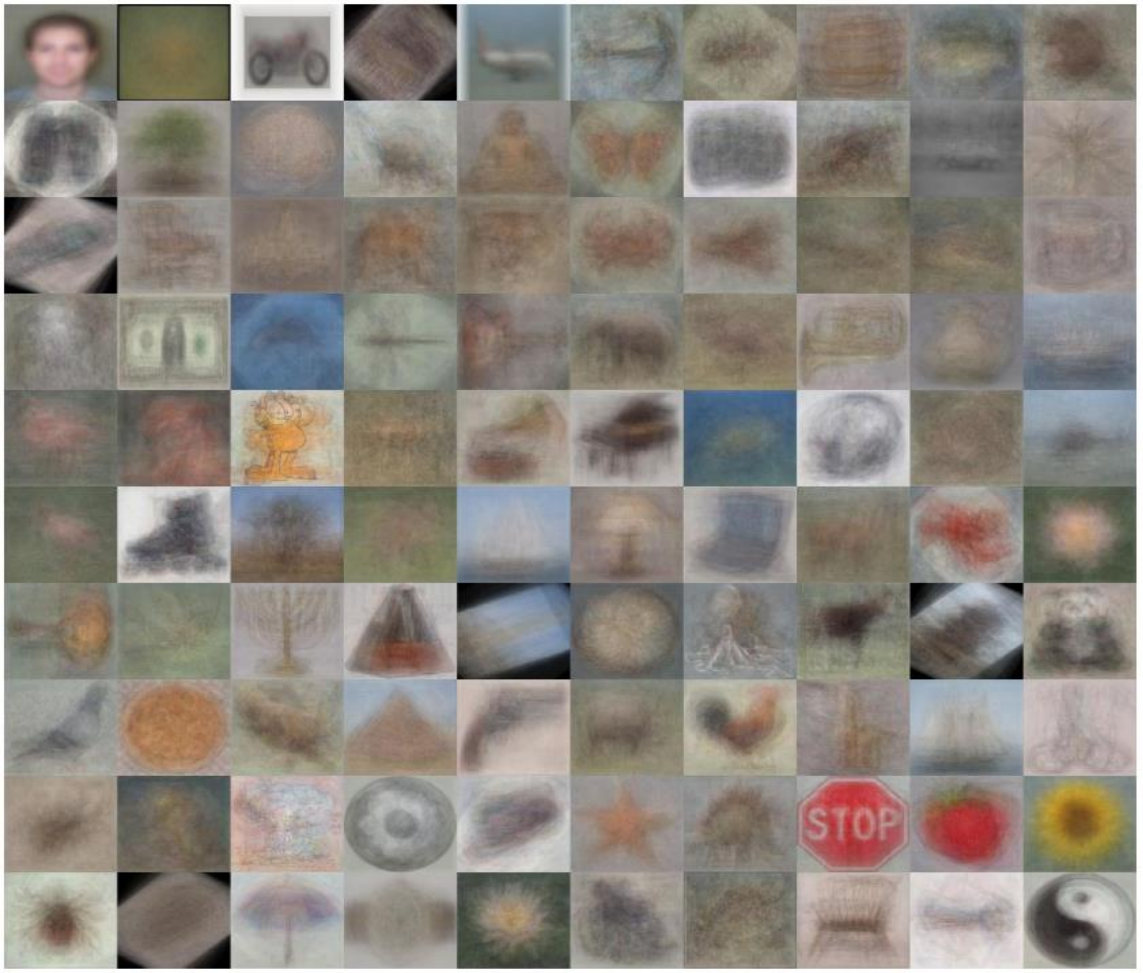

Caltech 101 Average Object Images

Image categorization

- Fine-grained recognition

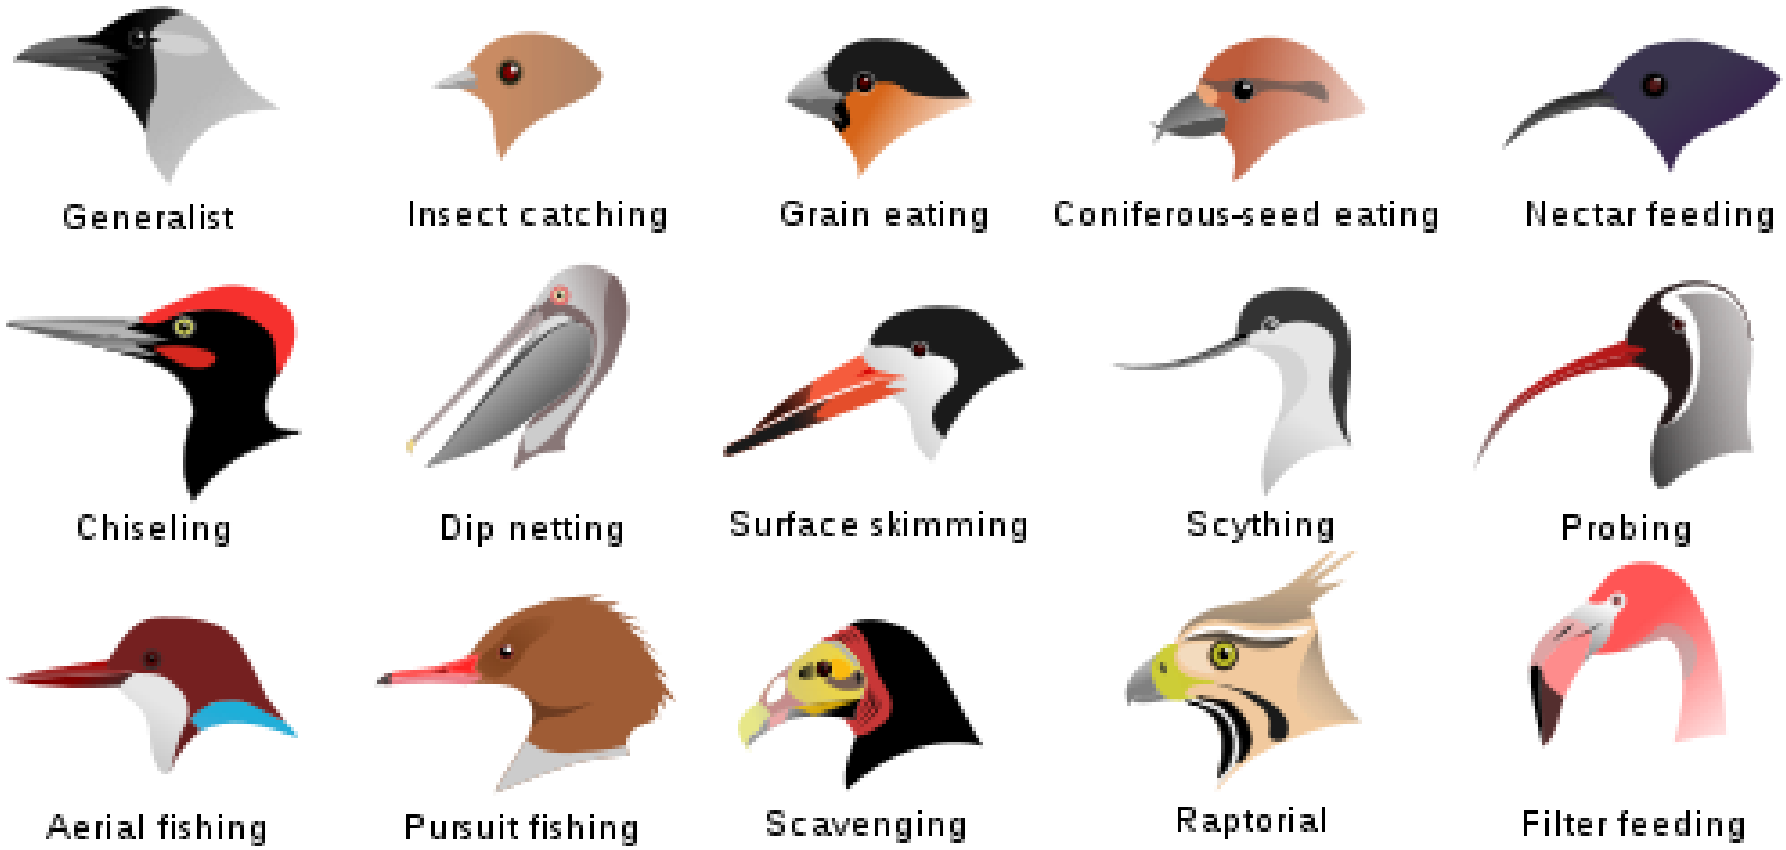

Visipedia Project

Image categorization

- Place recognition

Places Database [Zhou et al. NIPS 2014]

Image categorization

- Visual font recognition

[Chen et al. CVPR 2014]

Image categorization

- Dating historical photos

1940 1953 1966 1977

[Palermo et al. ECCV 2012]

Image categorization

- Image style recognition

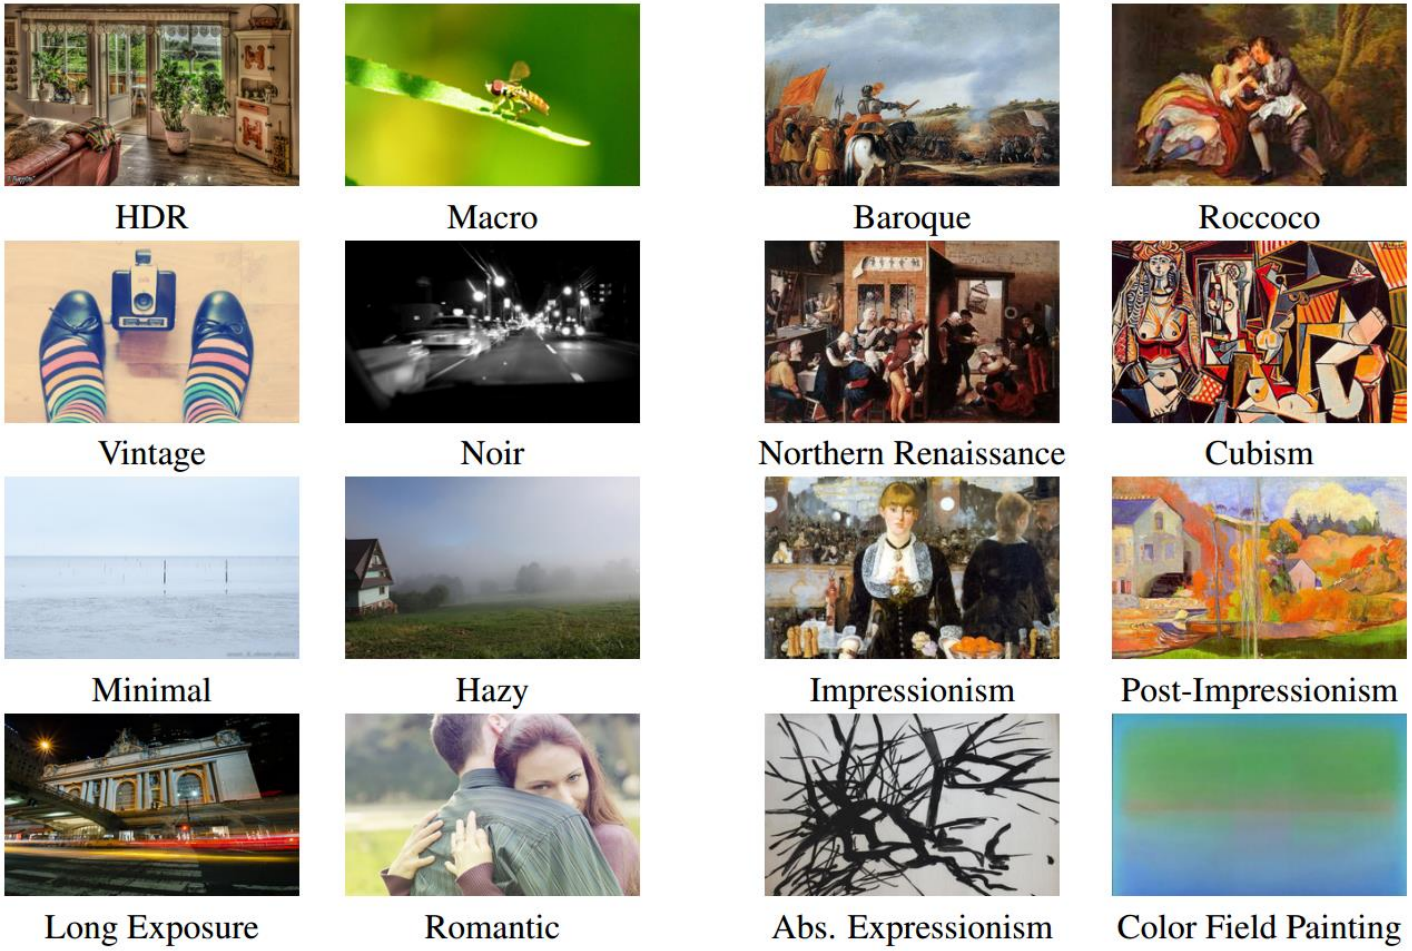

[Karayev et al. BMVC 2014]

Region categorization

- Layout prediction

Input

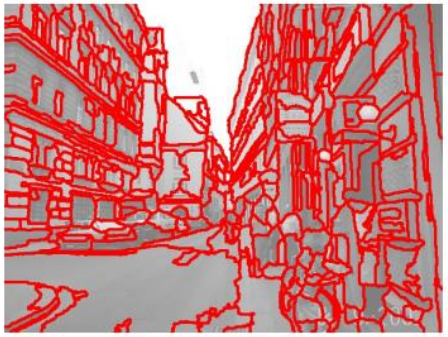

Superpixels

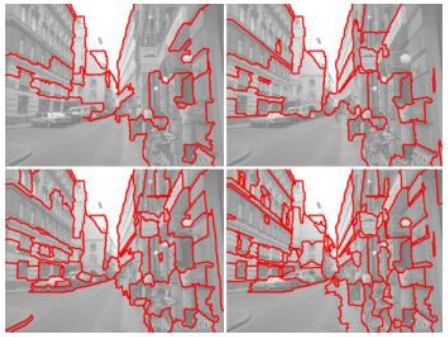

Multiple Segmentations

Surface Layout

Assign regions to orientation

Geometric context [Hoiem et al. IJCV 2007]

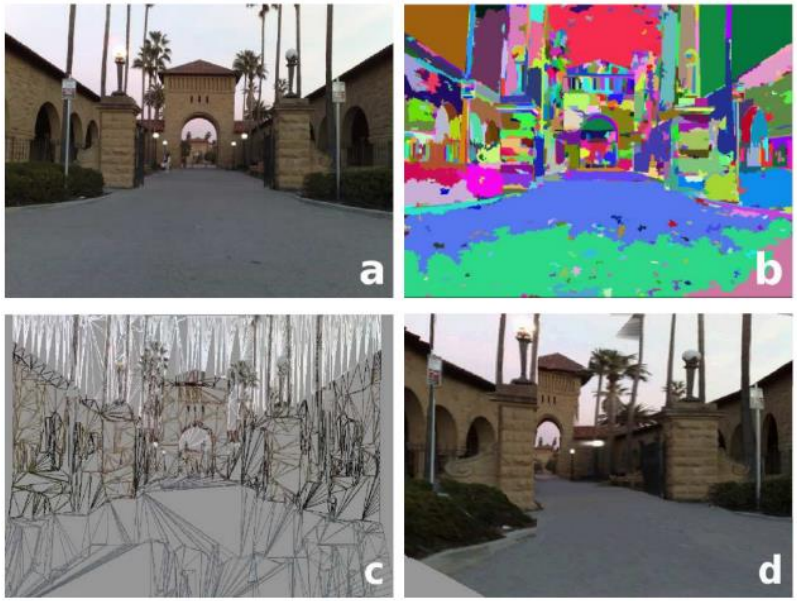

Assign regions to depth

Make3D [Saxena et al. PAMI 2008]

Region categorization

- Semantic segmentation from RGBD images

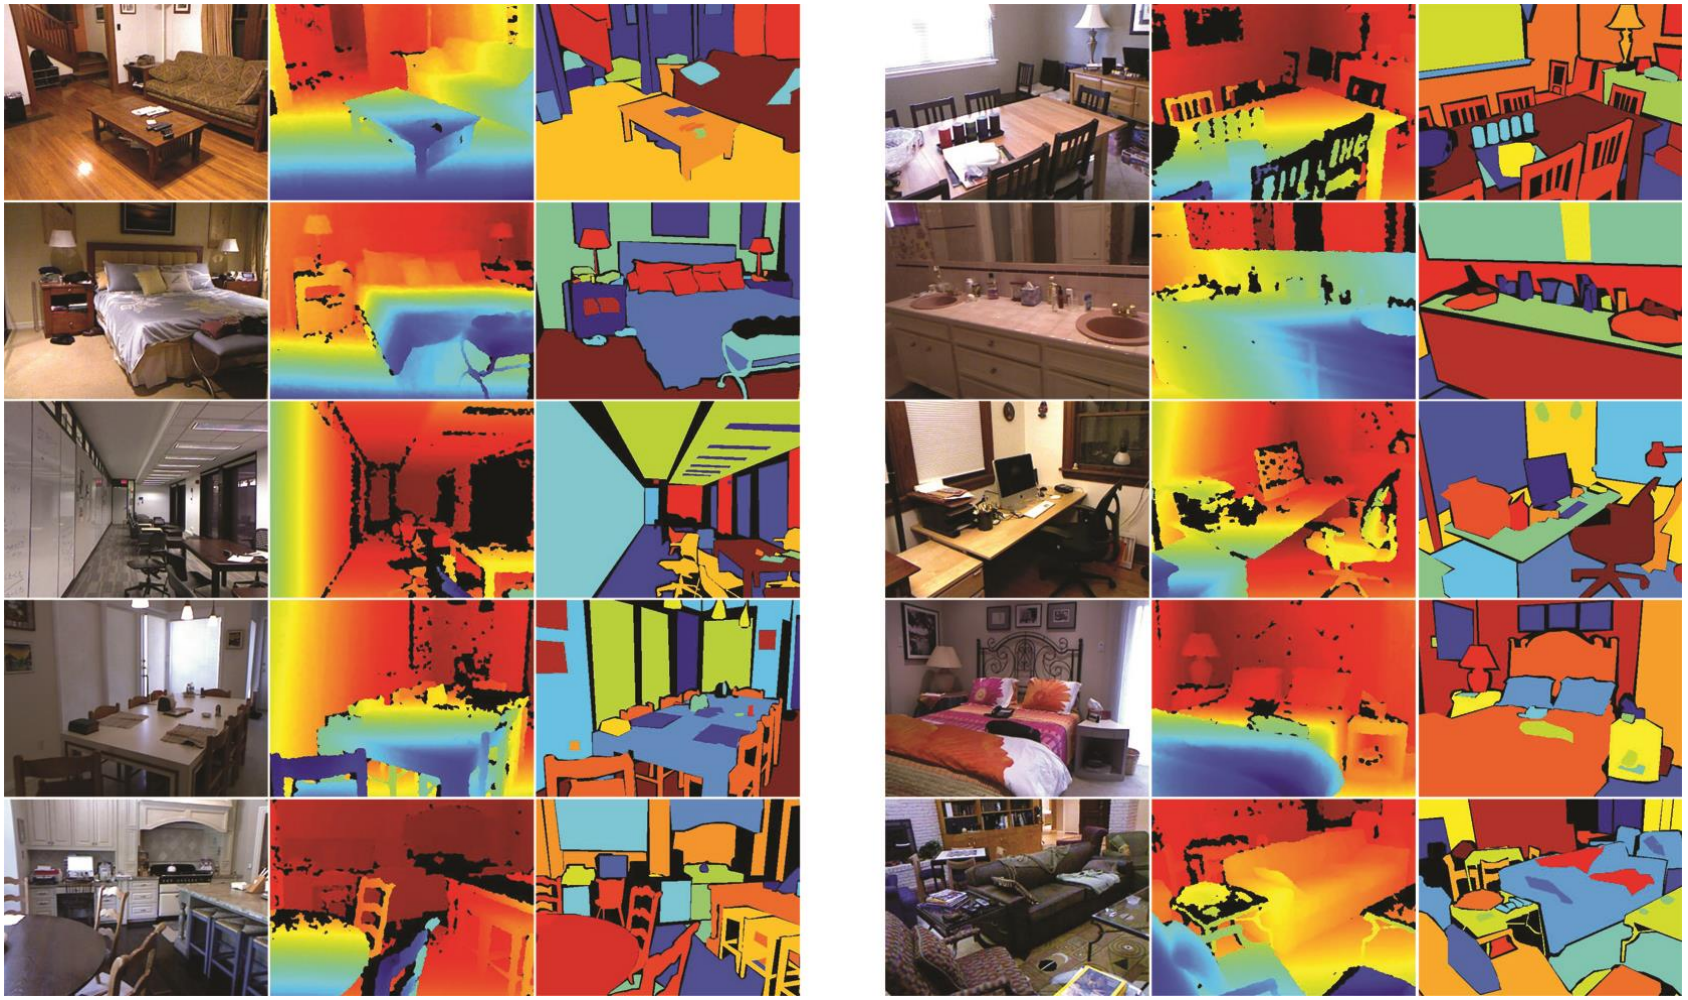

[Silberman et al. ECCV 2012]

Region categorization

- Material recognition

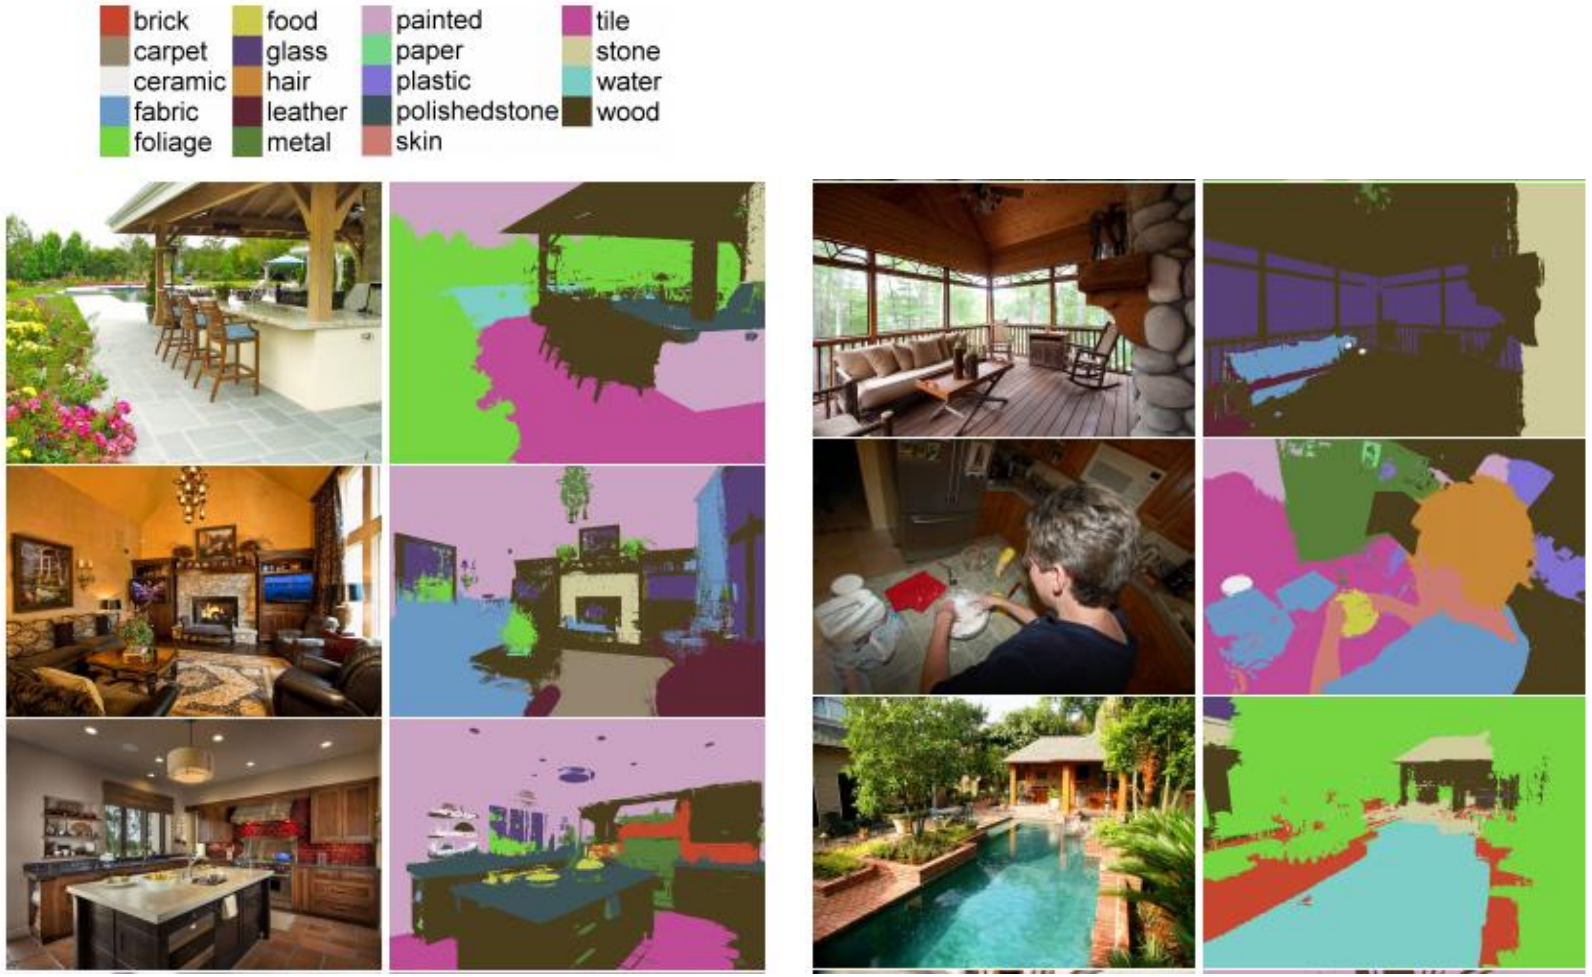

[Bell et al. CVPR 2015]

Supervised learning

Examples

+

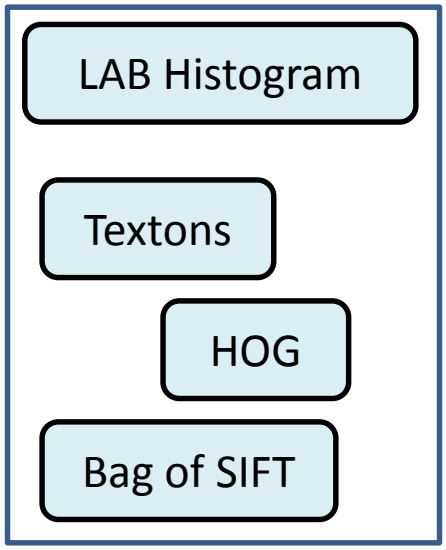

Image Features

+

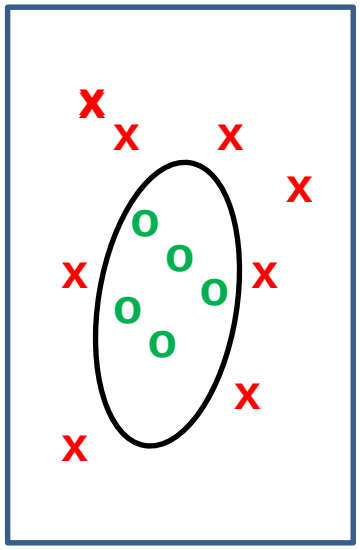

Classifier

= Category label

Training phase

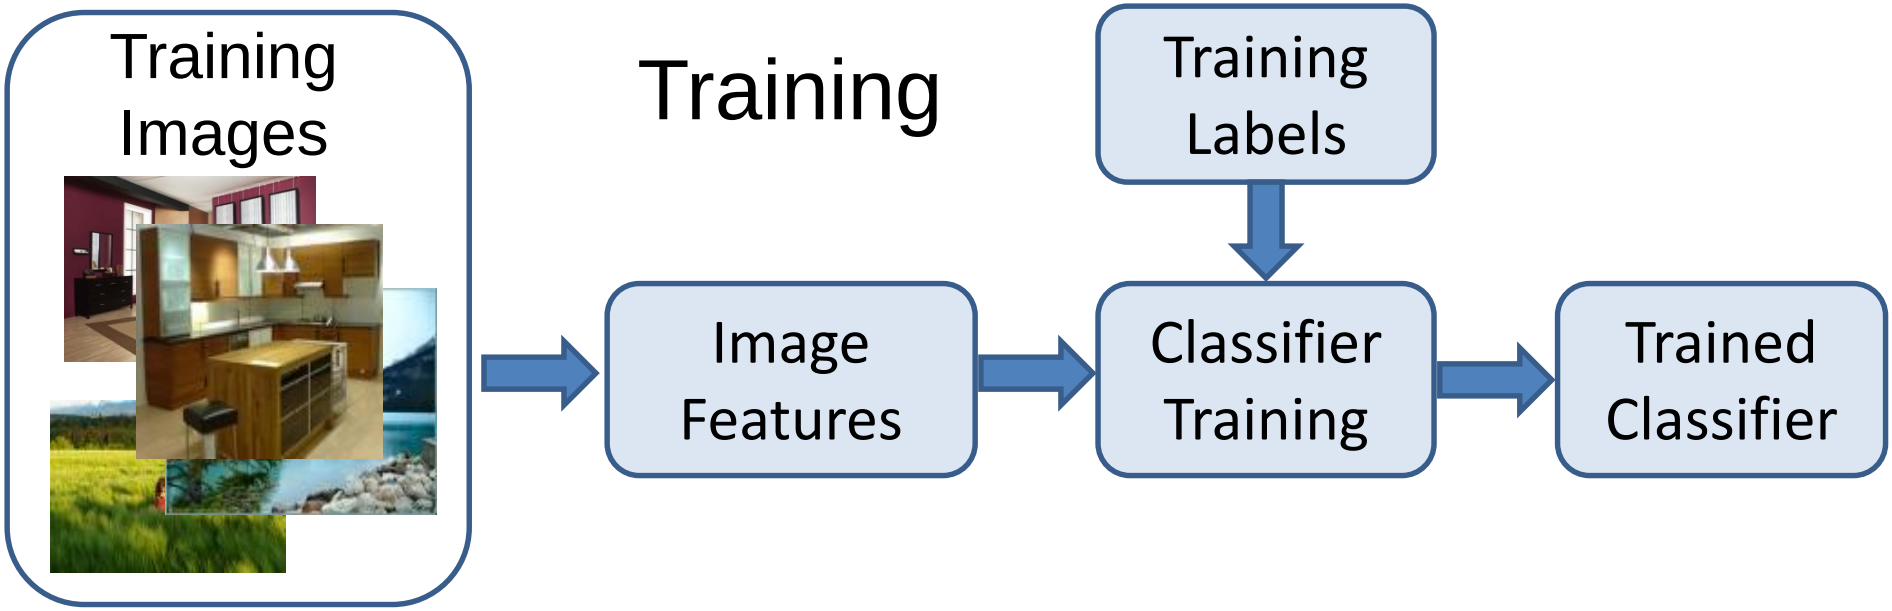

Testing phase

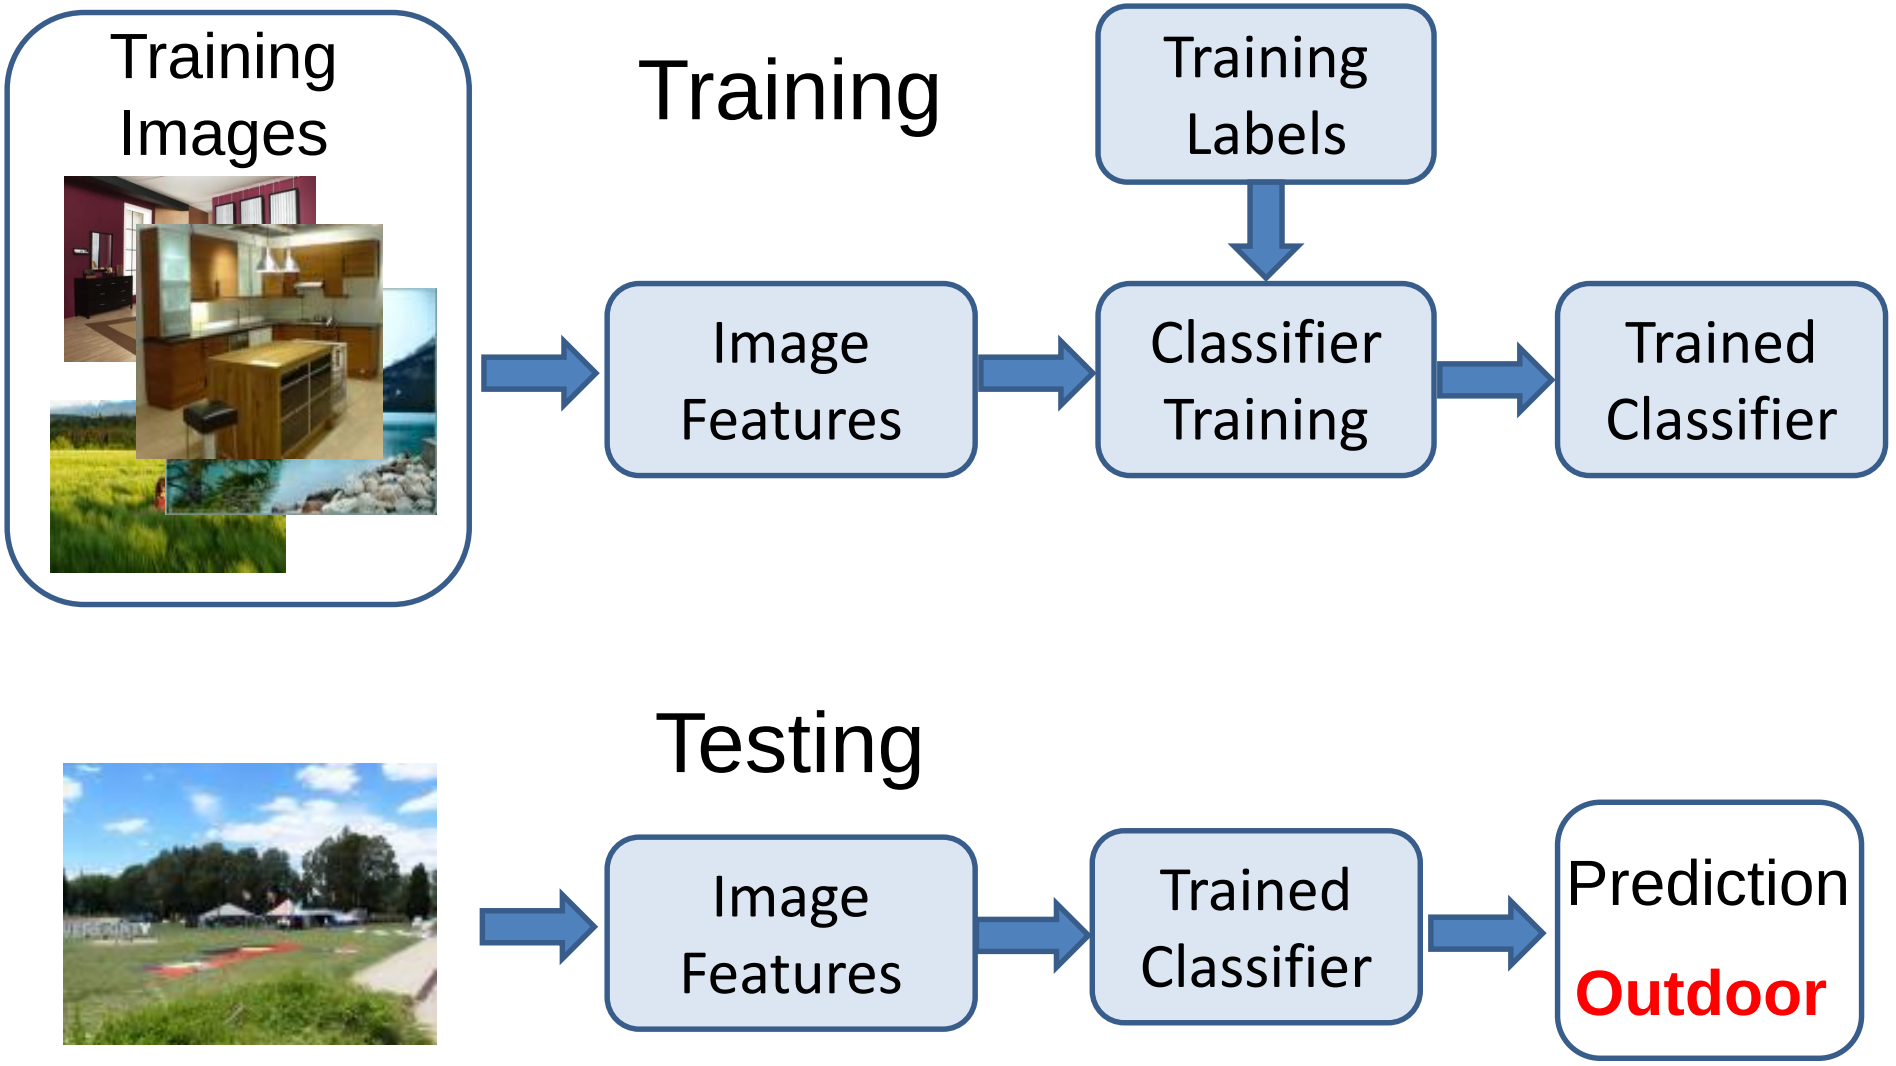

Testing phase

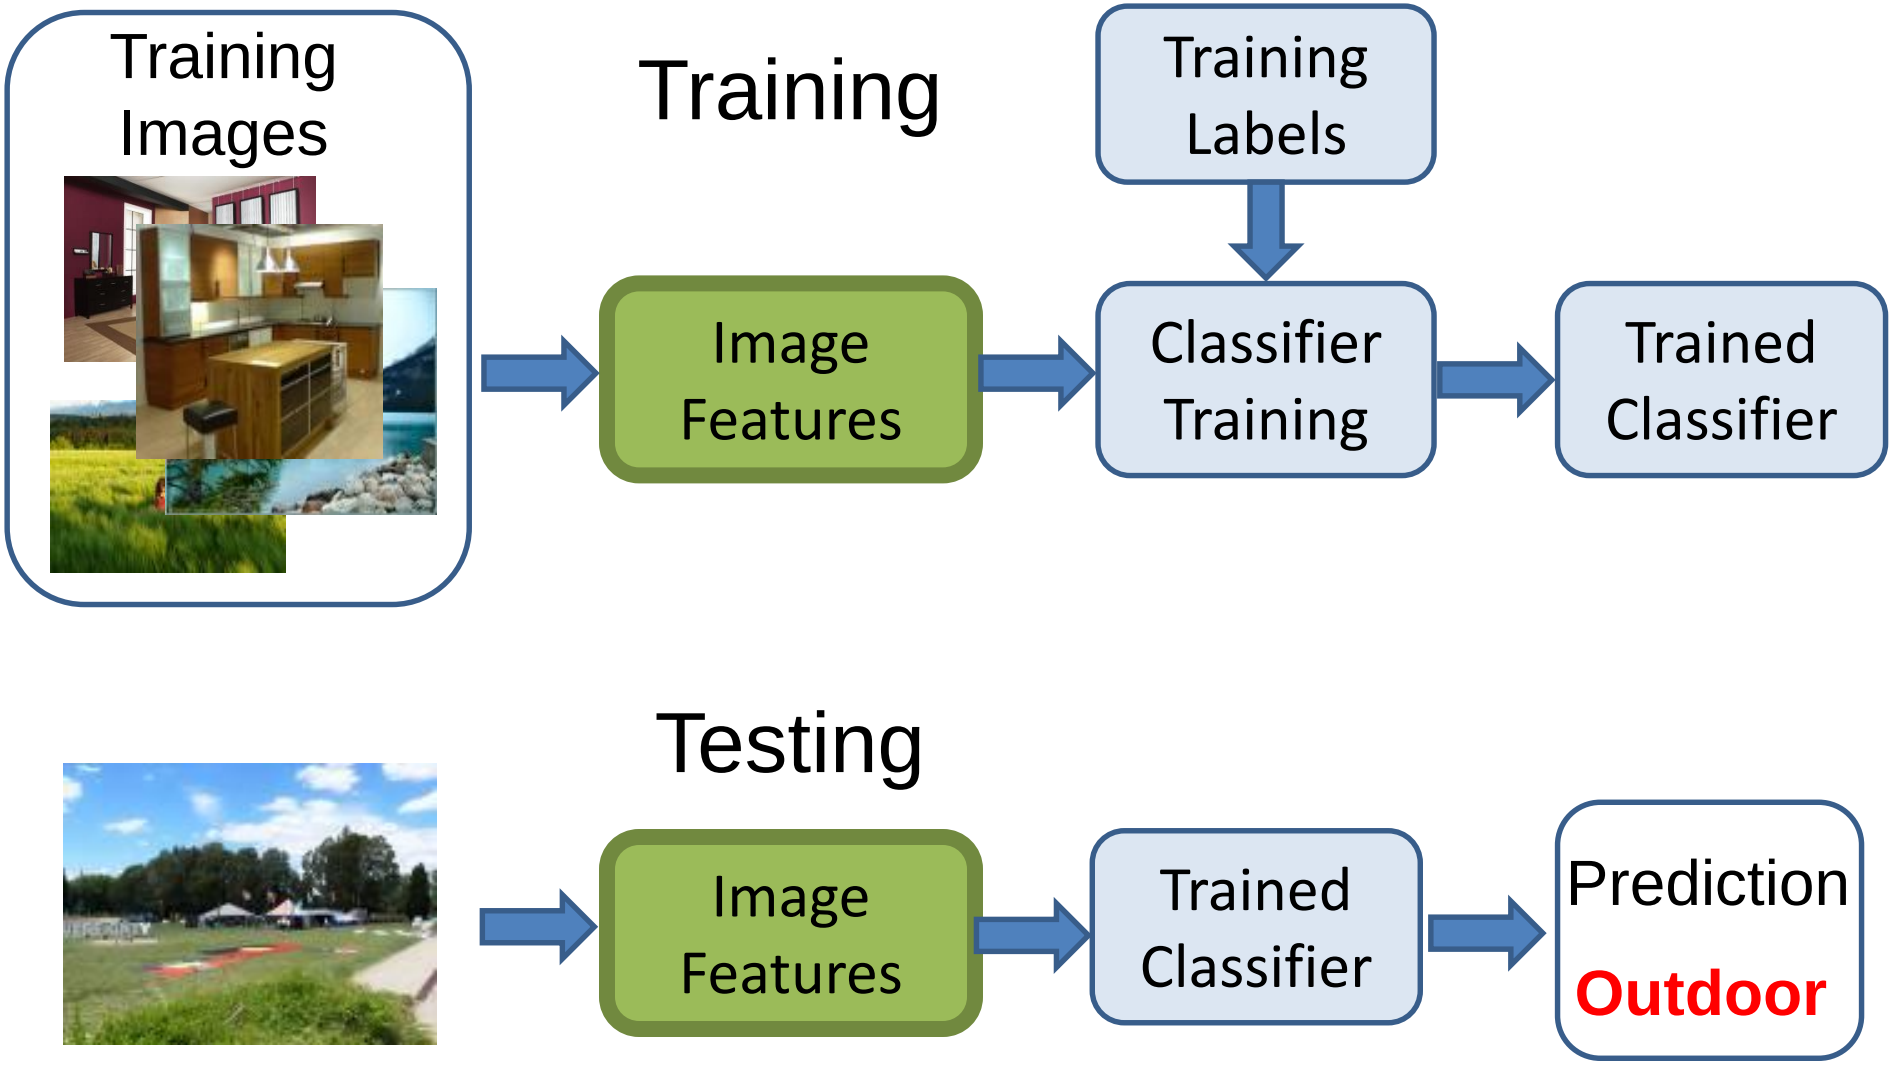

Q: What are good features for…

- recognizing a beach?






Q: What are good features for…

- recognizing cloth fabric?






Q: What are good features for…

- recognizing a mug?






What are the right features?

Depends on what you want to know!

- Object: shape
  - Local shape info, shading, shadows, texture
- Scene : geometric layout
  - linear perspective, gradients, line segments
- Material properties: albedo, feel, hardness
  - Color, texture
- Action: motion
  - Optical flow, tracked points

General principles of representation

- Coverage
  - Ensure that all relevant info is captured
- Concision
  - Minimize number of features without sacrificing coverage
- Directness
  - Ideal features are independently useful for prediction

Image representations

- Templates
  - Intensity, gradients, etc.
- Histograms
  - Color, texture, SIFT descriptors, etc.
- Average of features

Image Intensity

Gradient template

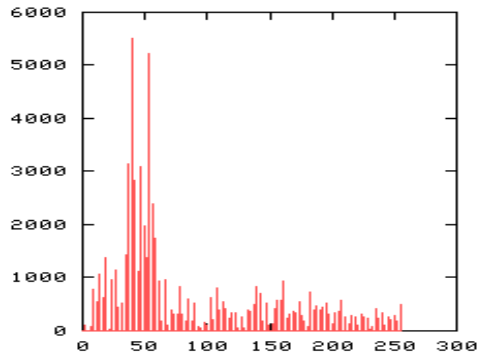

Image representations: histograms

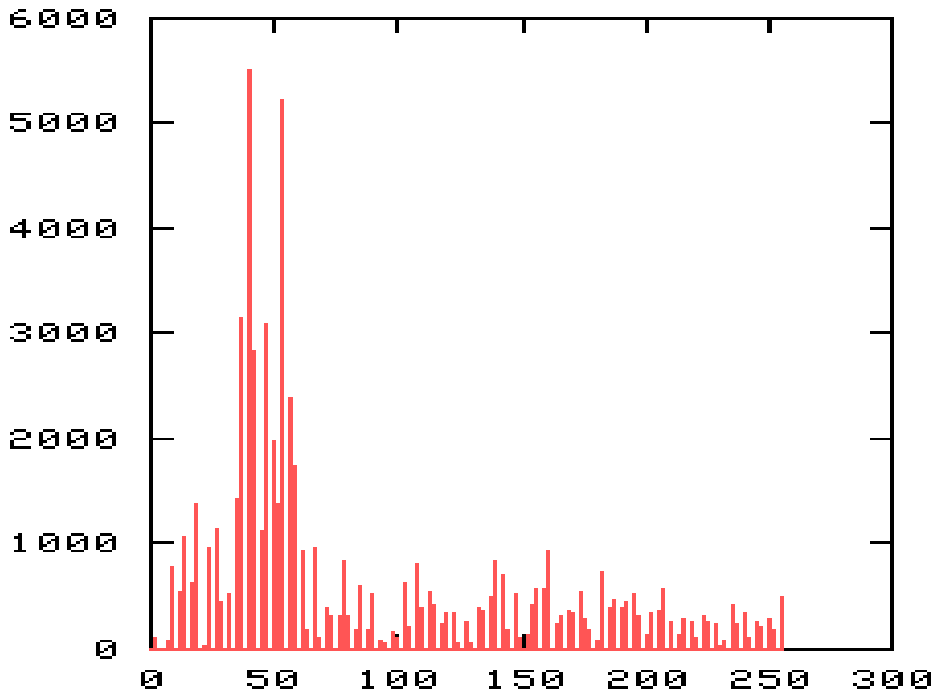

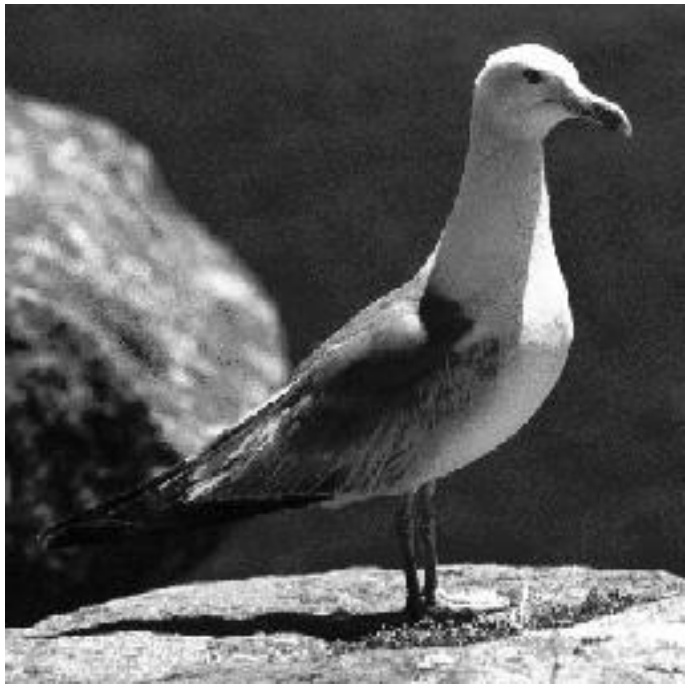

Global histogram

- Represent distribution of features
  - Color, texture, depth, …

Images from Dave Kauchak

Image representations: histograms

- Data samples in 2D

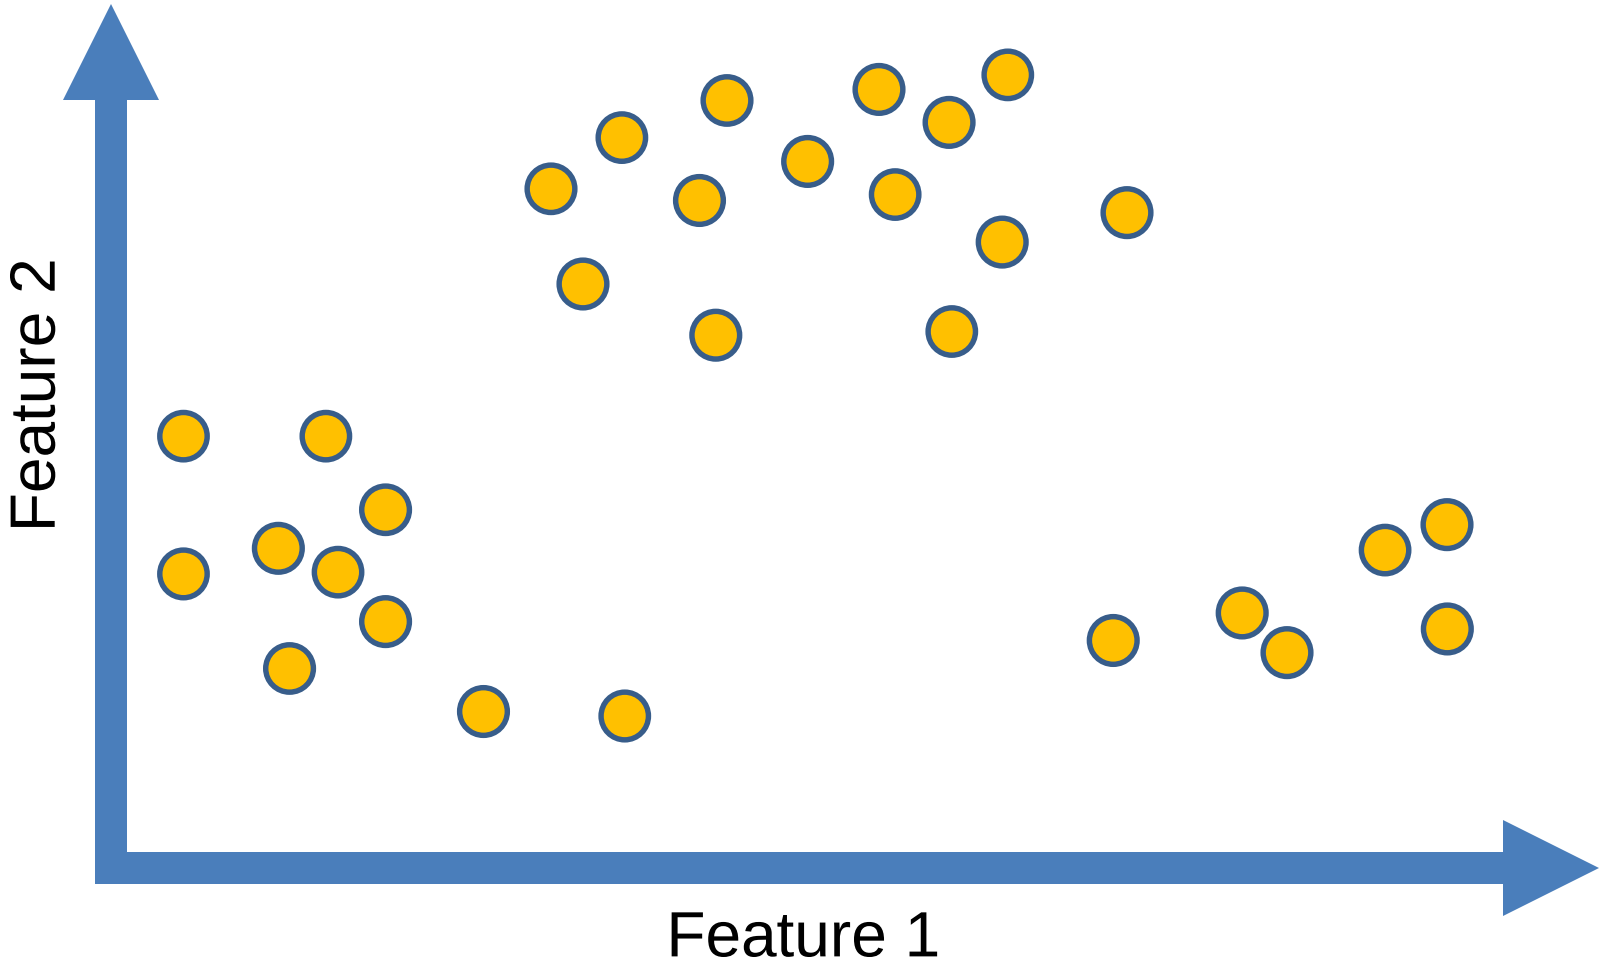

Image representations: histograms

- Probability or count of data in each bin
- Marginal histogram on feature 1

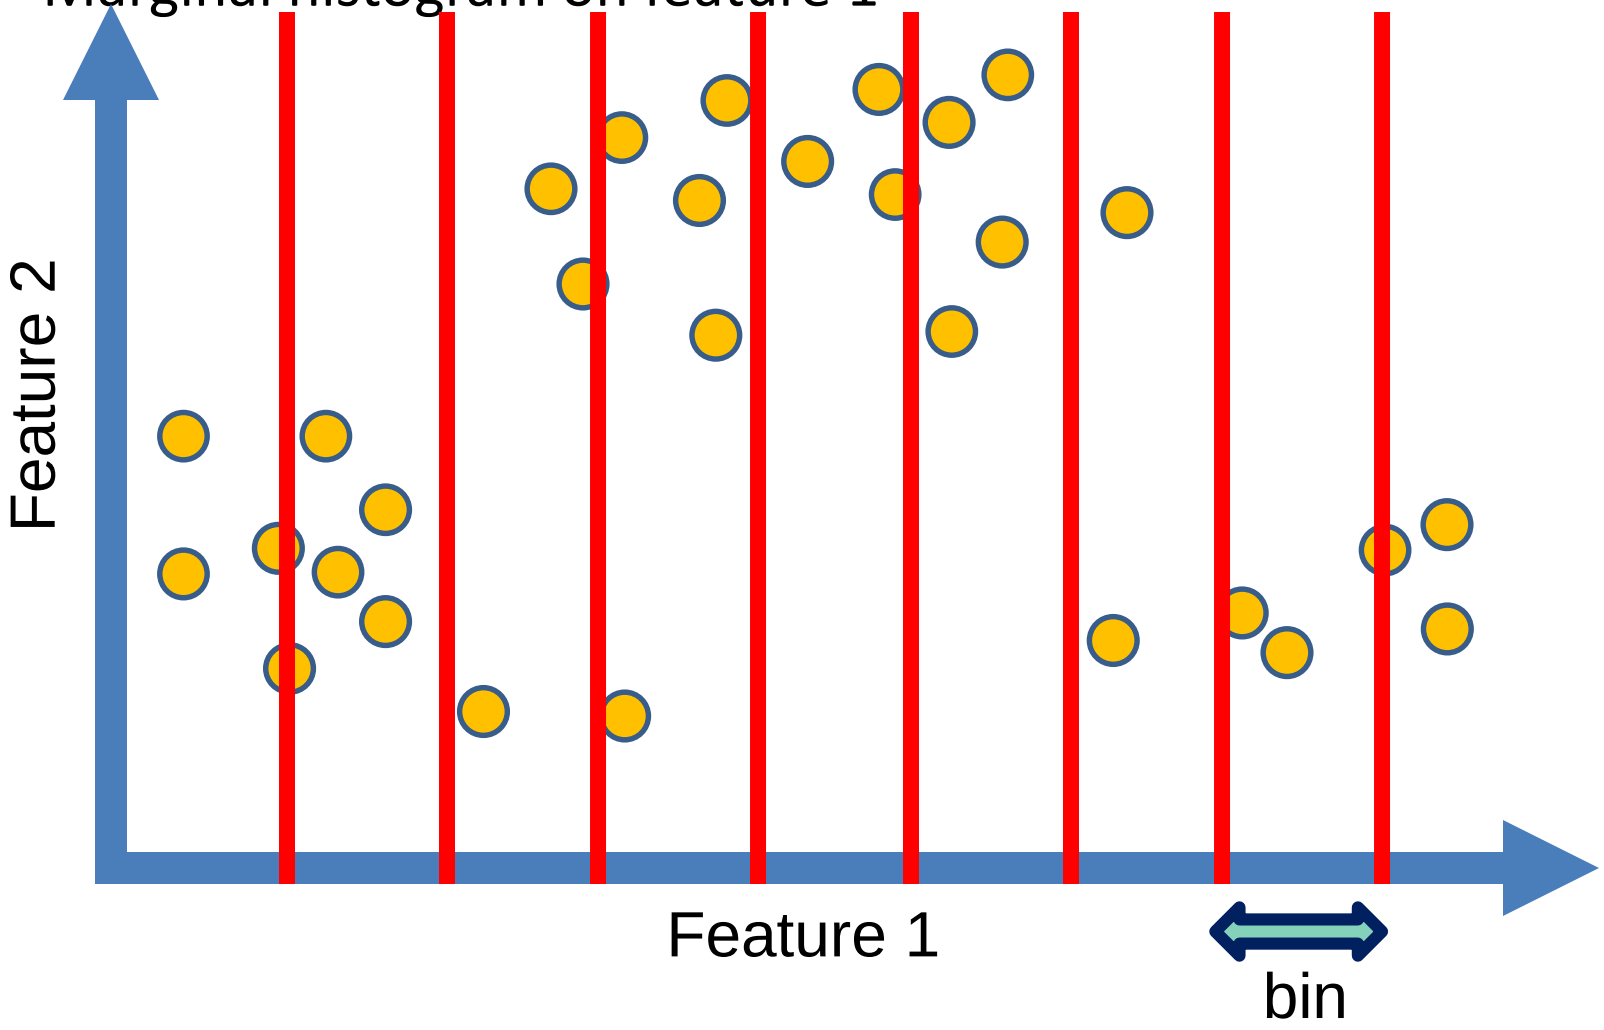

Image representations: histograms

- Marginal histogram on feature 2

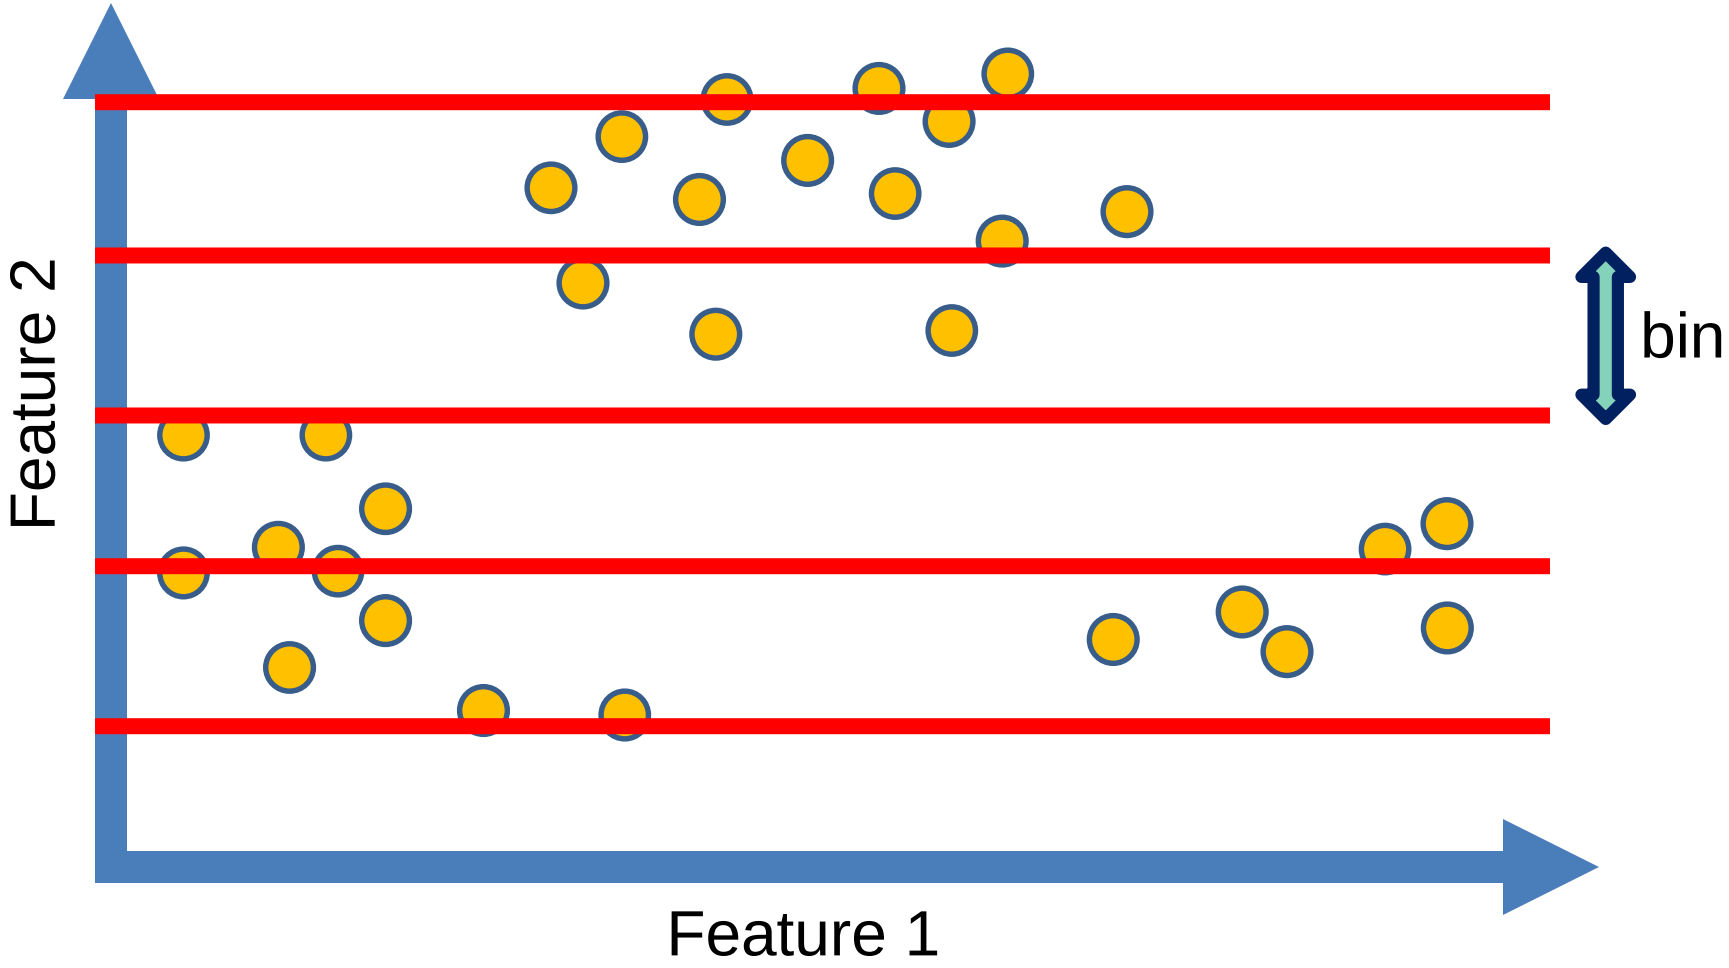

Image representations: histograms

- Joint histogram

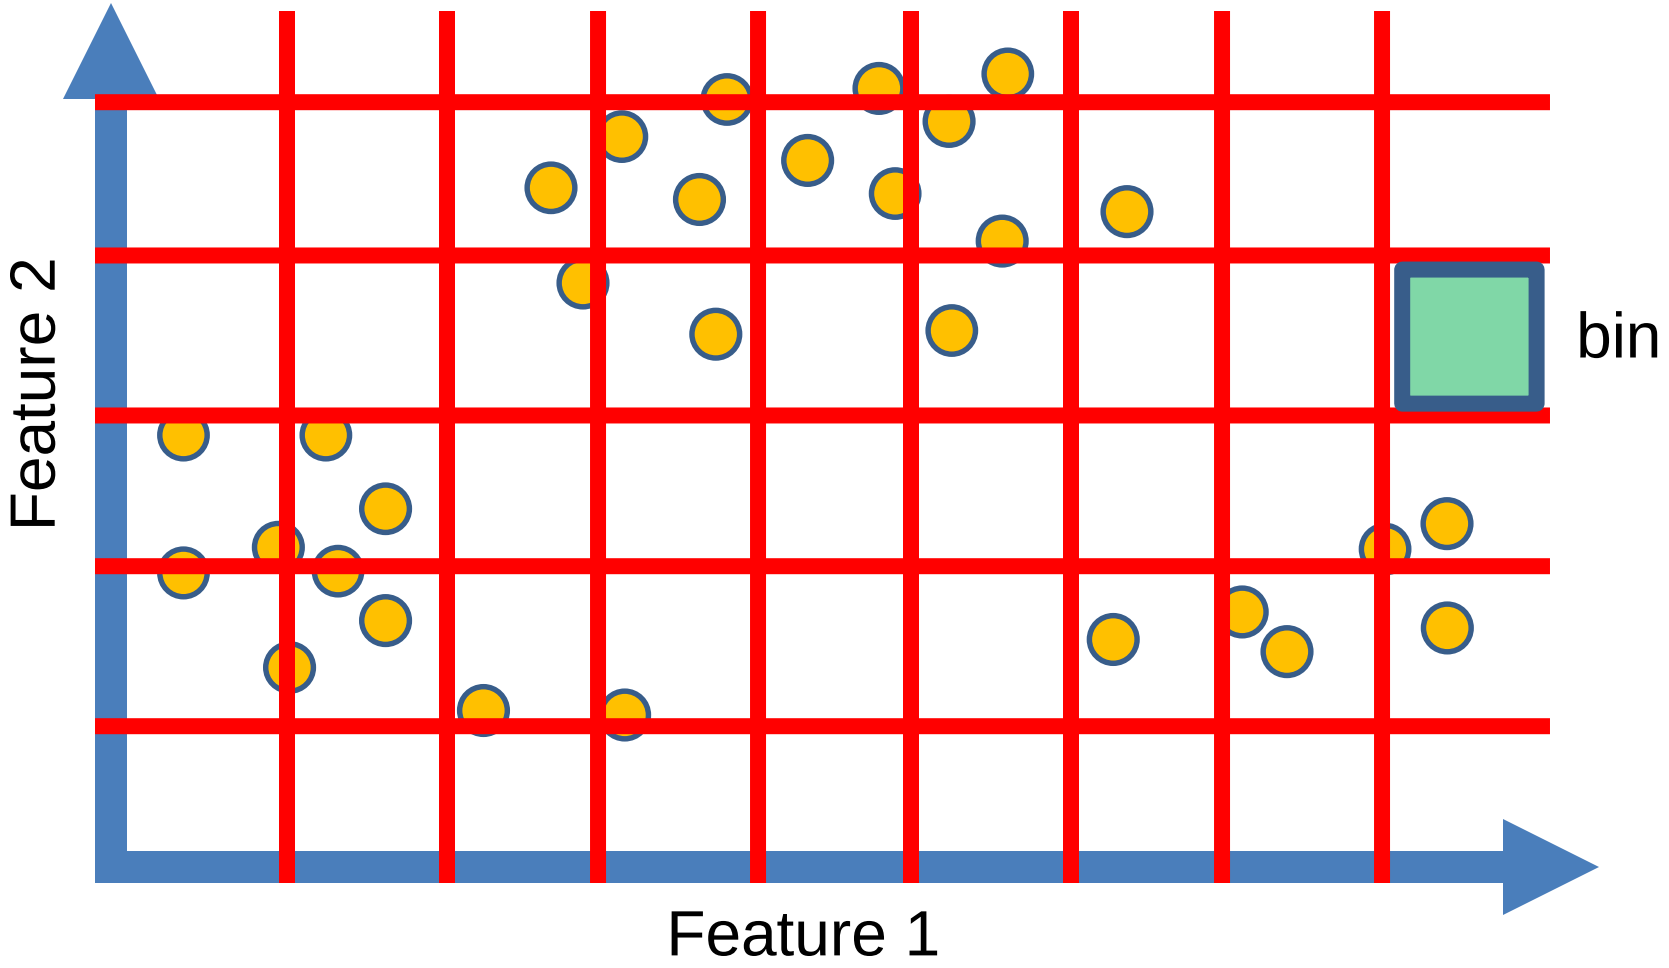

Modeling multi-dimensional data

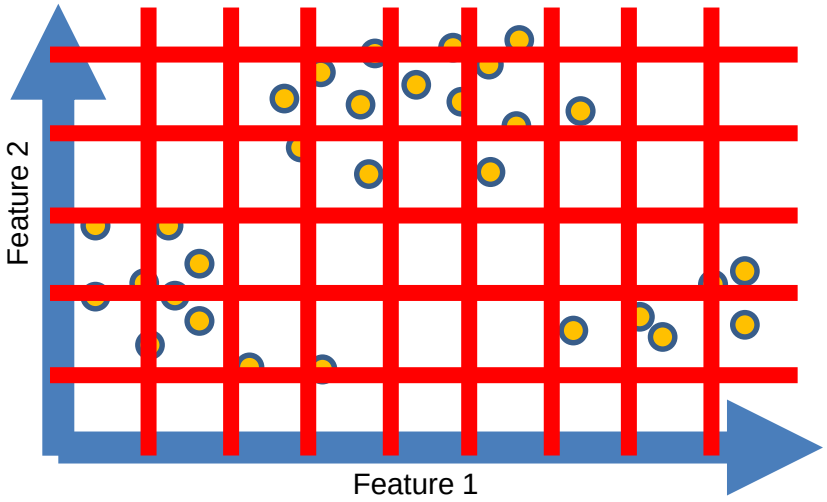

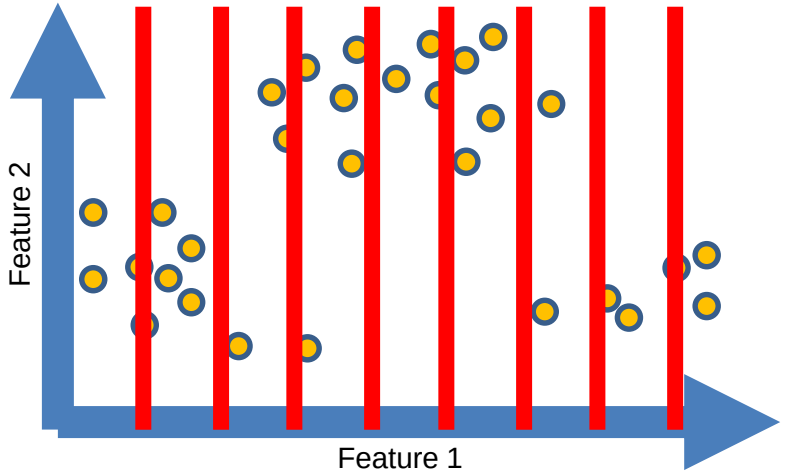

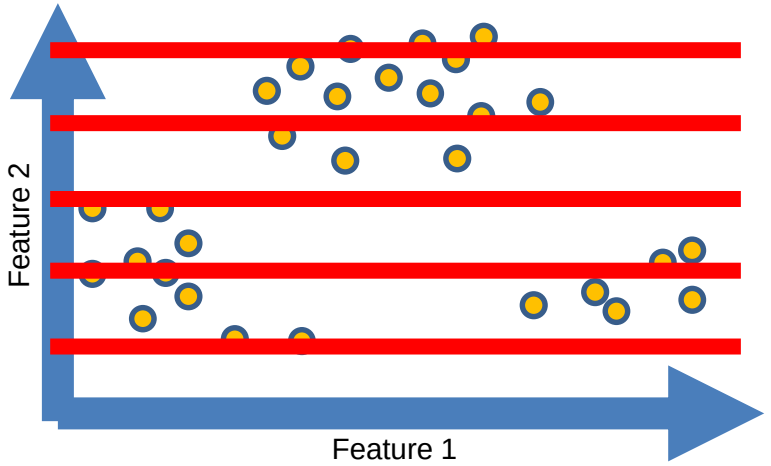

Joint histogram

- Requires lots of data
- Loss of resolution to avoid empty bins

Marginal histogram

- Requires independent features
- More data/bin than joint histogram

Modeling multi-dimensional data

- Clustering
- Use the same cluster centers for all images

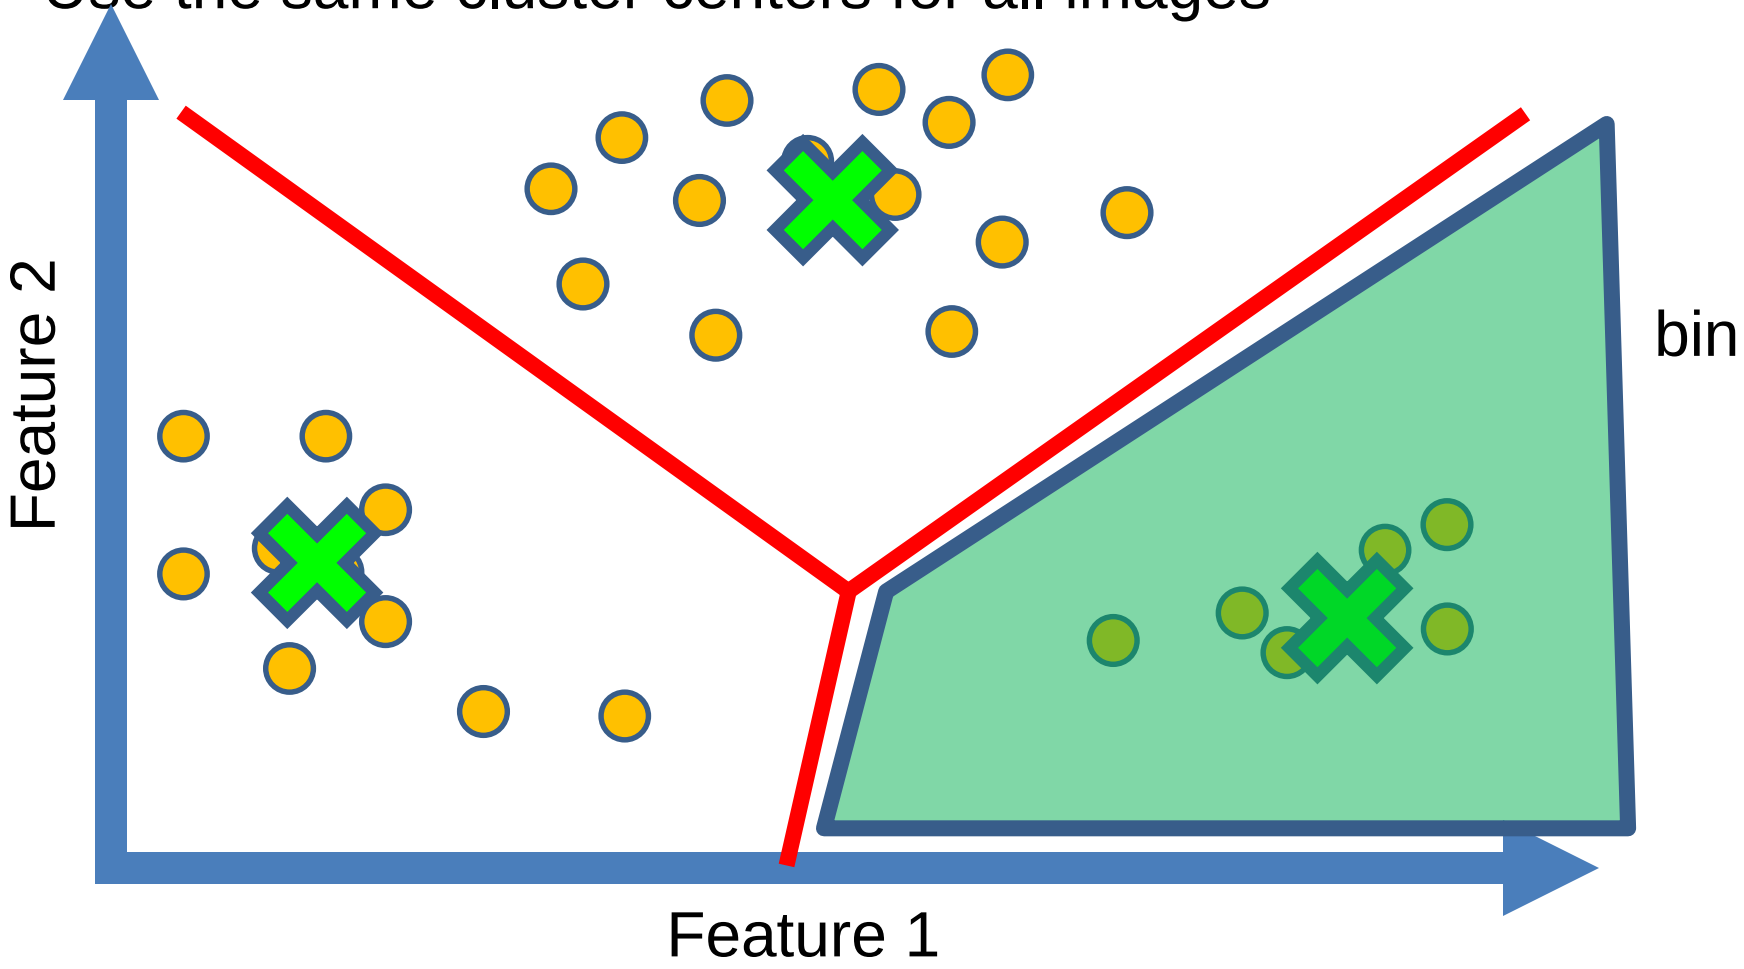

Computing histogram distance

- Histogram intersection

$$\text{histint}(h_i, h_j) = 1 - \sum_{m=1}^{K} \min\left(h_i(m), h_j(m)\right)$$

- Chi-squared Histogram matching distance

$$\chi^2(h_i, h_j) = \frac{1}{2} \sum_{m=1}^{K} \frac{[h_i(m) - h_j(m)]^2}{h_i(m) + h_j(m)}$$

- Earth mover's distance
  (Cross-bin similarity measure)
  - minimal cost paid to transform one distribution into the other

[Rubner et al. The Earth Mover's Distance as a Metric for Image Retrieval, IJCV 2000]

Histograms: implementation issues

- Quantization
  - Grids: fast but applicable only with few dimensions
  - Clustering: slower but can quantize data in higher dimensions

Few Bins
Need less data
Coarser representation

Many Bins
Need more data
Finer representation

- Matching
  - Histogram intersection or Euclidean may be faster
  - Chi-squared often works better
  - Earth mover's distance is good for when nearby bins represent similar values

What kind of things do we compute histograms of?

- Color

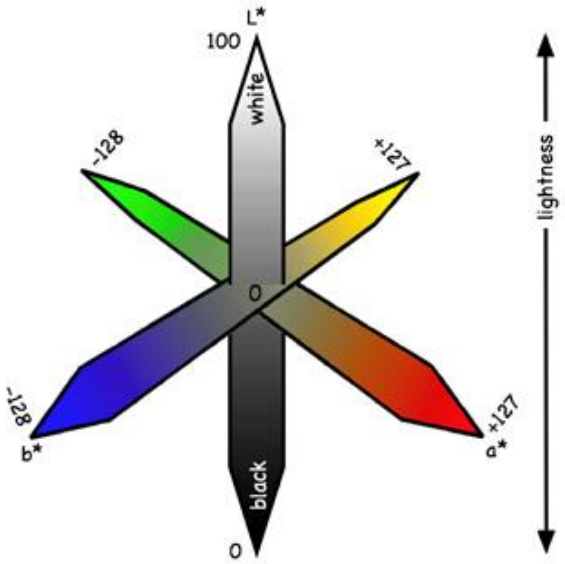

L*a*b* color space

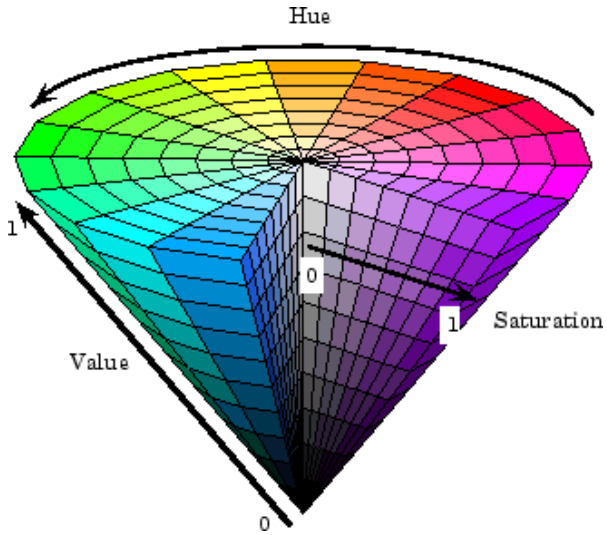

HSV color space

- Texture (filter banks or HOG over regions)

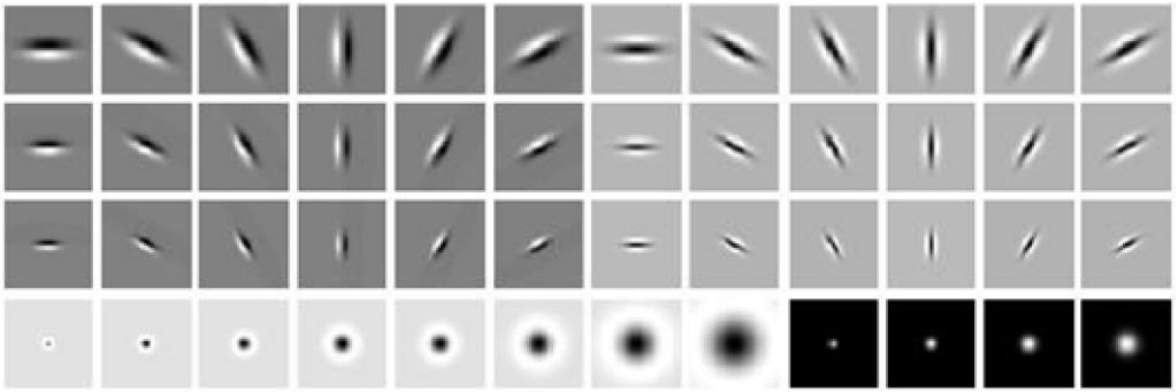

What kind of things do we compute histograms of?

- Histograms of descriptors

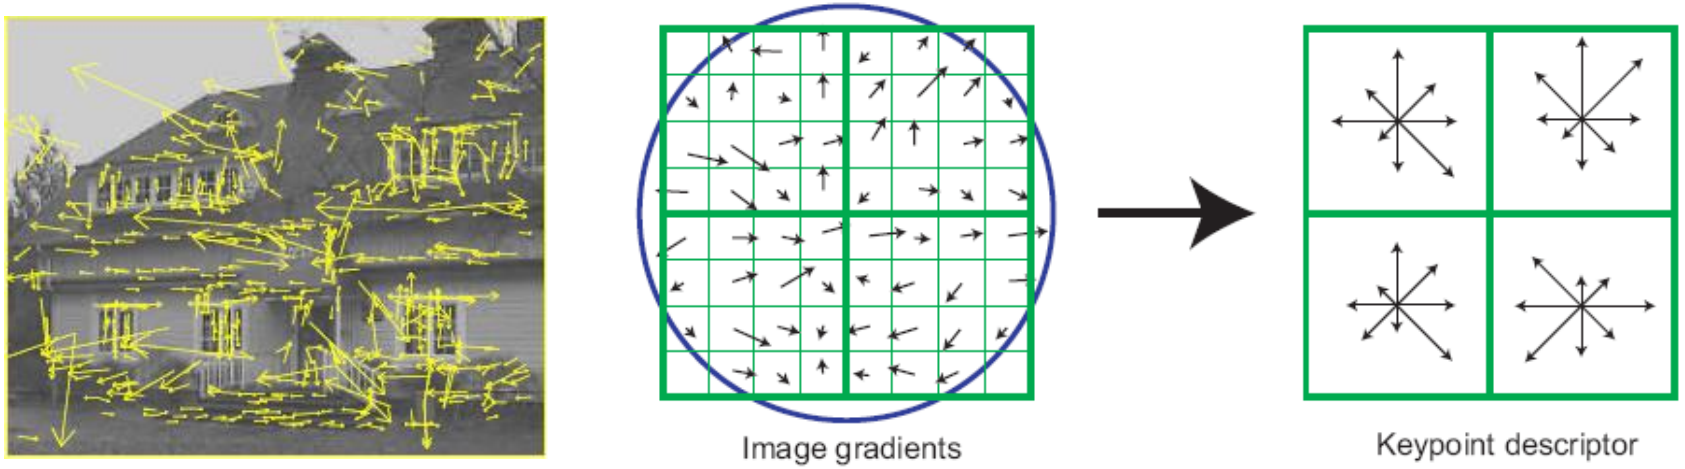

SIFT – [Lowe IJCV 2004]

- “Bag of visual words”

Analogy to documents

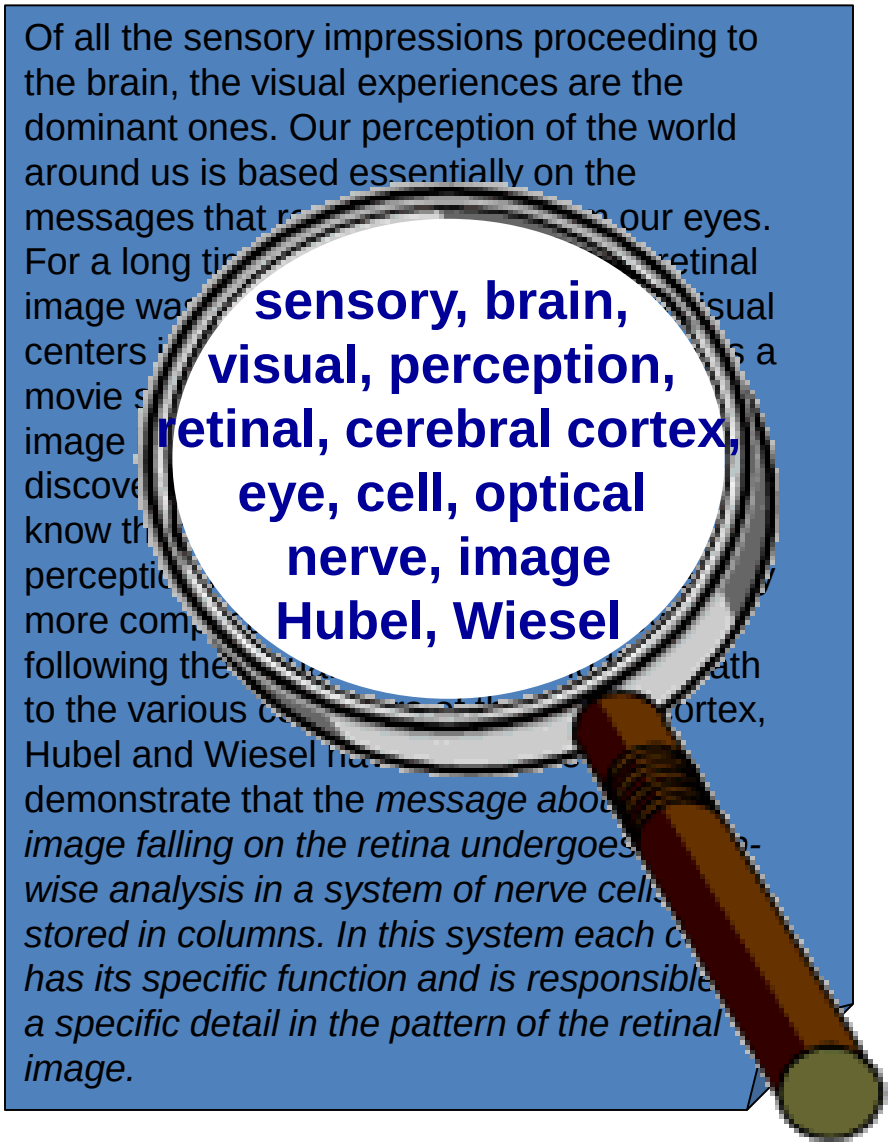

Of all the sensory impressions proceeding to the brain, the visual experiences are the dominant ones. Our perception of the world around us is based essentially on the messages that r… our eyes. For a long ti… retinal image wa… isual centers i… s a movie s… image … discove… know th… perceptio… more com… following the … ath to the various c… ortex, Hubel and Wiesel ha… demonstrate that the *message abou… image falling on the retina undergoes … wise analysis in a system of nerve cell… stored in columns. In this system each c… has its specific function and is responsible… a specific detail in the pattern of the retinal image.*

China is forecasting a trade surplus of $90bn (£51bn) to $100bn this year, a threefold increase on 2004's $32bn. The Commerce Ministry said the surplus would be created by a predicted 30% … $750bn, compared wi… $660bn. T… annoy th… China's … deliber… agrees … yuan is … governo… also need… demand so … country. China … e yuan against the doll… nd permitted it to trade within a narrow … but the US wants the yuan to be allowed … de freely. However, Beijing has made it cl… t it will take its time and tread carefully be… allowing the yuan to rise further in value.

Bag of *visual words*

- Image patches
- BoW histogram
- Codewords

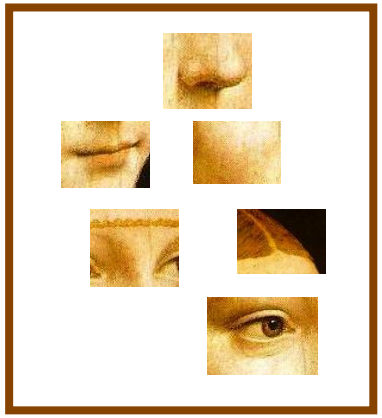

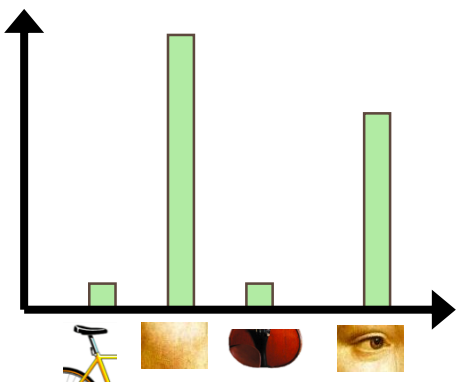

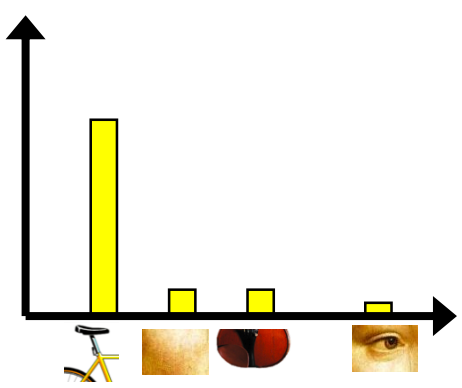

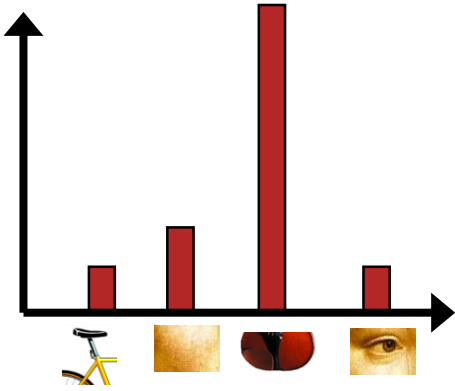

Image categorization with bag of words

Training

1. Extract keypoints and descriptors for all training images
2. Cluster descriptors
3. Quantize descriptors using cluster centers to get “visual words”
4. Represent each image by normalized counts of “visual words”
5. Train classifier on labeled examples using histogram values as features

Testing

1. Extract keypoints/descriptors and quantize into visual words
2. Compute visual word histogram
3. Compute label or confidence using classifier

Bag of visual words image classification

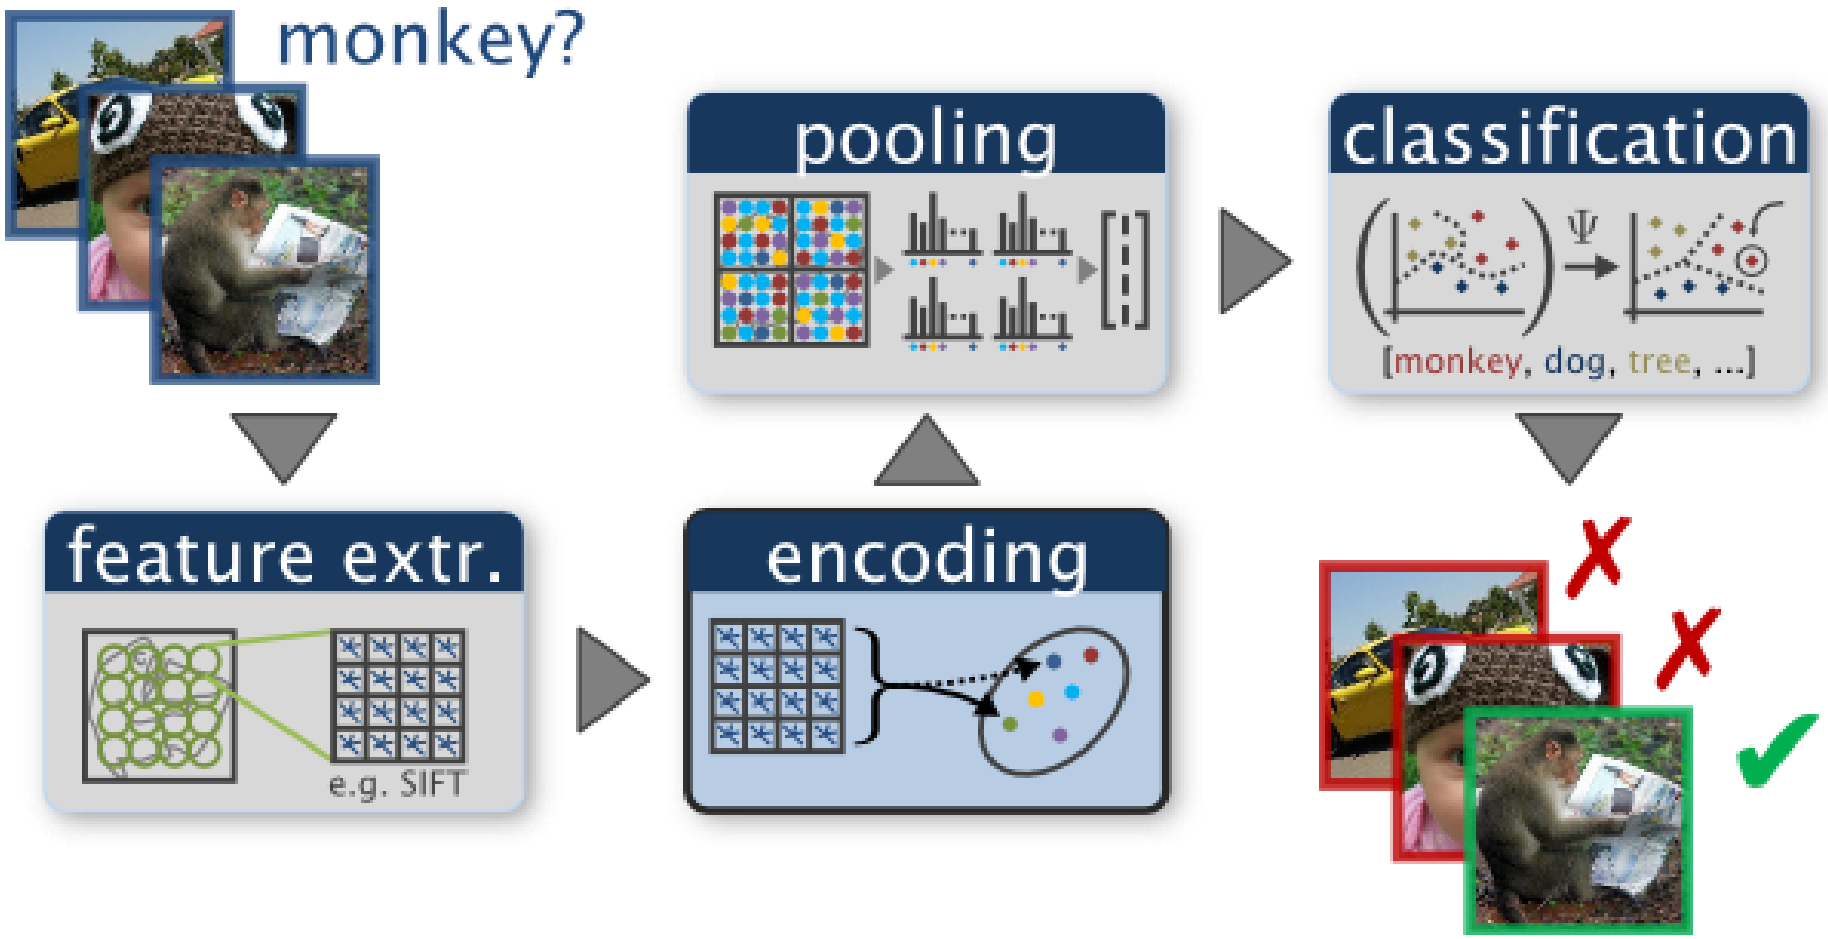

[Chatfieldet al. BMVC 2011]

Feature encoding

- Hard/soft assignment to clusters

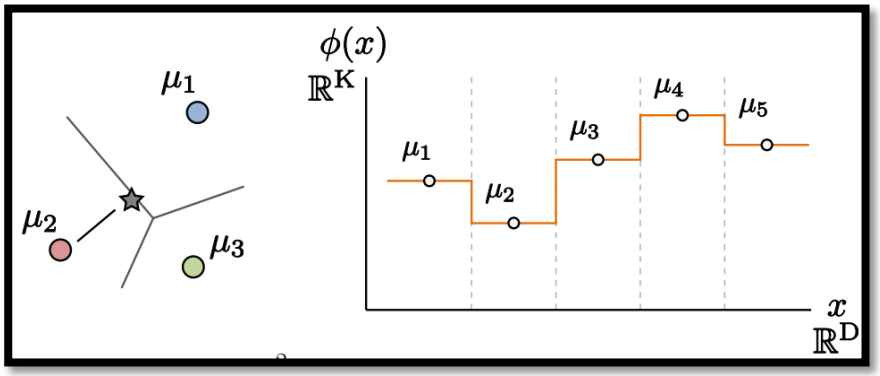

Histogram encoding (hard)

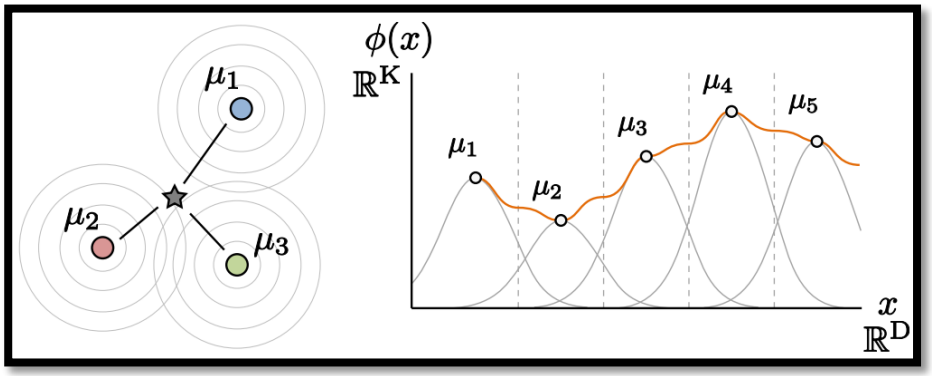

Kernel codebook encoding (soft)

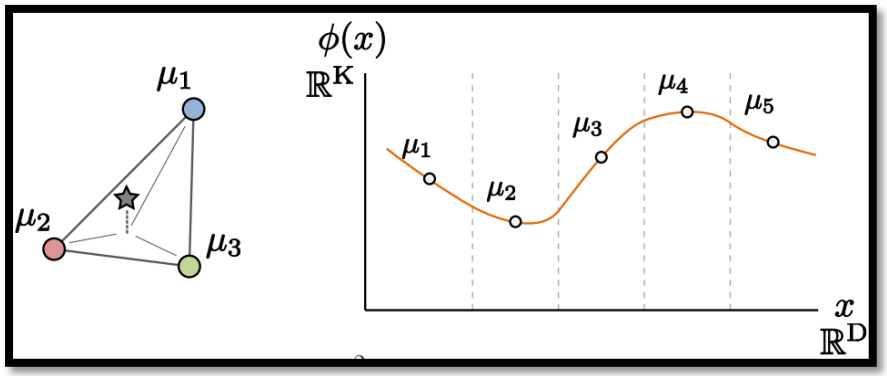

Locality constrained encoding

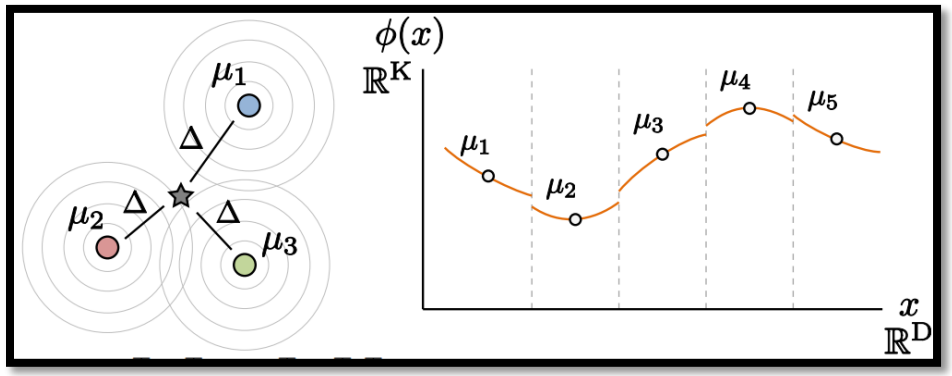

Fisher encoding (1st+2nd statistics)

[Chatfieldet al. BMVC 2011]

BOA vs VLAD vs Fisher

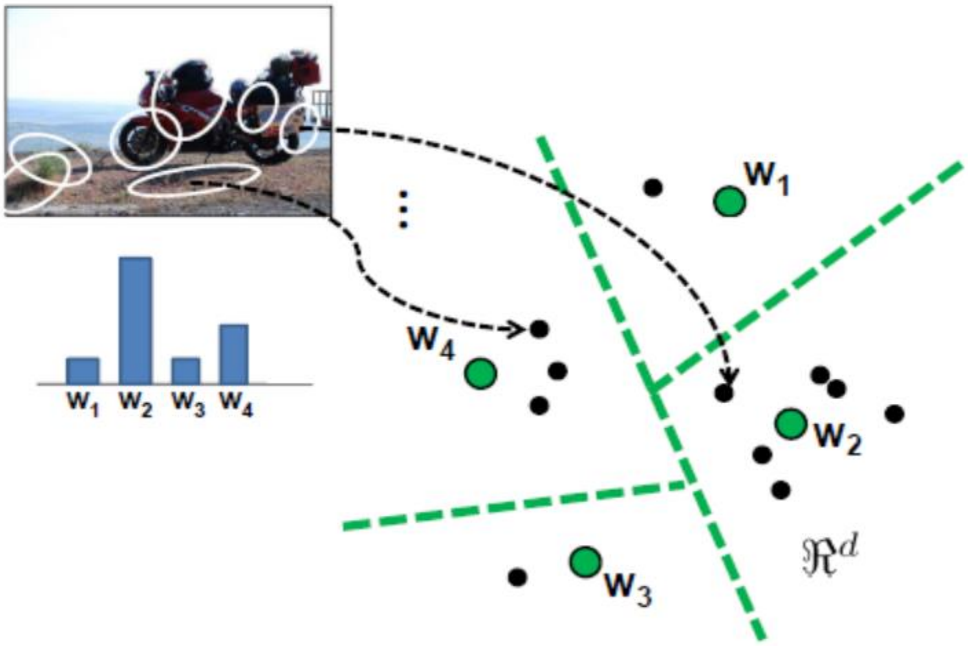

http://www.cs.utexas.edu/~grauman/courses/fall2009/papers/bag_of_visual_words.pdf

BOA vs VLAD vs Fisher

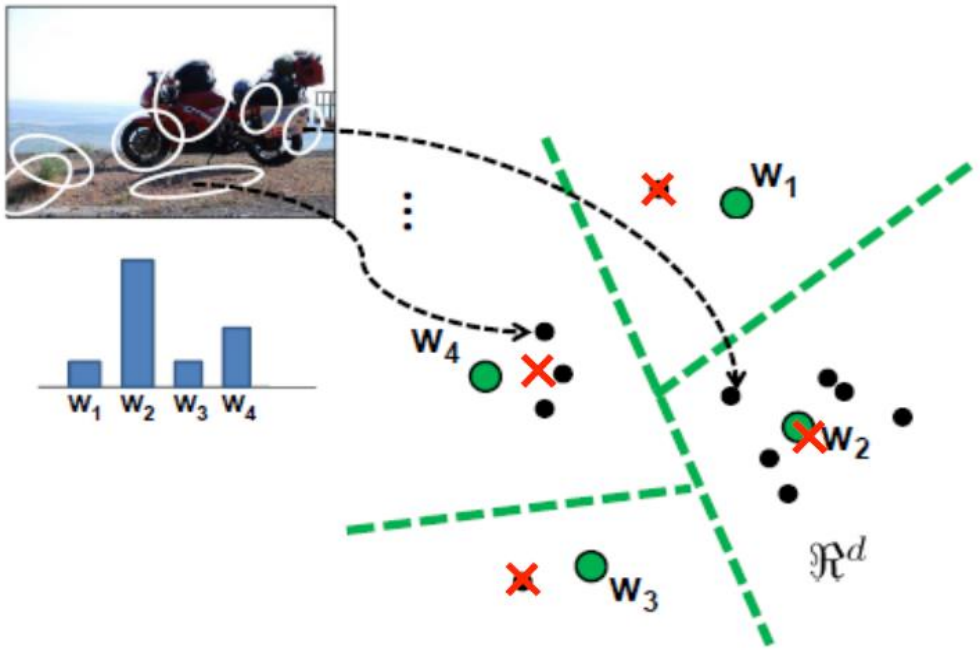

http://www.cs.utexas.edu/~grauman/courses/fall2009/papers/bag_of_visual_words.pdf

BOA vs VLAD vs Fisher

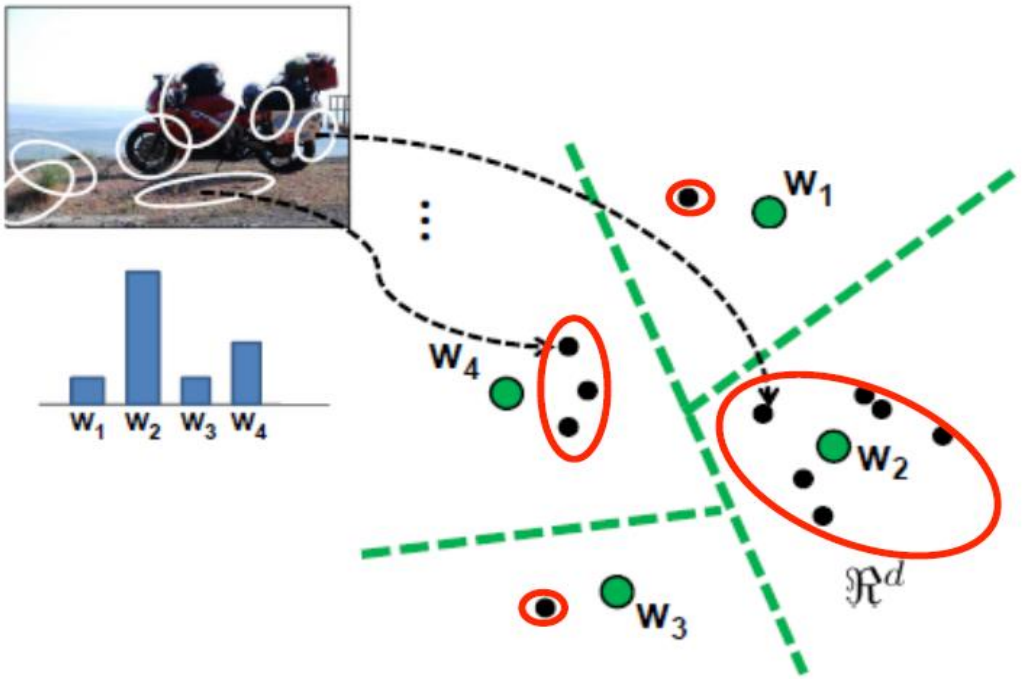

http://www.cs.utexas.edu/~grauman/courses/fall2009/papers/bag_of_visual_words.pdf

Fisher vector encoding

- Fit Gaussian Mixture Models

$$\Theta = (\mu_k, \Sigma_k, \pi_k : k = 1, \ldots, K)$$

- Posterior probability

$$q_{ik} = \frac{\exp\left[-\frac{1}{2}(\mathbf{x}_i - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_i - \mu_k)\right]}{\sum_{t=1}^{K} \exp\left[-\frac{1}{2}(\mathbf{x}_i - \mu_t)^T \Sigma_k^{-1} (\mathbf{x}_i - \mu_t)\right]}$$

- First and second order differences to cluster k

$$u_{jk} = \frac{1}{N\sqrt{\pi_k}} \sum_{i=1}^{N} q_{ik} \frac{x_{ji} - \mu_{jk}}{\sigma_{jk}},$$

$$v_{jk} = \frac{1}{N\sqrt{2\pi_k}} \sum_{i=1}^{N} q_{ik} \left[\left(\frac{x_{ji} - \mu_{jk}}{\sigma_{jk}} \right)^2 - 1 \right]$$

$$\Phi(I) = \begin{bmatrix} \vdots \\ \mathbf{u}_k \\ \vdots \\ \mathbf{v}_k \\ \vdots \end{bmatrix}$$

[Perronnin et al. ECCV 2010]

Performance comparisons

- Fisher vector encoding outperforms others
- Higher-order statistics helps

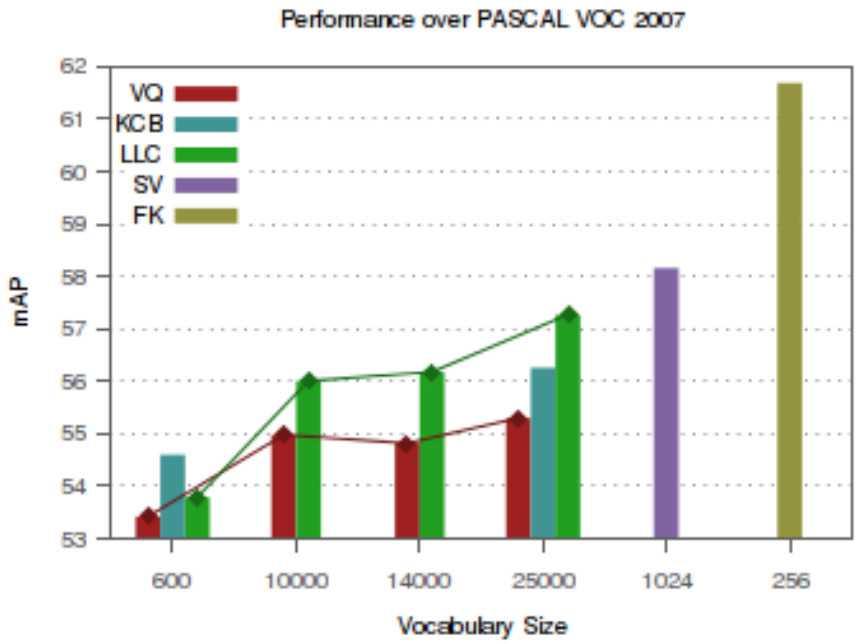

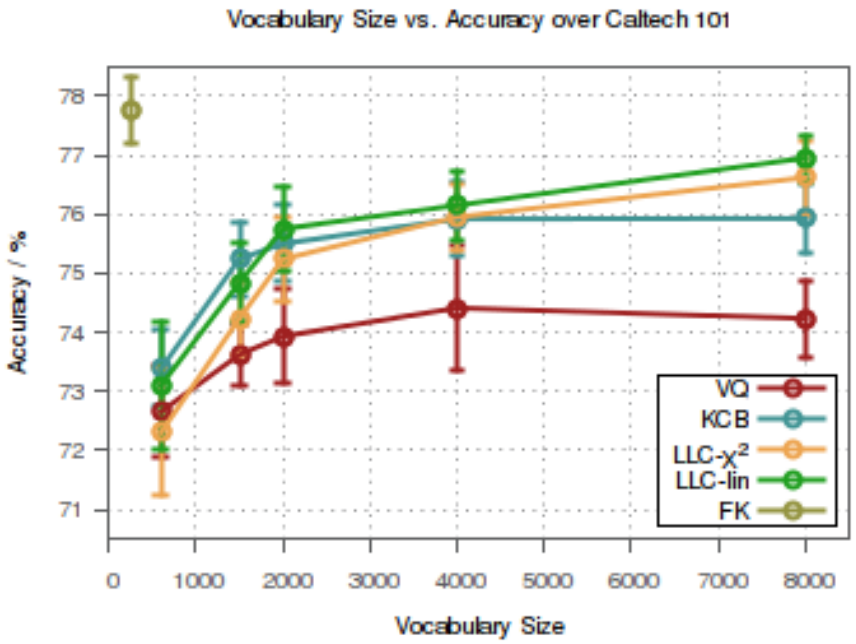

[Chatfieldet al. BMVC 2011]

But what about spatial layout?

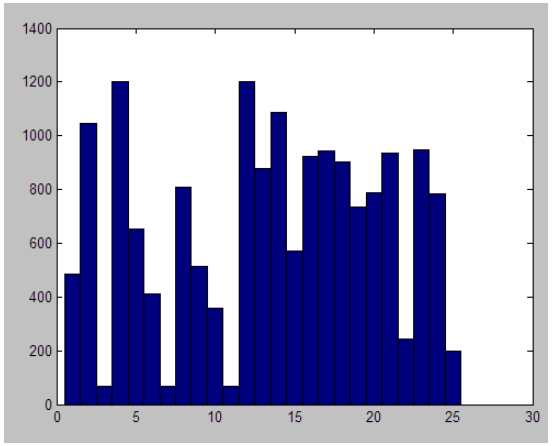

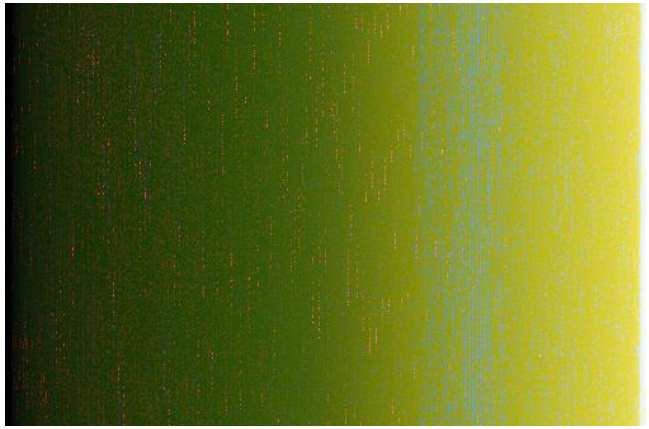

All of these images have the same color histogram

Spatial pyramid pooling

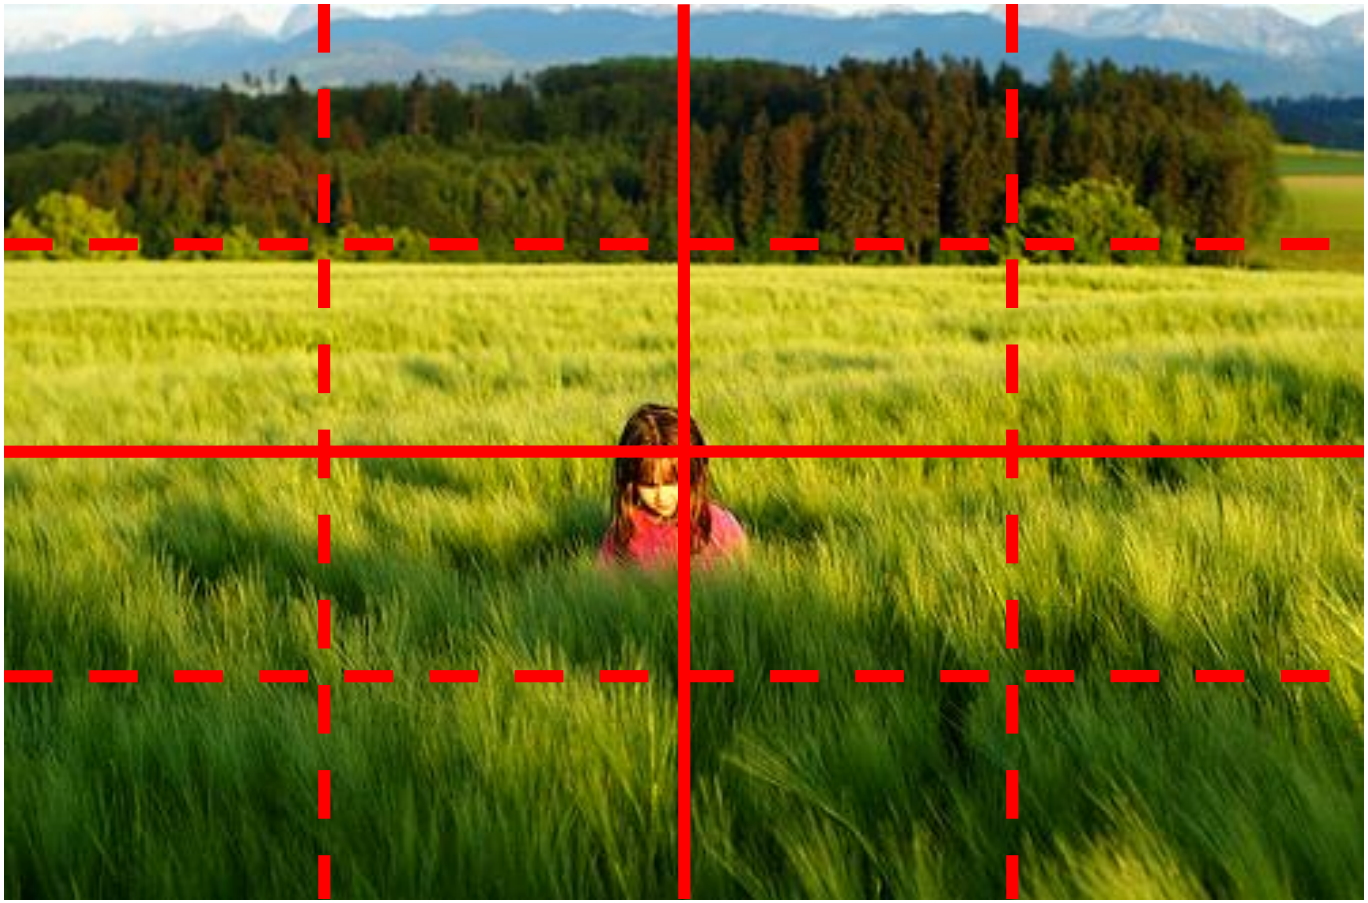

Compute histogram in each spatial bin

Spatial pyramid pooling

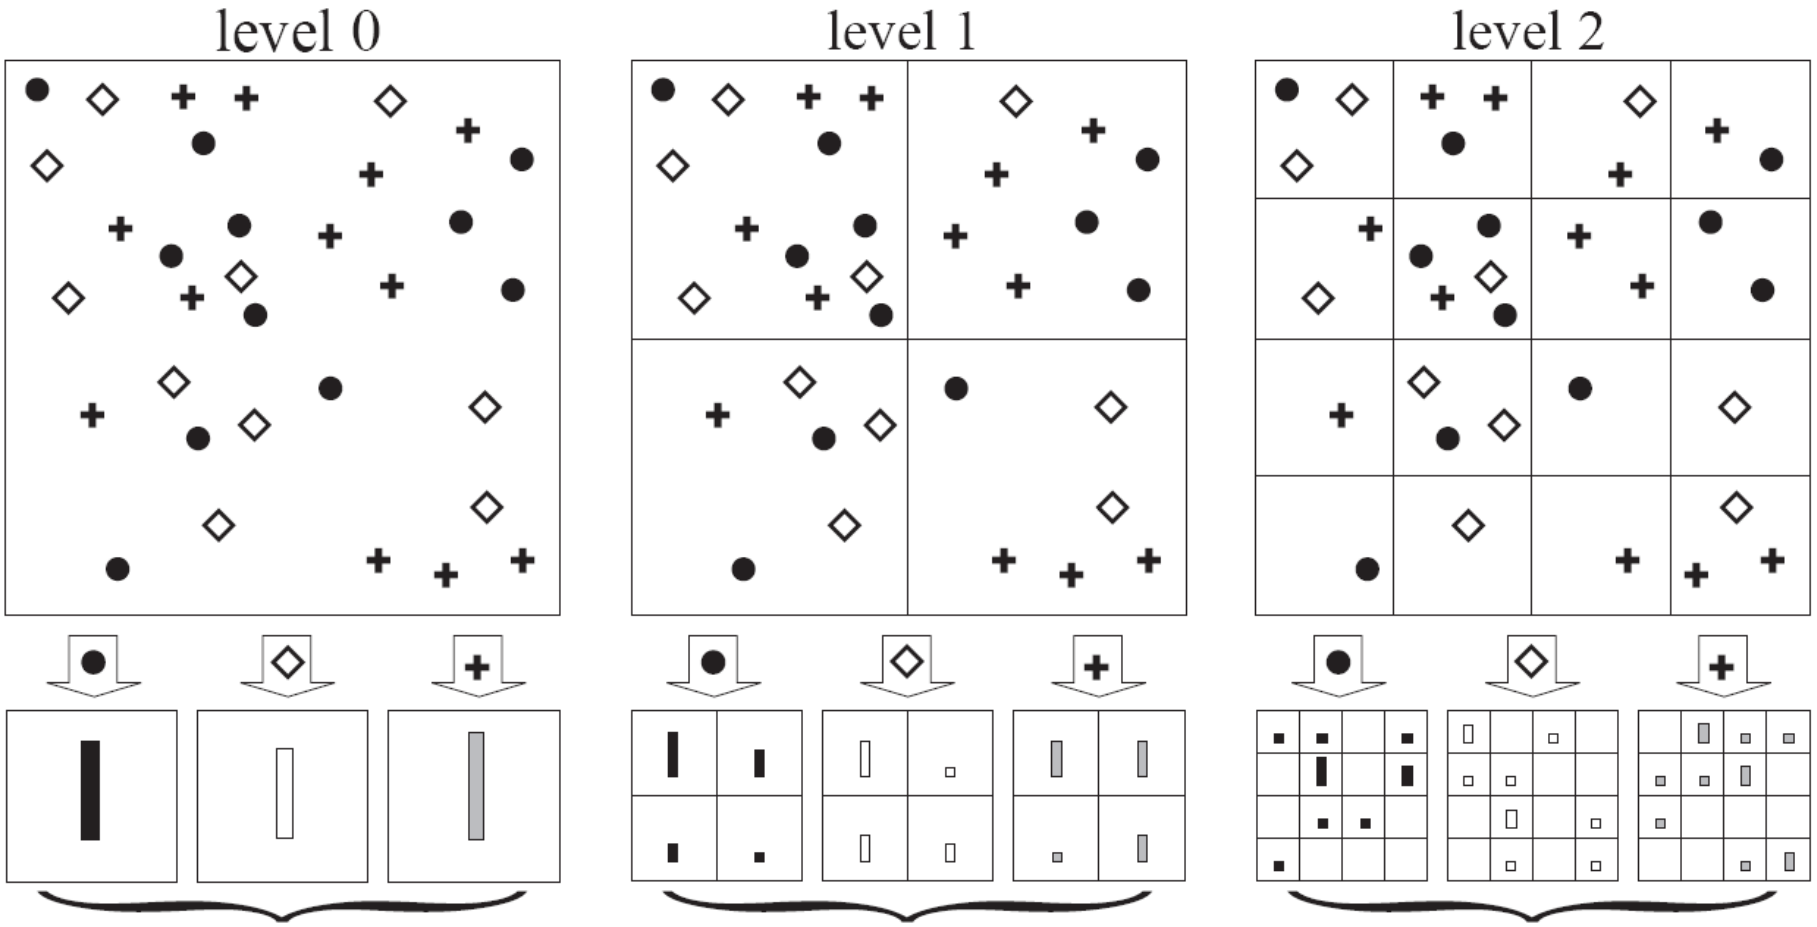

High number of features – PCA to reduce dimensionality

[Lazebnik et al. CVPR 2006]

Pooling

- Average/max pooling

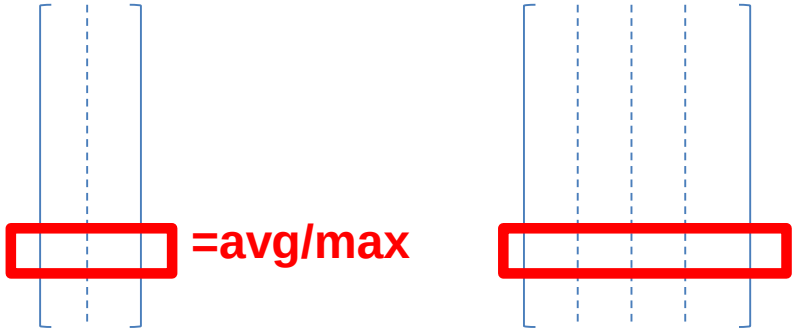

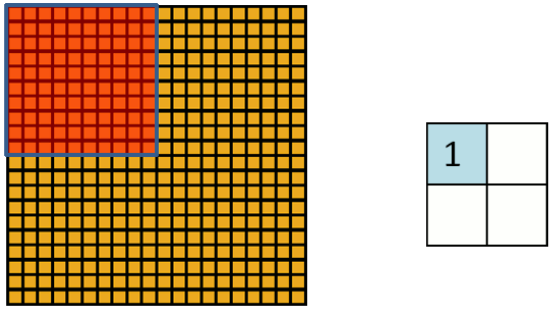

Source: Unsupervised Feature Learning and Deep Learning

- Second-order pooling [Joao et al. PAMI 2014]

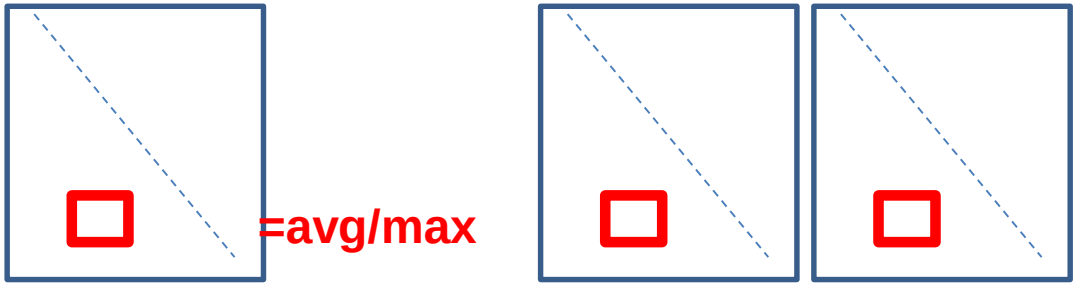

$$\mathbf{G}_{avg}(R_j) = \frac{1}{|F_{R_j}|} \sum_{i:(\mathbf{f}_i \in R_j)} \mathbf{x}_i \cdot \mathbf{x}_i^\top$$

$$\mathbf{G}_{max}(R_j) = \max_{i:(\mathbf{f}_i \in R_j)} \mathbf{x}_i \cdot \mathbf{x}_i^\top$$

2012 ImageNet 1K

(Fall 2012)

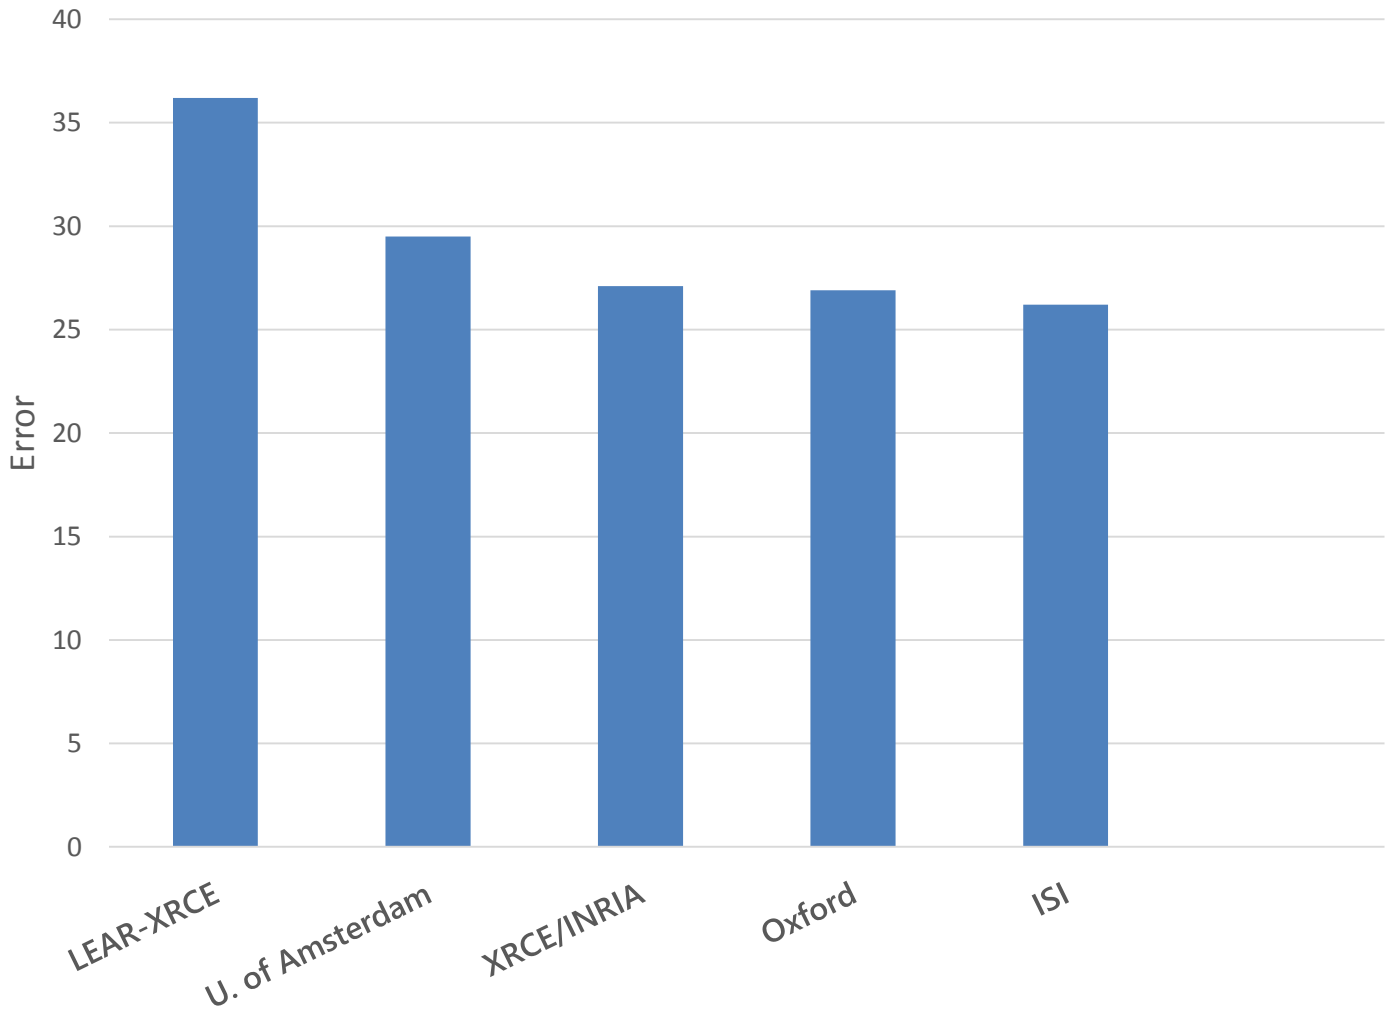

2012 ImageNet 1K

(Fall 2012)

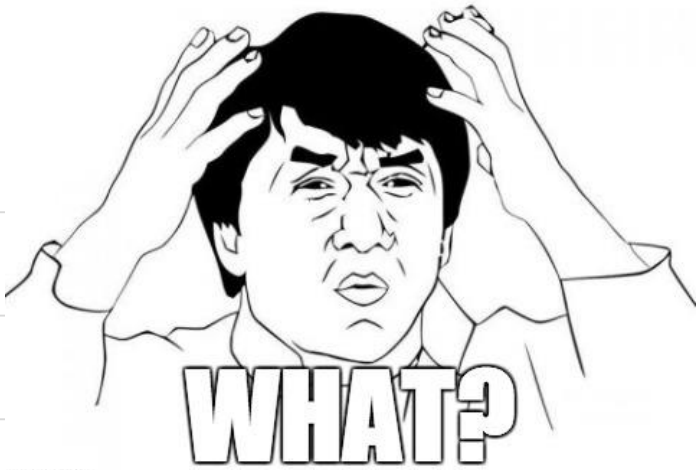

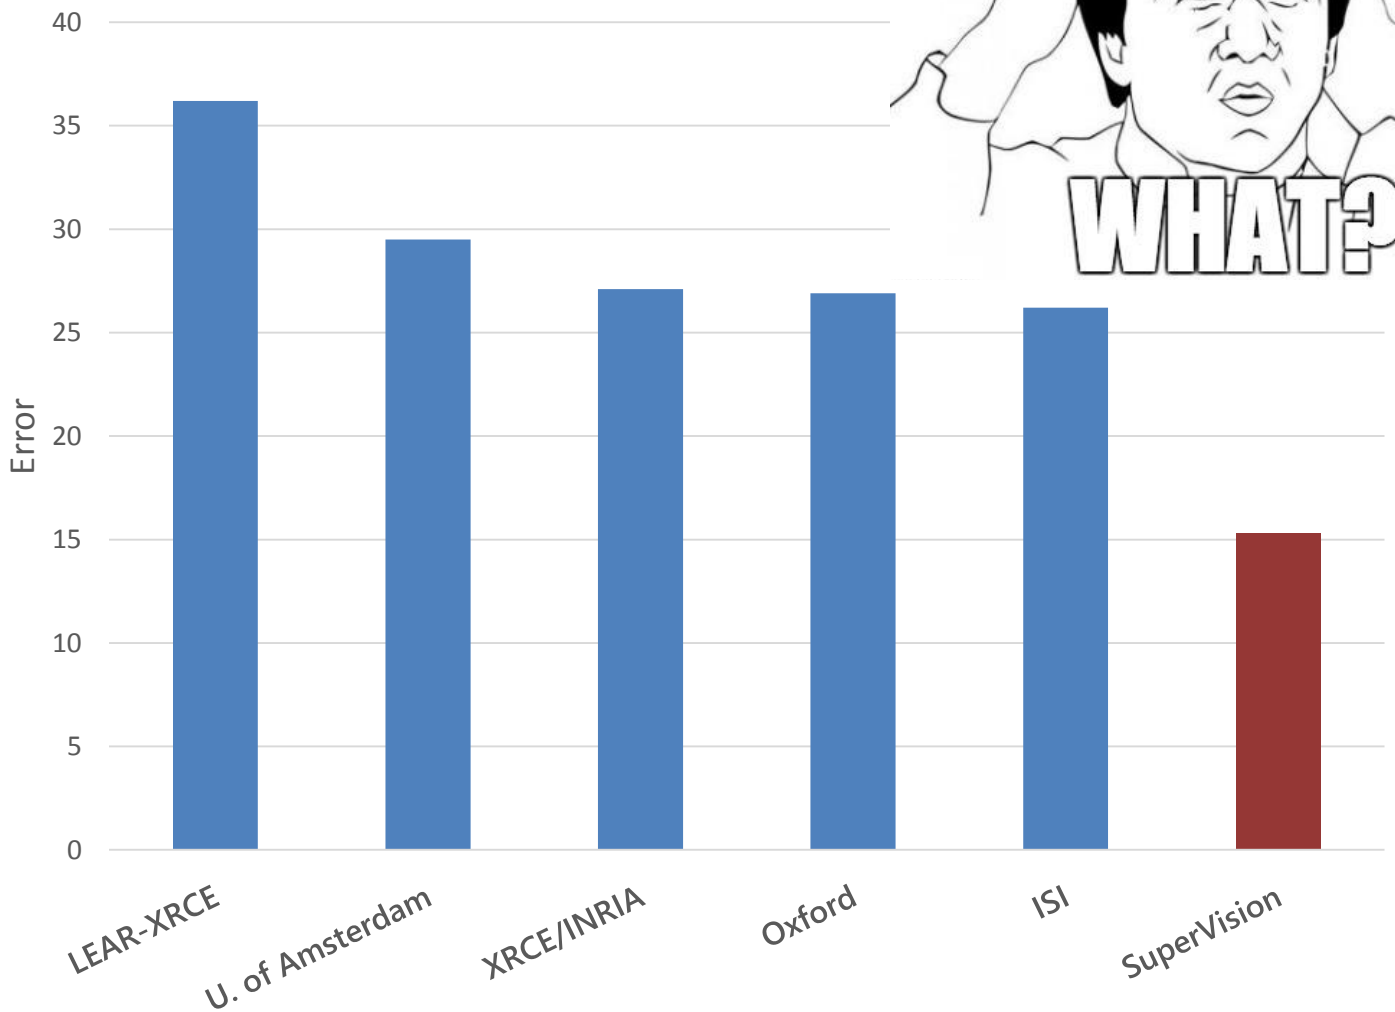

Shallow vs. deep learning

- Engineered vs. learned features

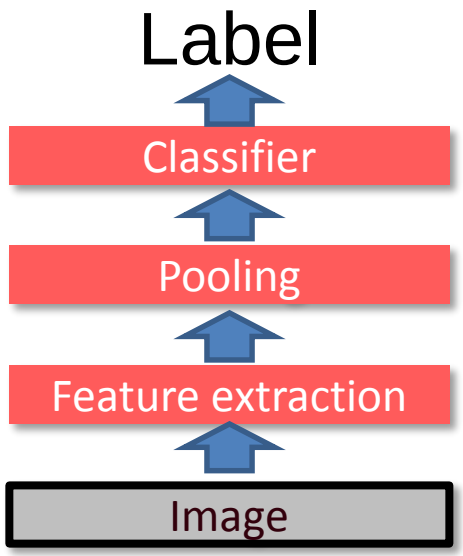

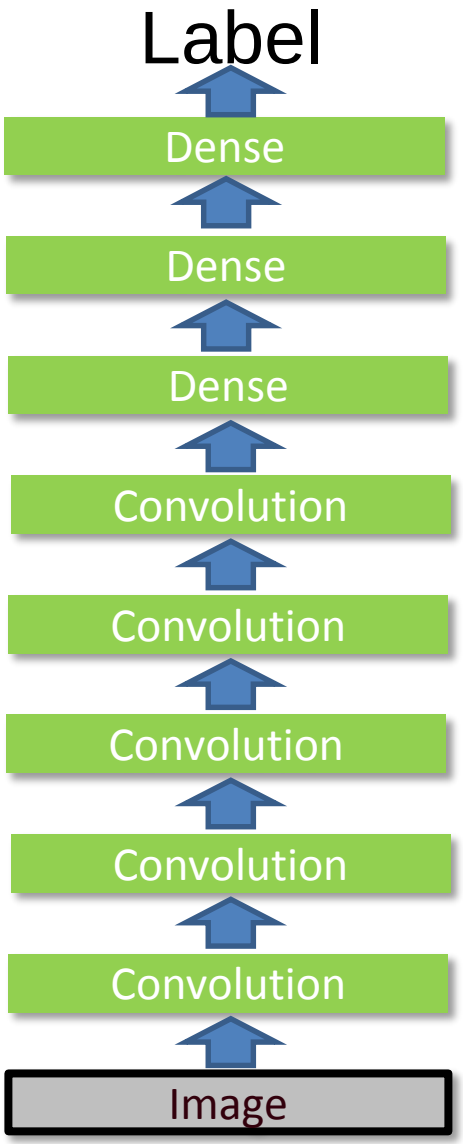

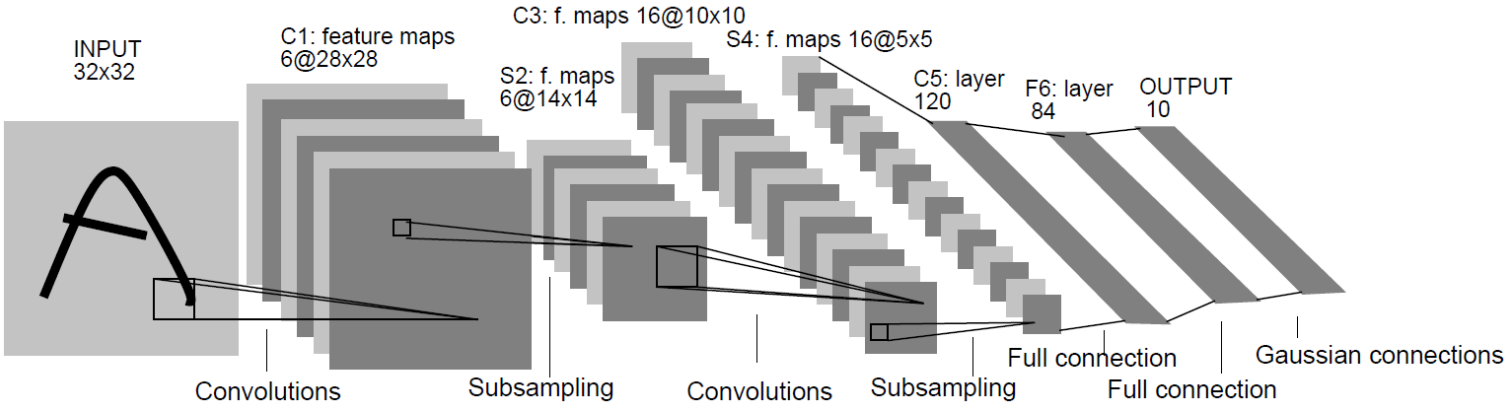

Gradient-Based Learning Applied to Document Recognition, LeCun, Bottou, Bengio and Haffner, Proc. of the IEEE, **1998**

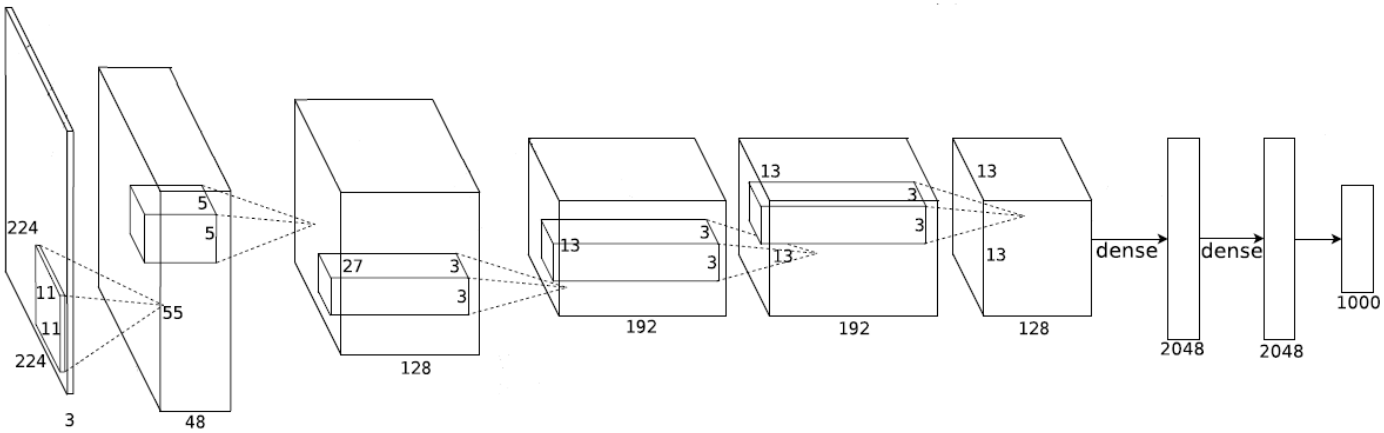

Imagenet Classification with Deep Convolutional Neural Networks, Krizhevsky, Sutskever, and Hinton, NIPS **2012**

Slide Credit: L. Zitnick

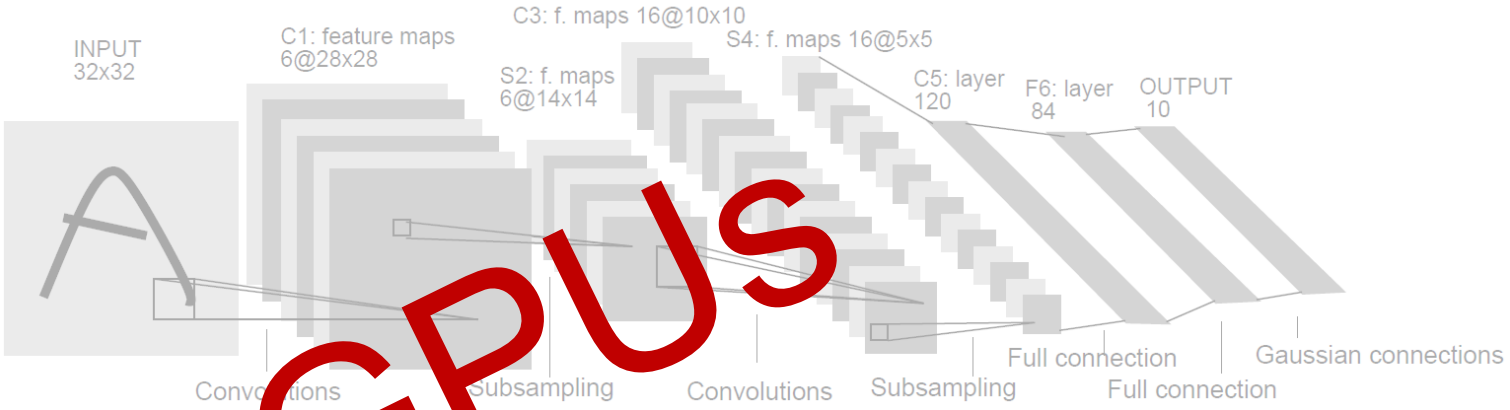

Gradient-Based Learning Applied to Document Recognition, LeCun, Bottou, Bengio and Haffner, Proc. of the IEEE, **1998**

GPUs + Data*

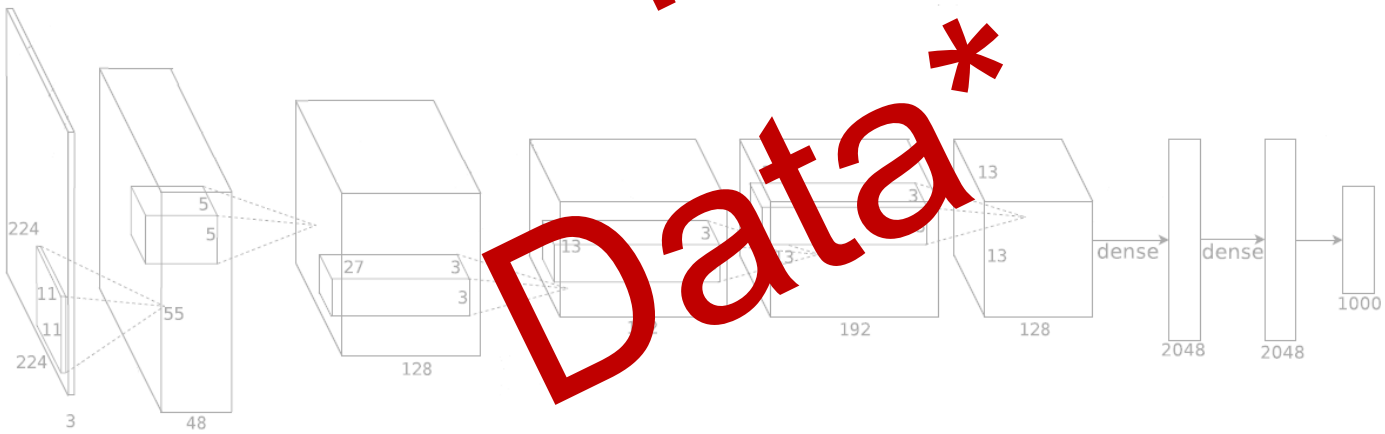

Imagenet Class Networks, Krizh

* Rectified activations and dropout

Slide Credit: L. Zitnick

Convolutional activation features

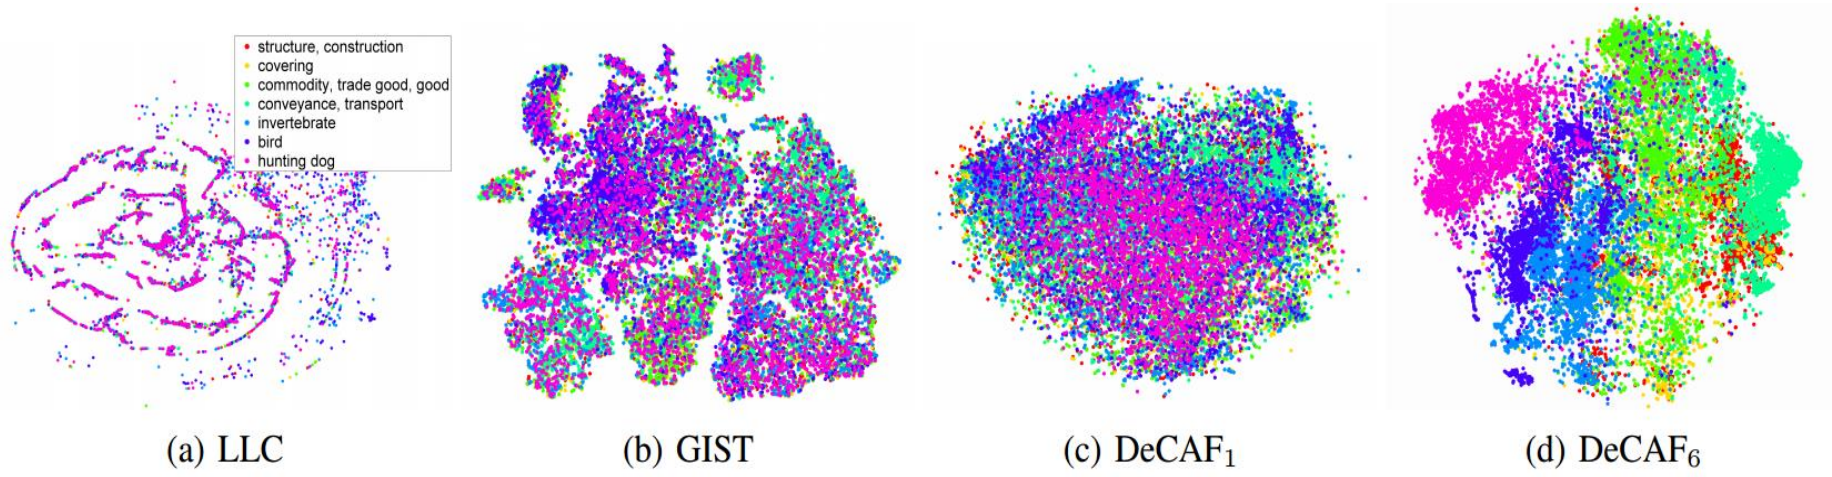

(a) LLC (b) GIST (c) $DeCAF_1$ (d) $DeCAF_6$

. This figure shows several t-SNE feature visualizations on the ILSVRC-2012 validation set. (a) LLC , (b) GIST, and features derived from our CNN: (c) DeCAF1, the first pooling layer, and (d) DeCAF6, the second to last hidden layer (best viewed in color).

[Donahue et al. ICML 2013]

Convolutional activation features

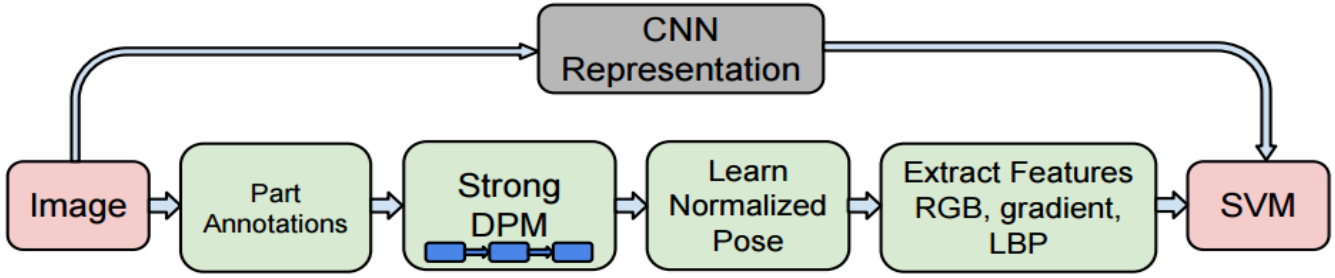

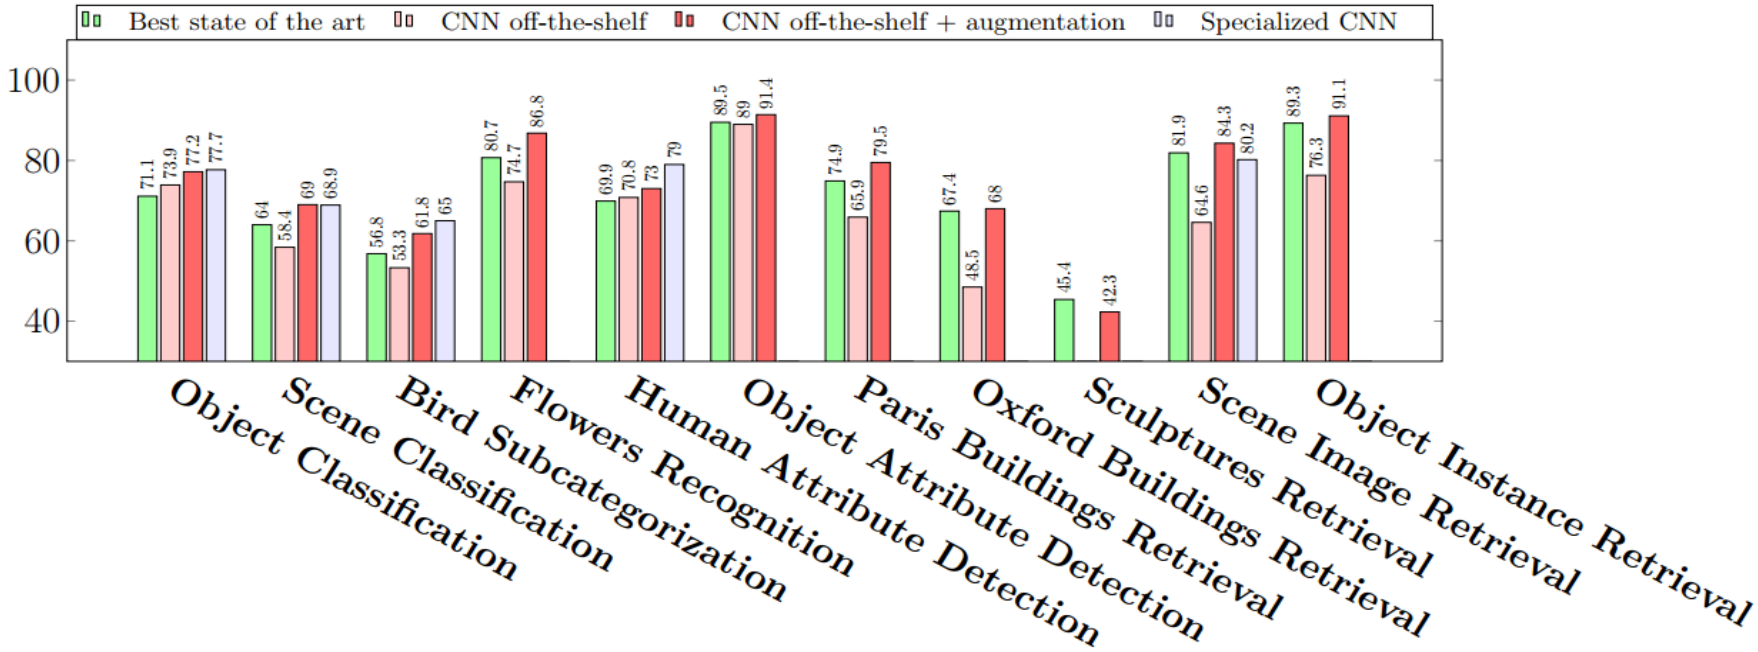

CNN Features off-the-shelf: an Astounding Baseline for Recognition [Razavian et al. 2014]

Things to remember

- Visual categorization help transfer knowledge

- Image features
  - Coverage, concision, directness
  - Color, gradients, textures, motion, descriptors
  - Histogram, feature encoding, and pooling
  - CNN as features

- Image/region categorization